

Time Series Forecasting

Part 1: Extracting Trend & Seasonal Components

BC2407 Seminar 9

Reference

- Avril Coghlan (2018), A Little Book of R for Time Series
 - <https://a-little-book-of-r-for-time-series.readthedocs.io/en/latest/>
- Chew C.H. (2020) Artificial Intelligence, Analytics and Data Science, Vol. 2.
 - Est. Q4 2020.

What is a Time Series?

- A vector of data measured across time.
- Examples:
 - Share Price at end of each **day** since IPO till today.
 - Height of a child vs **months** since birth till last month.
 - Revenue of a company by **quarters** since incorporation till last quarter.
 - GDP by **year** since independence till last year.

Difficulty in Forecasting Time Series

- Concerns:
 - To forecast (i.e. to extrapolate) from given data.
 - Circumstances in future might be different compared to existing time series data.
- Industry Practice:
 - Forecast is just a baseline.
 - To be adjusted based on
 - New data arrivals.
 - New events that has a significant impact.

Time Series Forecasting methods

1. Moving Average Method
2. Decomposition Methods
 1. Trend
 2. Seasonal
 3. Random Error
3. Exponential Smoothing Methods
4. Causal Methods
5. ARIMA method

Our course (part 1) will mainly cover the first 3 methods in depth and briefly mention 4 and 5.

Autocorrelation

- Can lagged data points be used to predict the next data points?
- Lag k : k period behind
- Example for monthly data:
 - Today is Mar 2020, Lag 1: Feb 2020, Lag 2: Jan 2020, etc.
- Autocorrelation by lag k
 - Correlation between the vector X_t and vector X_{t-k}

Example of Autocorrelation of Phone Sales by Lags 1, 2 & 3.

	A	B	C	D	E	F	G	H	I	J
1	Month	Sales	Lag 1	Lag 2	Lag 3					
2	Jan-2013	226								
3	Feb-2013	254	226							
4	Mar-2013	204	254	226						
5	Apr-2013	193	204	254	226					
6	May-2013	191	193	204	254					
7	Jun-2013	166	191	193	204					
8	Jul-2013	175	166	191	193					
9	Aug-2013	217	175	166	191					
10	Sep-2013	167	217	175	166					
11	Oct-2013	192	167	217	175		Lag1	Autocorrelation		
12	Nov-2013	127	192	167	217		Lag2	0.357		=CORREL(B3:B49,C3:C49)
13	Dec-2013	148	127	192	167		Lag3	0.084		=CORREL(B4:B49,D4:D49)
14	Jan-2014	184	148	127	167			0.089		=CORREL(B5:B49,E5:E49)
15	Feb-2014	200	184	148	127					
16	Mar-2014	185	200	184	148					
17	Apr-2014	245	185	200	184					
18	May-2014	177	245	185	200					
19	Jun-2014	185	177	245	185					
20	Jul-2014	175	185	177	245					
21	Aug-2014	181	175	185	177					
22	Sep-2014	179	181	175	185					
23	Oct-2014	181	179	181	179					
24	Nov-2014	175	181	179	181					
25	Dec-2014	181	175	181	175					
26	Jan-2015	175	181	175	181					
27	Feb-2015	181	175	181	175					
28	Mar-2015	175	181	175	181					
29	Apr-2015	181	175	181	175					
30	May-2015	175	181	175	181					
31	Jun-2015	181	175	181	175					
32	Jul-2015	175	181	175	181					
33	Aug-2015	181	175	181	175					
34	Sep-2015	175	181	175	181					
35	Oct-2015	181	175	181	175					
36	Nov-2015	175	181	175	181					
37	Dec-2015	181	175	181	175					
38	Jan-2016	175	181	175	181					
39	Feb-2016	181	175	181	175					
40	Mar-2016	175	181	175	181					
41	Apr-2016	181	175	181	175					
42	May-2016	175	181	175	181					
43	Jun-2016	181	175	181	175					
44	Jul-2016	175	181	175	181					
45	Aug-2016	181	175	181	175					
46	Sep-2016	175	181	175	181					
47	Oct-2016	185	175	181	179					
48	Nov-2016	245	185	175	181					
49	Dec-2016	177	245	185	175					

Note that each successive lag just "pushes the variable down" by a row. As an example, the lag 2 residual for July-2013 is just the residual 2 months ago, i.e., for May-2013.

Create Time Series Object in Base R via `ts()`

Time-Series Objects

Description

The function `ts` is used to create time-series objects.

`as.ts` and `is.ts` coerce an object to a time-series and test whether an object is a time series.

Usage

```
ts(data = NA, start = 1, end = numeric(), frequency = 1,  
    deltat = 1, ts.eps = getOption("ts.eps"), class = , names = )  
as.ts(x, ...)  
is.ts(x)
```

Arguments

<code>data</code>	a vector or matrix of the observed time-series values. A data frame will be coerced to a numeric matrix via <code>data.matrix</code> . (See also 'Details'.)
<code>start</code>	the time of the first observation. Either a single number or a vector of two integers, which specify a natural time unit and a (1-based) number of samples into the time unit. See the examples for the use of the second form.
<code>end</code>	the time of the last observation, specified in the same way as <code>start</code> .
<code>frequency</code>	the number of observations per unit of time.

Create ts object from HDB Sales Data

```
5 library("TTR")                # For MA via SMA()
6 library("forecast")           # To generate h-period ahead forecasts
7
8 setwd("C:/NC/Datasets/ML")
9
10 flatsales.df <- read.csv("5 room flat resale applications.csv")
11
12 # create time series object
13 flatsales.ts <- ts(flatsales.df$Sales.5rm, frequency = 4, start = c(2007,1))
```

	A	B
1	Quarter	Sales 5rm
2	2007-Q1	1402
3	2007-Q2	2305
4	2007-Q3	1901
5	2007-Q4	1667
6	2008-Q1	1574
7	2008-Q2	1997
8	2008-Q3	2172
9	2008-Q4	1578
10	2009-Q1	1506
11	2009-Q2	2712

Only one column
of data values

4 times per year

Data start in
2007, quarter 1

50	2019-Q1	1148
51	2019-Q2	1520
52	2019-Q3	1547
53	2019-Q4	1519

Time Series Object from `ts()`


```
> flatsales.ts
```

	Qtr1	Qtr2	Qtr3	Qtr4
2007	1402	2305	1901	1667
2008	1574	1997	2172	1578
2009	1506	2713	3422	2187
2010	2047	2240	1975	1477
2011	1582	1635	1415	1412
2012	1370	1594	1558	1288
2013	962	1070	969	784
2014	726	913	960	1024
2015	925	1210	1140	1113
2016	1023	1369	1283	1186
2017	1055	1407	1428	1403
2018	1111	1559	1797	1350
2019	1148	1520	1547	1519


Additive or Multiplicative Time Series Model

- Affects how to seasonally adjust (i.e. deseasonalize) the time series.
 - i.e. remove the effects of seasonality
 - Additive implies subtract seasonality effects
 - Multiplicative implies divide seasonality effects
- So as to more cleanly estimate the trend component.
- Finally, the forecast will include both trend and seasonality i.e. need to re-seasonalize
 - Additive implies add seasonality effects
 - Multiplicative implies multiply seasonality effects

Constant Fluctuations Over Time -> Additive Time Series




Changing Fluctuations Over Time -> Multiplicative Time Series



Flat sales time series is additive or multiplicative?

```
17 plot.ts(flatsales.ts, ylab = "Flats sold", xlab = "Quarter-Year",  
18         main = "Sales of 5 Room Resale Flats",  
19         sub = "Source: HDB")  
20
```



Moving Average (MA) Forecast

- A moving average is the mean of the observations in the past few periods, where the number of terms considered in the mean is called the span.
- The larger, the span, the more items are averaged, and thus looks more smoothed.
- Thus, Moving Avg is also considered as a (simple) Smoothing Method.

Setting the Span in Moving Avg


- Setting a span requires some judgment:
 - If you think fluctuations in the series are mainly due to random noise, use a relatively large span.
 - Otherwise, use a smaller span.

Moving Averages with Span = 3 vs 7


What's the differences?

```
24 m.ma3 <- SMA(flatsales.ts, n = 3)
25 plot(m.ma3, main = "Moving Avg Span 3", ylab = "MA3 Forecast")
26
27 m.ma7 <- SMA(flatsales.ts, n = 7)
28 plot(m.ma7, main = "Moving Avg Span 7", ylab = "MA7 Forecast")
```

Moving Avg Span 3



Moving Avg Span 7




MA7 is smoother (less fluctuations) and start later than MA3.

MA Based Decomposition of Time Series into Trend and Seasonal Components

```
30 # Classical Seasonal Decomposition by Moving Averages
31 m.ma.mul <- decompose(flatsales.ts, type = "multiplicative")
32 plot(m.ma.mul)
```

Decomposition of multiplicative time series



MA vs Exponential Smoothing

- MA essentially ignores data beyond the span and only considers the most recent “window” of data.
- In contrast, exponential smoothing considers all the data values since the start date in the dataset, but weighs each data value.
 - The more recent data has higher weights than older data.

Exponential Smoothing Methods

1. Simple Exponential Smoothing
 - for a series without trend and seasonality.
2. Holt's method
 - for a series with trend but no seasonality.
3. Winters' method (or Holt-Winters' method)
 - for a series with seasonality and possibly trend.

Simple Exponential Smoothing

- Every exponential model has at least one **smoothing parameter**, between 0 and 1.
- Simple exponential smoothing has a single smoothing constant denoted by α .
- The *level* of the series at time t (L_t) is an estimate of where the series would be at time t if there were no random noise.
- The simple exponential method is defined by the following two equations:

$$L_t = \alpha Y_t + (1 - \alpha)L_{t-1}$$

$$F_{t+k} = L_t$$

1. *What is the role played by α ?*
2. *How many data points are used in calculating L_t ?*
3. *Is L_t used to forecast Y_t ?*

- The k -period-ahead forecast, F_{t+k} , of Y_{t+k} made at period t is essentially the most recently estimated level, L_t .

Simple Exponential Smoothing on 5 Room Flat Sales

Est. 15 mins

- Using Excel, execute simple exponential smoothing on the 5 room resale flat data with:
 - $\alpha = 0.2$
 - $\alpha = 0.8$
- What value did you set for the first value of L_t to kickstart the model? *[Note: more than 1 way]*
- What did you observe when you use the different alpha values above? Which value is more suitable in your opinion?
- What is your one year ahead forecast of sales (i.e. 2020 Q1 to 2020 Q4) with each of the two alphas?

Holt's method

- When there is a trend in the series, Holt's method deals with it explicitly by including a trend term, T_t , and its corresponding smoothing constant β .
 - The interpretation of L_t is exactly as before.
 - The interpretation of T_t is that it represents an estimate of the *change* in the series from one period to the next.
- The equations for Holt's model are:

$$L_t = \alpha Y_t + (1 - \alpha)(L_{t-1} + T_{t-1})$$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

$$F_{t+k} = L_t + kT_t$$

What value can be set for the first value of T_t to kickstart the model? [More than 1 way]

Seasonality

- Seasonality is the consistent pattern that repeats each year. Eg: month-to-month, quarter-to-quarter, events.
 - The easiest way to check for seasonality is graphically: Look for a *regular* pattern of ups and/or downs.
- There are three methods for dealing with seasonality:
 - Winters' exponential smoothing model
 - Deseasonalize the data (then use any forecasting method to model the deseasonalized data and finally, “reseasonalize” these forecasts)
 - Multiple regression with dummy variables for the seasons

Winter's method (aka Holt-Winters' method)

- Winters' model is very similar to Holt's model, but it adds seasonal indexes and a corresponding smoothing constant γ .
 - The new smoothing constant controls how quickly the method reacts to observed changes in the seasonality.
 - If the constant is small, the method reacts slowly.
 - If it is large, the method reacts more quickly.
 - The equations for this method are shown below:

$$L_t = \alpha \frac{Y_t}{S_{t-M}} + (1 - \alpha)(L_{t-1} + T_{t-1})$$

$$T_t = \beta(L_t - L_{t-1}) + (1 - \beta)T_{t-1}$$

$$S_t = \gamma \frac{Y_t}{L_t} + (1 - \gamma)S_{t-M}$$

$$F_{t+k} = (L_t + kT_t)S_{t+k-M}$$

1. Is this an additive model or multiplicative model?
2. Why is the subscript for S : $t - M$ instead of $t - 1$?
3. What value can be set for the first values of S_t to kickstart the model? [More than 1 way]

Simple Exponential Smoothing

```
35 # Simple Exponential Smoothing -----
36 m.ses <- Holtwinters(flatsales.ts, seasonal = "multiplicative", beta=F, gamma=F)
37
38 m.ses
39 ## Optimal value of alpha = 0.5288797
40
41 plot(m.ses)
42
43 m.ses$fitted
44
45 m.ses$alpha*1519 + (1-m.ses$alpha)*1485.8267
46 ## verifying the meaning of coef = 1503.371 is  $L_{\{last\}t}$ 
47
48 RMSE.ses <- round(sqrt(m.ses$SSE/nrow(m.ses$fitted)))
49 ## 359
```

```
> m.ses
Holt-winters exponential smoothing without trend and with
out seasonal component.
```

```
Call:
Holtwinters(x = flatsales.ts, beta = F, gamma = F, season
al = "multiplicative")
```

```
Smoothing parameters:
alpha: 0.5288797
beta : FALSE
gamma: FALSE
```


```
Coefficients:
      [,1]
a 1503.371
```

How to get this?

What is this?

1. Min SSE of one period ahead forecast.
2. The last value of L_t

```
41 plot(m.ses, main = "Simple Exp Smoothing with optimized alpha = 0.5288797")
```



Which coloured line rep observed vs one period ahead SES forecast?

Black line rep observed; Red line rep one period ahead SES forecast.

Holt Smoothing Method

```
51 # Holt's method -----  
52 m.holt <- Holtwinters(flatsales.ts, seasonal = "multiplicative", gamma=F)  
53  
54 m.holt  
55 ## Optimal value of alpha = 1, beta = 0.2414013  
56  
57 plot(m.holt, main = "Holt Smoothing with optimized alpha = 1, Beta = 0.2414013")  
58  
59 RMSE.holt <- round(sqrt(m.holt$SSE/nrow(m.holt$fitted)))
```

```
> m.holt
```

Holt-Winters exponential smoothing with trend and without seasonal component.

call:

```
Holtwinters(x = flatsales.ts, gamma = F, seasonal = "multiplicative")
```


Smoothing parameters:

```
alpha: 1  
beta : 0.2414013  
gamma: FALSE
```

Coefficients:

```
      [,1]  
a 1519.00000  
b  21.43562
```

```
57 plot(m.holt, main = "Holt Smoothing with optimized alpha = 1, beta = 0.2414013")
```



Winters Smoothing Method

```
62 # winter's method -----
63 m.winters <- Holtwinters(flatsales.ts, seasonal = "multiplicative")
64
65 m.winters
66 ## Optimal value of alpha = 0.8981024, beta = 0, gamma = 1.
67
68 plot(m.winters, main = "Winters Smoothing with optimized alpha = 0.8981024, beta = 0, gamma = 1.")
69
70 RMSE.winters <- round(sqrt(m.winters$SSE/nrow(m.winters$fitted)))
71 ## Winters method has the lowest RMSE
```

```
> m.winters
```

Holt-Winters exponential smoothing with trend and multiplicative seasonal component.

Call:

```
Holtwinters(x = flatsales.ts, seasonal = "multiplicative")
```


Smoothing parameters:

```
alpha: 0.8981024
beta : 0
gamma: 1
```

Coefficients:

```
      [,1]
a 1617.4851043
b  -0.1250000
s1  0.8314238
s2  1.0986646
s3  1.0853405
s4  0.9391122
```

```
68 plot(m.winters, main = "Winters Smoothing with optimized alpha = 0.8981024, beta = 0, gamma = 1.")
```



Forecast h periods ahead


```
73 # Forecast 4 periods ahead -----  
74 m.winters.forecasts <- forecast(m.winters, h = 4)  
75  
76 m.winters.forecasts  
77  
78 plot(m.winters.forecasts, main = "4 Period Ahead Forecasts based on Winters Method")
```

```
> m.winters.forecasts
```

	Point Forecast	Lo 80	Hi 80	Lo 95	Hi 95
2020 Q1	1344.712	1016.6657	1672.758	843.0086	1846.415
2020 Q2	1776.799	1267.6998	2285.898	998.1990	2555.399
2020 Q3	1755.115	1171.8178	2338.412	863.0389	2647.191
2020 Q4	1518.530	-994.0947	4031.156	-2324.1981	5361.259


```
78 plot(m.winters.forecasts, main = "4 Period Ahead Forecasts based on winters Method")
```

4 Period Ahead Forecasts based on Winters Method



Blue Points: Point Forecast, Dark Grey Interval: 80% C.I., Light Grey Interval: 95% C.I.

Other Forecasting Methods

- Causal models
 - From theory or domain knowledge
 - Involve X s and t to forecast Y
- ARIMA models
 - Statistical analysis of autocorrelations
 - Requirement: Transform time series to a Stationary Time Series first before modelling the AR and MA components.
 - Stationary Time Series
 - Trend is constant over time.
 - Variance is constant over time.
 - Autocorrelation is constant over time.

Summary

- Forecasting is inherently extrapolation, unlike other models e.g. Linear Reg, CART, MARS, etc.
- MA models
- Exponential Smoothing models
 - SES
 - Holt
 - Winters
- Industry Practice
 - Model is just a baseline forecast
 - Adjust based on new data and new events.