

CE2001/ CZ2001: Algorithms

Graphs

Ke Yiping, Kelly

Greedy Algorithms

Learning Objectives

At the end of this lecture, students should be able to:

- Explain the strategy of Greedy algorithms
- Solve single-source shortest paths problem using Dijkstra's algorithm
- Prove the correctness of Dijkstra's algorithm
- Describe Prim's algorithm for finding minimum spanning trees (MSTs)
- Prove the correctness of Prim's algorithm

Greedy Algorithms

- In optimisation problems, the algorithm needs to make a series of choices whose overall effect is to minimise the total cost, or maximise the total benefit, of some system.
- There is a class of algorithms, called the **greedy algorithms**, in which we can find a solution by using only knowledge available at the time when the next choice (or guess) must be made.
- Each individual choice is the best within the knowledge available at the time.

Greedy Algorithms

- Each individual choice is not very expensive to compute.
- A choice cannot be undone, even if it is found to be a bad choice later.
- Greedy algorithms cannot guarantee to produce the optimal solution for a problem.

5

Dijkstra's Algorithm

Dijkstra's Algorithm

Shortest Path Problem:

The problem of finding the shortest path from one vertex in a graph to another vertex. "Shortest" may be the least number of edges, or the least total weight, etc.

Dijkstra's Algorithm:

This is an algorithm to find the shortest paths from a single source vertex to all other vertices in a **weighted, directed** graph. All weights must be **nonnegative**.

7

Dijkstra's Algorithm

Dijkstra's algorithm keeps two sets of vertices:

- **S** ---- the set of vertices whose shortest paths from the source node have already been determined [they form the tree]
- **V - S** ---- the remaining vertices

The other data structures needed are:

- **d** ---- array of estimates for the lengths of shortest paths from source node to all vertices
- **pi** ---- an array of predecessors for each vertex

8

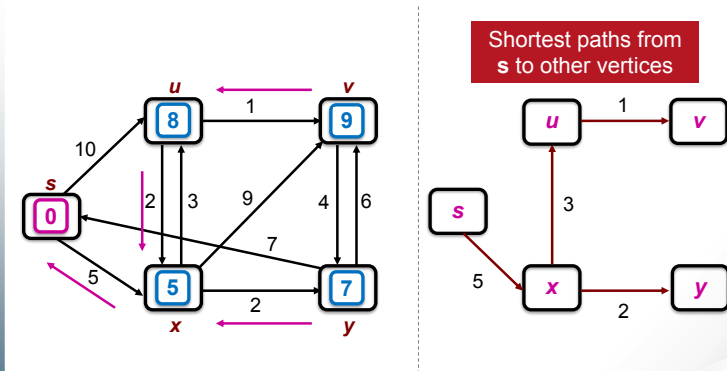
Basic Steps

The basic steps are:

1. Initialise **d** and **pi**
2. Set **S** to empty
3. While there are still vertices in **V - S**
 - i. Move **u**, the vertex in **V - S** that has the shortest path estimate from source, to **S**
 - ii. For all the vertices in **V - S** that are connected to **u**, update their estimates of shortest distances to the source

9

A Toy Example



10

Pseudocode of Dijkstra's Algorithm

```

Dijkstra_ShortestPath ( Graph G, Node source ) {
  for each vertex v {
    d[v] = infinity;
    pi[v] = null pointer;
    S[v] = 0;
  }
  d[source] = 0;
  put all vertices in priority queue, Q, in d[v]'s increasing order;
  while not Empty(Q) {
    u = ExtractCheapest(Q);
    S[u] = 1; /* Add u to S */
  }
}

```

11

Pseudocode of Dijkstra's Algorithm

```

for each vertex v adjacent to u
  if (S[v] ≠ 1 and d[v] > d[u] + w[u, v]) {
    remove v from Q;
    d[v] = d[u] + w[u, v];
    pi[v] = u;
    insert v into Q according to its d[v];
  }
} // end of while loop
}

```


Worst case time complexity of Dijkstra's algorithm is $O(|V|^2)$
(analysis not required).

12

Proof of Correctness

Property of Shortest Path

Lemma 1: In a weighted graph G , suppose that a shortest path from x to z consists of a path P from x to y followed by a path Q from y to z . Then P is a shortest path from x to y and Q is a shortest path from y to z .



Proof (By Contradiction):

Assume that P is not the shortest path from x to y . Then there will be another path from x to y , P' which is shorter than P . As a result P' followed by Q will be a path **shorter** than P followed by Q . But it was known that P followed by Q is the **shortest** path. Contradiction. Same can be said about Q .

13

Greedy choice is optimal

Theorem D1: Let $G = (V, E, W)$ be a weighted graph with nonnegative weights. Let S be a subset of V and let s be a member of S . Assume that $d[y]$ is the shortest distance in G from s to y , for each y in S . Let z be the next vertex chosen to go into S . If edge (y, z) is chosen to minimise $d[y] + W(y, z)$ over all edges with one vertex in S and one vertex in $V - S$, then the path consisting of a shortest path from s to y followed by the edge (y, z) is the shortest path from s to z .

Proof:

We will show that there is no other path from s to z that is shorter.

14

Greedy choice is optimal

Proof of Theorem D1 (continued)


$P: s \rightarrow y \rightarrow z$ (shortest path for z)

$$W(P) = d[y] + W(y, z)$$

$P': s \rightarrow y \rightarrow u \rightarrow \dots \rightarrow z$
(an alternative shortest path)

$$W(P') = d[y] + W(y, u) + \text{distance from } u \text{ to } z$$

Because $d[y] + W(y, u) \geq d[y] + W(y, z)$,
and distance from u to z is nonnegative,
therefore $W(P) \leq W(P')$.



Edge (y, z) is chosen to minimise $d[y] + W(y, z)$ over all edges with one vertex in S and one vertex in $V - S$.

15

Greedy choice is optimal

Proof of Theorem D1 (continued)

Let P be a shortest path from s to y followed by edge (y, z)

Let $W(P)$ = the distance travelled along P

Let $P' = \text{any shortest path different from } P, \text{ i.e., } P' = s, z_1, \dots, z_k, \dots, z$


Assume that z_k is the first vertex in P' not in set S .

$$W(P) = d[y] + W(y, z)$$

$$W(P') = d[z_{k-1}] + W(z_{k-1}, z_k) + \text{distance from } z_k \text{ to } z$$

Note that: $d[z_{k-1}] + W(z_{k-1}, z_k) \geq d[y] + W(y, z)$

Since distance from z_k to z is non-negative,
therefore, $W(P) \leq W(P')$.



16

Theorem D2 and Proof

Theorem D2: Given a directed weighted graph G with nonnegative weights and a source vertex s , Dijkstra's algorithm computes the shortest distance from s to each vertex of G that is reachable from s .

Proof (By induction):

We will show by induction that as each vertex v is added into set S , $d[v]$ is the shortest distance from s to v .

Basis:

The algorithm assigns $d[s]$ to zero when the source vertex s is added to S . So $d[s]$ is the shortest distance from s to s when S has the first vertex in it.

17


Theorem D2 and Proof (Continued)

Inductive Hypothesis:

Assume the theorem is true when S has k vertices. That is, assume $v_0, v_1, v_2, \dots, v_{k-1}$ are added where $d[v_1], d[v_2], \dots$ are the shortest distances.

When v_k is chosen by Dijkstra's algorithm, it means an edge (v_i, v_k) , where $i \in \{0, 1, 2, \dots, k-1\}$, is chosen to minimise $d[v_i] + W(v_i, v_k)$ among all edges with one vertex in S and one vertex not in S .

By Theorem D1, $d[v_k]$ is the shortest distance from source to v_k . So the theorem is true when S has $k + 1$ vertices.



18

Minimum Spanning Tree

Minimum Spanning Tree

Definition of Spanning Tree


A connected, acyclic subgraph containing all the vertices of a graph.

Definition of Minimum Spanning Tree

A minimum-weight spanning tree in a weighted graph.


20

Spanning Tree



21

Minimum Spanning Tree



22

Main Idea of Prim's Algorithm

Prim's Algorithm

- It works on undirected graph.
- It builds upon a single partial minimum spanning tree, at each step adding an edge connecting the vertex nearest to but not already in the current partial minimum spanning tree.
- At first a vertex is chosen, this vertex will be the first node in T .
- Set P is initialised: P = set of vertices not in tree T but are adjacent to some vertices in T .

23

Main Idea of Prim's Algorithm


Prim's Algorithm (Cont.)

- In every iteration in the Prim's Algorithm, a new vertex u from set P will be connected to the tree T . The vertex u will be deleted from the set P . The vertices adjacent to u and not already in P will be added to P .
- When all vertices are connected into T , P will be empty. This means the end of the algorithm.
- The new vertex in every iteration will be chosen by using **greedy** method, i.e. among all vertices in P which are connected to some vertices already inserted in the tree T but themselves are not in T , we choose one with the minimum cost.

24

An Example of Prim's Algorithm

Prim's MST



Black vertices: unseen vertices Pink vertices: tree vertices Blue vertices: fringe vertices

3 subsets of vertices


Prim's Algorithm classifies vertices into three disjoint categories:

- **Tree vertices** - in the tree being constructed so far
- **Fringe vertices** - not in the tree but adjacent to some vertices in the tree
- **Unseen vertices** - all others

Greedy choice of Prim's Algo

- Key step in the algorithm is the selection of a vertex from the fringe (which, of course, depends on the weights on incident edges).
- Prim's Algorithm always chooses a minimum weight edge from **tree** vertex to **fringe** vertex.

Main Idea of Prim's Algorithm



Choose $\min(w, x, y, z)$

Pseudocode of Prim's Algo

```

primMST(G, s, n) // outline of Prim's algorithm
{
  Initialise all vertices as unseen.
  Reclassify s as tree vertex.
  Reclassify all vertices adjacent to s as fringe.
  While (there are fringe vertices)
  {
    Select an edge of minimum weight between a tree
    vertex t and a fringe vertex v;
    Reclassify v as tree; add edge tv to the tree;
    Reclassify all unseen vertices adjacent to v as fringe.
  }
}
  
```

29

Implementing Prim's Algo

Data Structures Used:

- Array **d**: distance of a fringe vertex from the tree
- Array **pi**: vertex connecting a fringe vertex to a tree vertex
- Array **S**: whether a vertex is in the minimum spanning tree being built
- Priority queue **pq**: queue of fringe vertices in the order of the distances from the tree

At the end of the algorithm, array **pi** has the minimum spanning tree.

30

Implementing Prim's Algo

```


primMST(G, s, n) {
  initialise priority queue pq as empty;
  for each vertex v {
    d[v] = infinity; S[v] = 0;
    pi[v] = null pointer; }
  d[s] = 0; S[s] = 1;
  insert(pq, s, 0);
  while (pq is not empty) {
    u = getMin(pq); deleteMin(pq);
    S[u] = 1;
    updateFringe(pq, G, u); }
}
  
```

31

Update Fringe Set of Vertices

```

updateFringe(pq, G, v) {
  for all vertices w adjacent to v {
    if (S[w] != 1) { //if w is not a tree vertex
      newWgt = weight of edge vw;
      if (d[w] == infinity) {
        d[w] = newWgt; pi[w] = v;
        insert(pq, w, newWgt);
      } else if (newWgt < d[w]) {
        d[w] = newWgt; pi[w] = v;
        decreaseKey(pq, w, newWgt);
      }
    } // if w is not a tree vertex
  } // for all vertices
}
  
```



32

MST Property

Minimum Spanning Tree Property

Let T be a spanning tree of G , where $G = (V, E, W)$ is a connected, weighted graph. Suppose that for every edge (u, v) of G that is not in T , if (u, v) is added to T it creates a cycle such that (u, v) is a maximum-weight edge on that cycle. Then T has the **Minimum Spanning Tree Property** (or **MST Property**, in short).

33

Lemma 1 and Proof

Lemma 1: In a connected weighted graph $G = (V, E, W)$, if T_1 and T_2 are two spanning trees that have the MST property, then they have the same total weight.

Proof by induction on k , the number of edges in T_1 but not T_2 (there are also k edges in T_2 but not in T_1).

Basis:

$k = 0$; i.e. $T_1 = T_2$. Therefore, they have the same weight.

34

Proof of Lemma 1 (continued)

Inductive hypothesis: For $k > 0$, assume the lemma holds when there are j differing edges where $0 \leq j < k$.

Let uv be the minimum weight edge among the differing edges (assume uv is in T_2 but not T_1).


Look at unique path in T_1 from u to v .

Suppose it is made up of w_0, w_1, \dots, w_p where $w_0 = u, \dots, w_p = v$.

This path must contain some edge different from T_2 's.

Let $w_i w_{i+1}$ be this differing edge.

By MST property of T_1 , $w_i w_{i+1}$ cannot be $> uv$'s weight.



35

Proof of Lemma 1 (continued)

But since uv was chosen to be the minimum weight among differing edges, $w_i w_{i+1}$ cannot have weight less than uv .

Therefore, $W(w_i w_{i+1}) = W(uv)$.

Add uv to T_1 (creating a cycle). Remove $w_i w_{i+1}$ leaving tree T'_1 (which has the same weight as T_1).

But T'_1 and T_2 differ only on $k - 1$ edges.

So by inductive hypothesis, T'_1 and T_2 have the same total weight. Therefore, T_1 and T_2 have same weight.


36

Theorem 1 and Proof

Theorem 1: In a connected weighted graph $G = (V, E, W)$, a tree T is a minimum spanning tree if and only if T has the MST property.

Proof (Only if): Assume T is an MST for graph G .

Suppose T does not satisfy the MST property, i.e. there is some edge uv that is not in T such that adding uv creates a cycle, in which some other edge xy has weight $W(xy) > W(uv)$.



Then, by removing xy and adding uv , we create a new spanning tree whose total weight is $< W(T)$; This contradicts the assumption that T is an MST.

37

Proof of Theorem 1 (continued)

Theorem 1: In a connected weighted graph $G = (V, E, W)$, a tree T is a minimum spanning tree if and only if T has the MST property.

Proof (Only if): Assume T is an MST for graph G .

(Cont.)

(If) Assume T has MST property.

If T_{\min} is an MST, then T_{\min} has MST property by the first half of the proof.

By Lemma 1, $W(T) = W(T_{\min})$, so T is also an MST.

38

Prim's Algorithm is Optimal

Lemma 2: Let $G = (V, E, W)$ be a connected weighted graph; Let T_k be the tree with k vertices constructed by Prim's Algorithm, for $k = 1, 2, \dots, n$; and let G_k be the subgraph induced by the vertices of T_k . Then T_k has the MST property in G_k . (**Proof is not required**)

Theorem 2: Prim's Algorithm outputs a minimum spanning tree.

Proof:

- From Lemma 2, T_n has the MST property.
- By Theorem 1, T_n is a minimum spanning tree.

39

Priority Queue for MST (Optional)

- Inserted by order of priority (not chronological, as in 'normal' queues – FIFO)
- Elements to be inserted have a 'key' - contains the priority; element with highest priority will be selected first. [priority can be largest value (e.g. if we're computing max profit) or smallest value (e.g. if we're interested in min cost)]
- Think of pq as a sequence of pairs: $(id_1, w_1), (id_2, w_2), \dots, (id_k, w_k)$. The order is in increasing w_i and id is a unique identifier for an element

40

Methods of Priority Queue (Optional)

The Priority Queue consists of:

Create: Constructor to set up PQ

isEmpty; getMin; getPriority: Access functions

insert; deleteMin; decreaseKey: Manipulation procedures

Insert(pq, id, w): Inserts (id, w) into an existing pq - position depends on w

decreaseKey(pq, id, neww): Rearranges pq based on new wt of element id

getMin(pq): Returns id_1 ;

getPriority(pq): Returns weight of min element

Summary

- Greedy algorithm is a general strategy to solve optimization problems
- Dijkstra's algorithm finds single-source shortest paths in a weighted graph of nonnegative edge weights
- Prim's algorithm finds the minimum spanning trees in weighted graphs
- Both are greedy algorithms, and use priority queue