

## SC3000/CZ3005 Artificial Intelligence

### Intelligent Agents and Search

Prof Bo AN

[www.ntu.edu.sg/home/boan](http://www.ntu.edu.sg/home/boan)


Email: [boan@ntu.edu.sg](mailto:boan@ntu.edu.sg)

Office: N4-02b-55



## Lesson Outline

- How can one describe the task/problem for the agent?
- What are the properties of the task environment for the agent?
- Problem formulation
- Uninformed search strategies
- Informed search strategies: greedy search, A \* search




## Agent



An **agent** is an entity that

- **Perceives** through sensors (e.g. eyes, ears, cameras, infrared range sensors)
- **Acts** through effectors (e.g. hands, legs, motors)



## Rational Agents



- A rational agent is one that does the **right** thing
- **Rational action**: action that maximises the expected value of an objective **performance** measure given the percept sequence to date
- Rationality depends on
  - performance measure
  - everything that the agent has perceived so far
  - built-in knowledge about the environment
  - actions that can be performed



## Example: Google X2: Driverless Taxi

- **Percepts:** video, speed, acceleration, engine status, GPS, radar, ...
- **Actions:** steer, accelerate, brake, horn, display, ...
- **Goals:** safety, reach destination, maximise profits, obey laws, passenger comfort, ...
- **Environment:** Singapore urban streets, highways, traffic, pedestrians, weather, customers, ...



Image source: [https://en.wikipedia.org/wiki/Waymo#/media/File:Waymo\\_Chrysler\\_Pacifica\\_in\\_Los\\_Altos,\\_2017.jpg](https://en.wikipedia.org/wiki/Waymo#/media/File:Waymo_Chrysler_Pacifica_in_Los_Altos,_2017.jpg)



## More Examples

Agent Type	Percepts	Actions	Goals	Environment
Medical diagnosis system	Symptoms, findings, patient's answers	Questions, tests, treatments	Healthy patient, minimize costs	Patient, hospital
Satellite image analysis system	Pixels of varying intensity, color	Print a categorization of scene	Correct categorization	Images from orbiting satellite
Part-picking robot	Pixels of varying intensity	Pick up parts and sorts into bins	Place parts in correct bins	Conveyor belt with parts
Refinery controller	Temperature, pressure readings	Open, close valves; adjust temperature	Maximize purity, yield, safety	Refinery
Interactive English tutor	Typed words	Print exercises, suggestions, corrections	Maximize student's score on test	Set of students



## Types of Environment

only you (know the card) Go vs poker

Accessible (vs inaccessible)	Agent's sensory apparatus gives it access to the complete state of the environment
Deterministic (vs nondeterministic)	The next state of the environment is completely determined by the current state and the actions selected by the agent
Episodic (vs Sequential)	Each episode is not affected by the previous taken actions
Static (vs dynamic)	Environment does not change while an agent is deliberating
Discrete (vs continuous)	A limited number of distinct percepts and actions



## Example: Driverless Taxi

Accessible?	No. Some traffic information on road is missing
Deterministic?	No. Some cars in front may turn right suddenly
Episodic?	No. The current action is based on previous driving actions
Static?	No. When the taxi moves, Other cars are moving as well
Discrete?	No. Speed, Distance, Fuel consumption are in real domains



## Example: Chess

Accessible?	Yes. All positions in chessboard can be observed
Deterministic?	Yes. The outcome of each movement can be determined
Episodic?	No. The action depends on previous movements
Static?	Yes. When there is no clock, when are you considering the next step, the opponent can't move; <b>Semi</b> . When there is a clock, and time is up, you will give up the movement
Discrete?	Yes. All positions and movements are in discrete domains



## More Examples

Environment	Accessible	Deterministic	Episodic	Static	Discrete
Chess with a clock	Yes	Yes	No	Semi	Yes
Chess without a clock	Yes	Yes	No	Yes	Yes
Poker	No	No	No	Yes	Yes
Backgammon	Yes	No	No	Yes	Yes
Taxi driving	No	No	No	No	No
Medical diagnosis system	No	No	No	No	No
Image-analysis system	Yes	Yes	Yes	Semi	No
Part-picking robot	No	No	Yes	No	No
Refinery controller	No	No	No	No	No
Interactive English tutor	No	No	No	No	Yes



## Design of Problem-Solving Agent

### Idea

- Systematically considers the **expected outcomes** of different possible sequences of actions that lead to states of known value
- Choose the best one
  - shortest journey from A to B?
  - most cost effective journey from A to B?



## Design of Problem-Solving Agent


### Steps

- Goal formulation
- Problem formulation
- Search process
  - No knowledge → uninformed search
  - Knowledge → informed search
- Action execution (follow the recommended route)



## Example: Romania

- ❑ **Goal:** be in Bucharest
- ❑ **Formulate problem:**
  - **states:** various cities
  - **actions:** drive between cities
- ❑ **Solution:**
  - sequence of cities, e.g., Arad, Sibiu, Fagaras, Bucharest



## Example: Vacuum Cleaner Agent




- ❑ Robotic vacuum cleaners move autonomously
- ❑ Some can come back to a docking station to charge their batteries
- ❑ A few are able to empty their dust containers into the dock as well



## Example: A Simple Vacuum World

- ❑ Two locations, each location may or may not contain dirt, and the agent may be in one location or the other




- ❑ 8 possible world states
- ❑ Possible actions: **left**, **right**, and **suck**
- ❑ Goal: clean up all dirt → Two goal states, i.e. {7, 8}



## Single-State Problem

- **Accessible** world state (sensory information is available)
- **Known** outcome of action (deterministic)




- e.g.: start in #5
- Solution: **right, suck** 5-6-8



## Multiple-State Problem

- **Inaccessible** world state (with limited sensory information):
  - agent only knows which sets of states it is in
- **Known** outcome of action (deterministic)



- e.g.: start in {1, 2, 3, 4, 5, 6, 7, 8}
  - Action **right** goes to {2, 4, 6, 8}
  - Solution: **right, suck, left, suck**



## Well-Defined Formulation

<b>Definition of a problem</b>	The information used by an agent to decide what to do
<b>Specification</b>	<ul style="list-style-type: none"> <li>• Initial state</li> <li>• Action set, i.e. available actions (successor functions)</li> <li>• State space, i.e. states reachable from the initial state               <ul style="list-style-type: none"> <li>• Solution path: sequence of actions from one state to another</li> </ul> </li> <li>• Goal test predicate               <ul style="list-style-type: none"> <li>• Single state, enumerated list of states, abstract properties</li> </ul> </li> <li>• Cost function               <ul style="list-style-type: none"> <li>• Path cost <math>g(n)</math>, sum of all (action) step costs along the path</li> </ul> </li> </ul>
<b>Solution</b>	A path (a sequence of operators leading) from the Initial-State to a state that satisfies the Goal-Test



## Measuring Problem-Solving Performance

### Search Cost

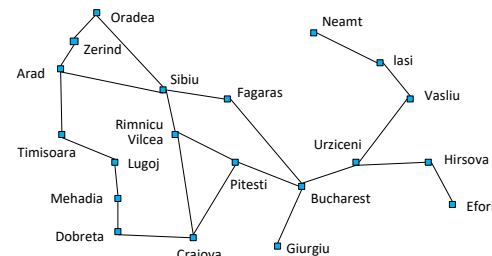
- What does it cost to find the solution?
  - e.g. How long (time)? How many resources used (memory)?

### Total cost of problem-solving

- **Search cost** ("offline") + **Execution cost** ("online")
- Trade-offs often required
  - Search a **very long time** for the **optimal solution**, or
  - Search a **shorter time** for a **"good enough"** solution



## Single-State Problem Example




- **Initial state:** e.g., "at Arad"
- Set of possible **actions** and the corresponding next states
  - e.g., Arad → Zerind
- **Goal test:**
  - explicit (e.g.,  $x$  = "at Bucharest")
- **Path cost** function
  - e.g., sum of distances, number of operators executed solution: a sequence of operators leading from the initial state to a goal state

## Example: Vacuum World (Single-state Version)



- **Initial state:** one of the eight states shown previously
- **Actions:** left, right, suck
- **Goal test:** no dirt in any square
- **Path cost:** 1 per action



## Multiple-State Problem Formulation



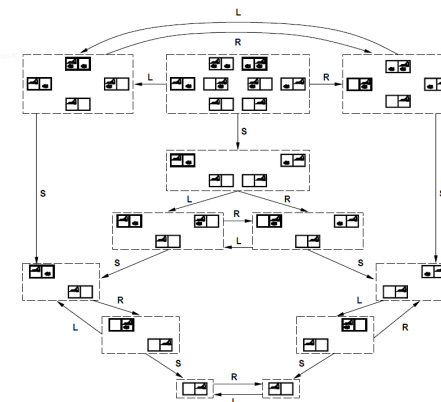
- Initial state **set**
  - Set of possible actions and the corresponding **sets of** next states
  - Goal test
  - Path cost function
- **Solution:**
    - a path (connecting sets of states) that leads to a set of states all of which are goal states

## Example: Vacuum World (Multiple-state Version)



- **States:** subset of the eight states
- **Operators:** left, right, suck
- **Goal test:** all states in state set have no dirt
- **Path cost:** 1 per operator

## Example: Vacuum World (Multiple-state Version)





## Example: 8-puzzle

- **States:** integer locations of tiles
  - number of states =  $9!$
- **Actions:** move blank left, right, up, down
- **Goal test:** = goal state (given)
- **Path cost:** 1 per move

Start state	5	4	
	6	1	8
	7	3	2

Goal state	1	2	3
	8		4
	7	6	5



## Real-World Problems

Route finding problems:

- Routing in computer networks
- Robot navigation
- Automated travel advisory
- Airline travel planning




Touring problems:

- Traveling Salesperson problem
- "Shortest tour": visit every city exactly once



## Question to Think About

The stone puzzle is characterised as follows. Two white stones and two black stones are initially positioned on a board, as illustrated below. The board is composed of a single row of five squares, each of which can be either empty or hold one stone. The white stones can only move to the right and the black stones only to the left, either into an immediately adjacent empty square or by jumping over an adjacent stone of opposite colour into an empty square.



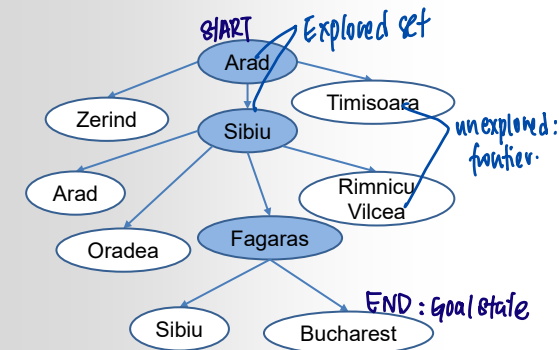
The problem is to exchange the respective positions of the white and black stones, moving them one at a time according to the rules above. Answer the following to solve this puzzle using a problem-solving approach:

- Give a *well-defined formulation* of the problem in terms of states, operators, goal test predicate, and path cost.



## Search Algorithms

- Exploration of state space by generating successors of already-explored states
  - **Frontier:** candidate nodes for expansion
  - Explored set





# Search Strategies

- A **strategy** is defined by picking the order of node expansion.
- Strategies are evaluated along the following dimensions:

Completeness	Does it always find a solution if one exists?
Time Complexity	How long does it take to find a solution: the number of nodes generated / explored
Space Complexity	Maximum number of nodes in memory unexplored (i.e. white nodes)
Optimality	Does it always find the best (least-cost) solution?

leads to  
rls b/w optimality & completeness  
may not exist

store frontier n  
necessarily explored.

# Uninformed vs Informed

## Uninformed search strategies

- Use **only** the information available in the problem definition
- Breadth-first search
  - Uniform-cost search
  - Depth-first search
  - Depth-limited search
  - Iterative deepening search

## Informed search strategies

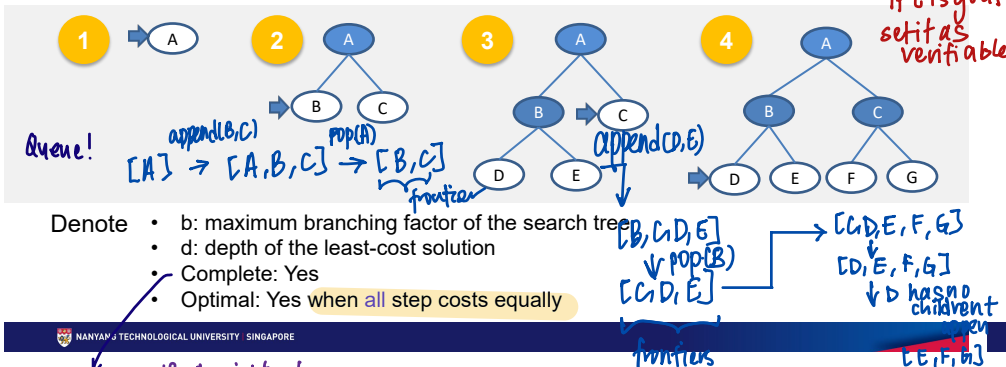
- Use **problem-specific knowledge** to guide the search (inductive bias)
- Usually more efficient

knows knowledge about the task  
To find min  $x^2$  btw  $[-1, 1]$   
uninformed:  $[-1, 1] \forall \epsilon \in \mathbb{R}$   
agent searches every value in the search space

informed: use derivative -  
 $\frac{d}{dx} x^2 = 2x$   
to find min  $2x = 0$   
Agent knows about derivatives

# Breadth-First Search

Expand **shallowest** unexpanded node which can be implemented by a **First-In-First-Out (FIFO)** queue



e.g. if C is goal state  $\Rightarrow$  Algo will eventually reach C eventually find  
if all weights = 1  
BFS = shortest path finding algo.

# Complexity of BFS

- Hypothetical state-space, where every node can be expanded into  $b$  new nodes, solution of path-length  $d$
- Time:  $1 + b + b^2 + b^3 + \dots + b^d = O(b^{d+1})$
- Space: (keeps every node in memory)  $O(b^d)$

Always expands to the branching factor  
Geometric series  
how many layers  
branching factor  
means store every node until last layer.

Depth	Nodes	Time	Memory
0	1	1 millisecond	100 bytes
2	111	0.1 seconds	1 kilobytes
11	11111	11 seconds	11 kilobytes
18	10 <sup>6</sup>	18 minutes	111 megabyte
8	10 <sup>8</sup>	31 hours	11 gigabytes
10	10 <sup>10</sup>	128 days	1 terabyte
12	10 <sup>12</sup>	35 years	111 terabytes
14	10 <sup>14</sup>	3500 years	11111 terabytes

depth # Blue nodes explored total # nodes explored

1 0 6  
2 1 1  
3 2 3  
4 3 7


if E is goal  
set it as verifiable  
[A]  $\rightarrow$  [A,B,C]  $\rightarrow$  [B,C]  
[B,C,D,E]  $\rightarrow$  [C,D,E,F,G]  
[C,D,E,F,G]  $\rightarrow$  [D,E,F,G]  
[D,E,F,G]  $\rightarrow$  [E,F,G]  
[E,F,G]  $\rightarrow$  [F,G]  
[F,G]  $\rightarrow$  [G]



## Uniform-Cost Search

To consider **edge costs**, expand unexpanded node with the **least** path cost  $g$


- Modification of breath-first search
- Instead of First-In-First-Out (FIFO) queue, using a priority queue with path cost  $g(n)$  to order the elements
- BFS = UCS with  $g(n) = \text{Depth}(n)$



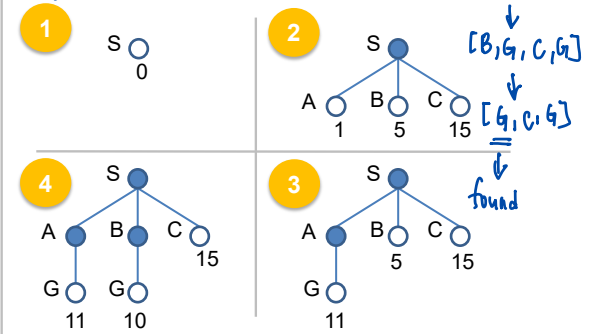
## Uniform-Cost Search

uses priority queue

$$\begin{aligned} S \rightarrow A &= 0 + 1 = 1 \\ S \rightarrow B &= 0 + 5 = 5 \\ S \rightarrow C &= 0 + 15 = 15 \end{aligned}$$



$[S] \rightarrow [S, A, B, C] \rightarrow [A, B, C] \rightarrow [A, B, C, G] \rightarrow [B, G, C]$




## Uniform-Cost Search

Complete	Yes
Time	# of nodes with path cost $g \leq \text{cost of optimal solution}$ (eqv. # of nodes pop out from the priority queue)
Space	# of nodes with path cost $g \leq \text{cost of optimal solution}$
Optimal	Yes

in the worst case still have to explore all the nodes

## Depth-First Search

Expand deepest unexpanded node which can be implemented by a Last-In-First-Out (LIFO) stack. Backtrack only when no more expansion



$[G, M, L] \rightarrow$



## Depth-First Search

Denote

- $m$ : maximum depth of the state space

Complete	<ul style="list-style-type: none"> <li>infinite-depth spaces: No</li> <li>finite-depth spaces with loops: No                     <ul style="list-style-type: none"> <li>with repeated-state checking: Yes</li> </ul> </li> <li>finite-depth spaces without loops: Yes</li> </ul>
Time	$O(b^m)$ If solutions are dense, may be much faster than breadth-first
Space	$O(bm)$
Optimal	No

only done  
first part

→ only if the cost is 1 and the answer is at the bottom DFS then only DFS is optimal

real world problems



## Depth-Limited Search

To avoid infinite searching, Depth-first search with a **cutoff** on the max depth  $l$  of a path

Complete	Yes, if $l \geq d$
Time	$O(b^l)$
Space	$O(bl)$
Optimal	No



## Iterative Deepening Search

Iteratively estimate the max depth  $l$  of DLS one-by-one

Limit = 0




Limit = 1



## Iterative Deepening Search

Iteratively estimate the max depth  $l$  of DLS one-by-one

Limit = 2



At every depth limit you apply DFS

When this iteration is finished ⇒ color blue ⇒ flushed out of the memory.

⇒ At any point DLS ⇒ bxd.

many completions of DFS


optimal space comp of DFS



## Iterative Deepening Search

Iteratively estimate the max depth / of DLS one-by-one

Limit = 3



## Iterative Deepening Search...

```

Function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution sequence
  inputs: problem, a problem

  for depth 0 to ∞ do
    if DEPTH-LIMITED-SEARCH(problem, depth) succeeds then return its result
  end
  return failure
    
```

Complete Yes

Time  $O(b^d)$

Space  $O(bd)$

Optimal Yes



## Summary (we make assumptions for optimality)

Criterion	Breadth-first	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Time	$b^d$	$b^d$	$b^m$	$b^l$	$b^d$	$b^{d/2}$
Space	$b^d$	$b^d$	$bm$	$bl$	$bd$	$b^{d/2}$
Optimal	Yes	Yes	No	No	Yes	Yes
Complete	Yes	Yes	No	Yes, if $l \geq d$	Yes	Yes



## General Search

Uninformed search strategies

- **Systematic** generation of new states ( $\rightarrow$  Goal Test)
- **Inefficient** (exponential space and time complexity)

Informed search strategies

- Use **problem-specific** knowledge
  - To decide the order of node expansion
- Best First Search: expand the most desirable unexpanded node
  - Use an **evaluation function** to estimate the "desirability" of each node



## Evaluation function

- Path-cost function  $g(n)$ 
  - Cost from initial state to current state (search-node)  $n$
  - No information on the cost **toward the goal**
- Need to estimate cost to the closest goal
- "Heuristic" function  $h(n)$  *ultimate goal fct  $\Rightarrow$  estimate a good heuristic func*
  - Estimated cost of the cheapest path from  $n$  to a goal state  $h(n)$  *"base of domain know"*
    - Exact cost cannot be determined
  - depends only on the state at that node
  - $h(n)$  is not larger than the real cost (admissible)



## Greedy Search

- Expands the node that **appears** to be closest to goal
- Evaluation function  $h(n)$ : estimate of cost from  $n$  to **goal**
  - Function Greedy-Search(problem) returns solution
    - Return Best-First-Search(problem,  $h$ ) //  $h(goal) = 0$


**Question:** How to estimation the cost from  $n$  to goal?

**Answer:** Recall that we want to use problem-specific knowledge



## Example: Route-finding from Arad to Bucharest

$h(n)$  = straight-line distance from  $n$  to Bucharest



Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

- Useful but potentially fallible (**heuristic**)
- Heuristic functions are problem-specific



## Example

- a) The initial state




Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

## Example

b) After expanding Arad

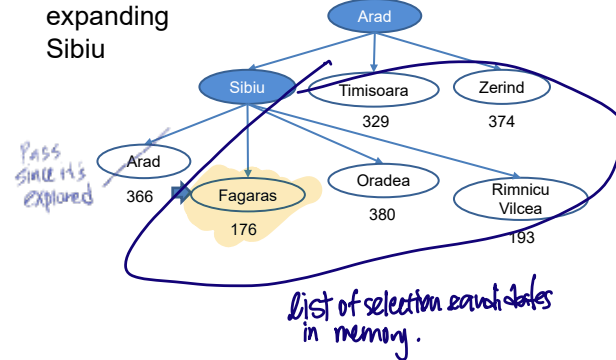


Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

## Example

c) After expanding Sibiu

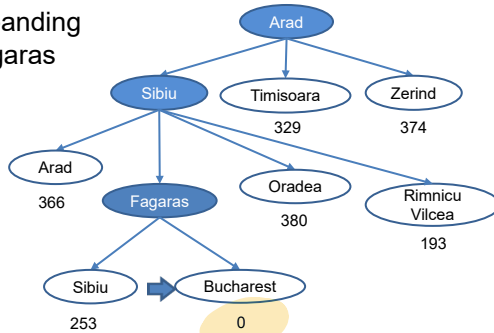


Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

## Example

d) After expanding Fagaras



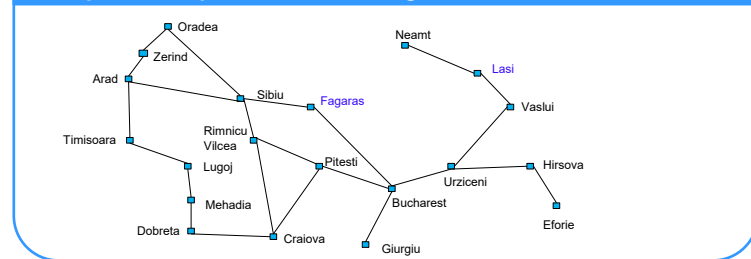
Straight-line distance to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Efoire	161
Fagaras	176
Giurgiu	77
Hirsova	151
Lasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

## Complete?

Question: Is this approach complete?

Example: Find a path from Lasi to Fagaras



Answer: No

In this vanilla version of greedy.  
If you reach a node and the node has no children  $\Rightarrow$  you will consider

## Greedy Search...

- $m$ : maximum depth of the search space

Complete	No
Time	$O(b^m)$
Space	$O(b^m)$ (keeps all nodes in memory)
Optimal	No

## A \* Search

- Uniform-cost search
  - $g(n)$ : cost to reach  $n$  (Past Experience)
  - optimal and complete, but can be very inefficient
- Greedy search
  - $h(n)$ : cost from  $n$  to goal (Future Prediction)
  - neither optimal nor complete, but cuts search space considerably

*combines*

## A \* Search

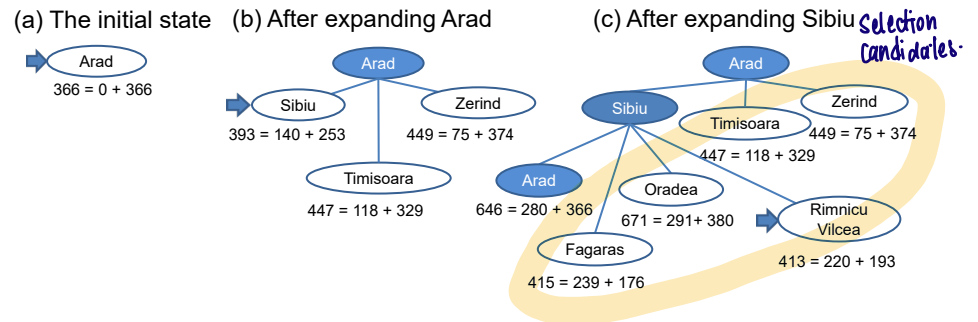
Idea: Combine Greedy search with Uniform-Cost search

Evaluation function:  $f(n) = g(n) + h(n)$

- $f(n)$ : estimated total cost of path through  $n$  to goal (Whole Life)
- If  $g = 0 \rightarrow$  greedy search; If  $h = 0 \rightarrow$  uniform-cost search
- Function  $A^* Search(problem)$  returns solution
  - Return  $Best\text{-}First\text{-}Search(problem, g + h)$

## Example: Route-finding from Arad to Bucharest

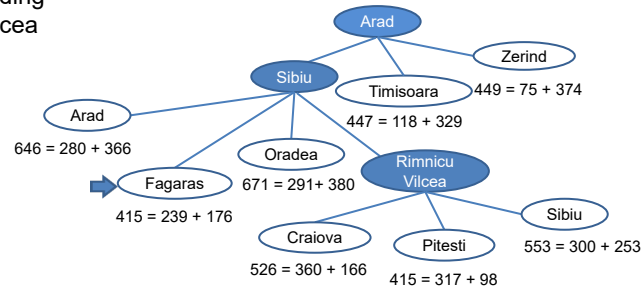
Best-first-search with evaluation function  $g + h$



## Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function  $g + h$

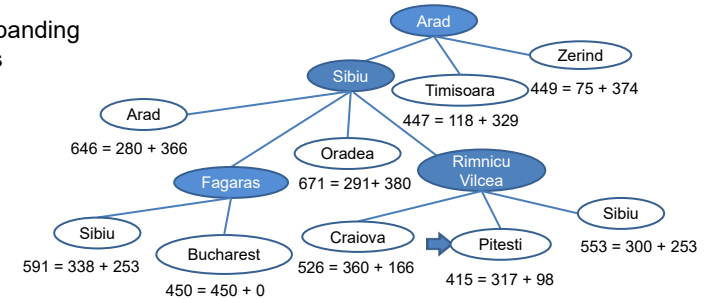
(d) After expanding Rimnicu Vilcea



## Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function  $g + h$

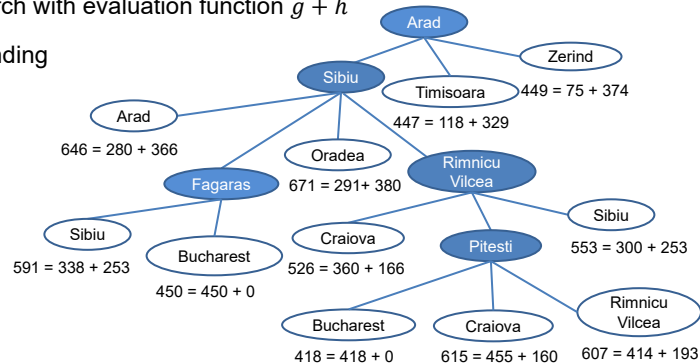
(e) After expanding Fagaras




## Example: Route-finding from Arad to Bucharest

Best-first-search with evaluation function  $g + h$

(f) After expanding Pitesti




## Example: Route-finding in Manhattan




Manhattan  
Distance  
Heuristic




## Example: Route-finding in Manhattan (Greedy)




## Example: Route-finding in Manhattan (UCS)



## Example: Route-finding in Manhattan (A\*)



## Complexity of A\*



Time	Exponential in length of solution
Space	(all generated nodes are kept in memory) Exponential in length of solution

With a good heuristic, significant savings are still possible compared to uninformed search methods