

CZ3005 Artificial Intelligence

Week 8a – Logical Agent

Yu Han

han.yu@ntu.edu.sg

Nanyang Assistant Professor
School of Computer Science and Engineering
Nanyang Technological University



Second Half Lecturer



Nanyang Assistant Professor Yu Han

- Email: han.yu@ntu.edu.sg
- Office: N4-02c-109
- Website: <http://hanyu.sg/>

2

Course Overview (2nd Half)

Week	Date	Venue: LT1A	Topic	Tutorial	Lab
8	09/10/2023	1	Logical Agent (a)		
8	10/10/2023	2	Logical Agent (b)		
9	16/10/2023	3	Propositional Logic (a)		
9	17/10/2023	4	Propositional Logic (b)		
10	23/10/2023	5	First-Order Logic (a)	Tutorial 4	One Lab session in Week 10/11 (Lab Attendance is <u>NOT</u> <u>Compulsory</u>)
10	24/10/2023	6	First-Order Logic (b)		
11	30/10/2023	7	Default Logic (a)		
11	31/10/2023	8	Default Logic (b)		
12	06/11/2023	9	Fuzzy Logic (a)	Tutorial 5	
12	07/11/2023	10	Fuzzy Logic (b)		
13	No Class				

3

Why Logic?

- The earliest way people imagined how to build AI system
- Still a relevant tool to guide today's AI with human knowledge
- Good interpretability (which deep learning-based AI lacks)

4

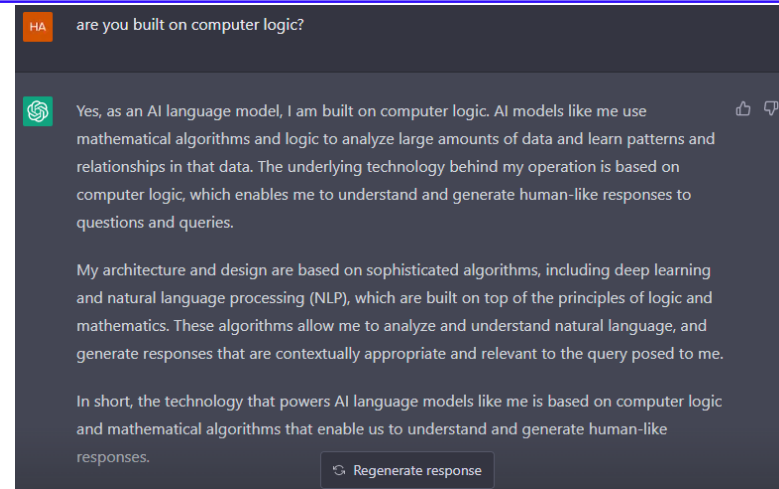
An Example Logical Agent

https://www.youtube.com/watch?v=Mxl_tzOs6ME



5

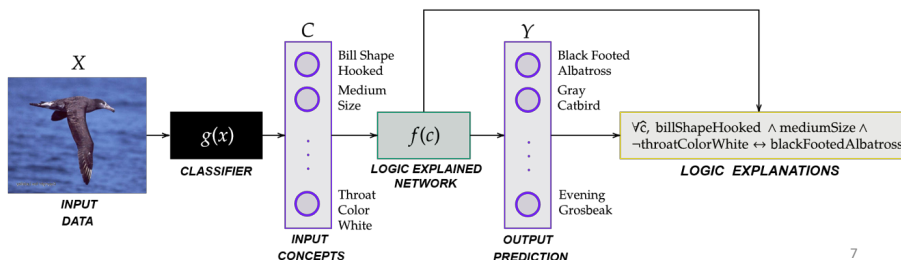
Why Logic?



6

Why Logic?

- New life for an old technology:
 - Deep Learning and Logic Programming are starting to merge today for building Explainable AI (XAI) systems.
 - Logic Explained Networks:




7

From Search to Logical Reasoning

- Representation of Knowledge and the Reasoning Processes are central to the entire field of AI
- A knowledge-based system (KBS) is a computer program that reasons and uses a knowledge base to solve complex problems.

8

Knowledge and Intelligence



9

The Knowledge-Based Approach

- Agents that **know**
 - Achieve competence by being told new knowledge or by learning
 - Achieve adaptability by updating their knowledge
 - > **Knowledge representation**
 - State of the world, properties and evolution of the world; **goals** of the agent, actions and their effect
- Agents that **reason**
 - Use knowledge to deduce course of actions
 - > **Knowledge inference**

Logic

10

Knowledge-Based Agents

- **Knowledge base (KB)**
 - Set of **sentences** i.e., representations of facts (DB)
 - Knowledge representation language
- Adding and querying knowledge
 - **Tell**: add a sentence to the KB
 - **Ask**: retrieve knowledge from the KB
 - Answers **must follow** from what has been **Tell**'ed (told)
- Inference mechanism
 - Determine what follows from the KB

11

Problem Formulation of KBS

- Knowledge Based System
 - **States**: Instances of the KB (**sets of sentences**)
 - Use **Tell** to build the KB
 - e.g. Tell(KB, "Smoke \Rightarrow Fire")
 - Tell(KB, "Fire \Rightarrow Call_999")
 - ...
 - Tell(KB, "Smoke") → *there's smoke*
 - **Operators**: Add / Infer a new sentence
 - **Goal**: Answer a query
 - Use **Ask** to query the KB
 - e.g. Ask(KB, "? Call_999")

12

A Generic Knowledge-Based Agent

```

function KB-Agent (percept) returns action
  static KB,           // a knowledge base
         t             // a time counter, initially 0

  Tell (KB, Make-Percept-Sentence (percept, t))
  action ← Ask (KB, Make-Action-Query (percept, t))
  Tell (KB, Make-Action-Sentence (action, t))
  t ← t + 1
  return action
    
```

- 3 steps: interpretation, inference, execution
- KB: background knowledge (observed)
+ acquired information (deduced)

Human
context is obvious to
human.

Logical
encoding of
knowledge

13

Levels of Knowledge

• Epistemological level


- Declarative description of knowledge
 - e.g. facts: “there is smoke in the kitchen”, “it is not warm enough”
 - rules: “if there is smoke then there must be a fire”

• Logical level

- Logical encoding of knowledge (into sentences)
 - e.g. facts: **Smoke**; rules: **Implies(Smoke, Fire)**


• Implementation level

- Physical representation of knowledge (sentences)
 - e.g. - the string “Implies(Smoke, Fire)”, or
 - a “1” entry in a 2-dimensional array: **Implies[X,Y]**



14

Example: the Wumpus World



RL: don't know rules & explore

Rules are obvious, logic is used to solve this problem.

15

Example: the Wumpus World

• Problem description (PAGE)

– Environment

- Grid of squares, walls;
- Agent, gold, pits, wumpus.

– Goal


- Find the gold, return to S at [1,1].

– Percepts

- A list of 5 symbols, e.g. [Stench, Breeze, Glitter, Bump, Scream];
- Agent's location *not* perceived.

– Actions

- Go-Forward, Turn-Left, Turn-Right, Grab, Shoot (1 arrow only), Climb.



16

The Wumpus World

Acting and Reasoning in the Wumpus World

• Problem description (cont'd)

– Initial state

- Agent at [1,1]; gold, pits and wumpus in random squares.

– Path-cost

- Climbing out with the gold: +1000 (without: 0)
- Each action: -1
- Getting killed (pit or wumpus): -10000


– Knowledge

- "In all squares adjacent to the one where the wumpus is, the agent will perceive a stench."
- "In all squares adjacent to a pit, the agent will perceive a breeze."
- In the square where the gold is, the agent will perceive a glitter."
- When walking into a wall, the agent will perceive a bump."
- When the wumpus is killed, the agent will perceive a scream."

17

(0) Initial state

[nil, nil, nil, nil, nil]




A = Agent G = Glitter, Gold
B = Breeze OK = Safe square

Initial state:
nothing sensed.

(1) after {F}

[nil, Breeze, nil, nil, nil]



RL will tell us to take leap of faith but here we BACKTRACK feed a breeze

P = Pit
S = Stench V = Visited
W = Wumpus


18

Acting and Reasoning in the Wumpus World

Acting and Reasoning in the Wumpus World

(1) after {F}


[nil, Breeze, nil, nil, nil]



A = Agent G = Glitter, Gold
B = Breeze OK = Safe square

(6) after {F, L, L, F, R, F}

[Stench, nil, nil, nil, nil]




P = Pit
S = Stench V = Visited
W = Wumpus

19

(6) after {F, L, L, F, R, F}


[Stench, nil, nil, nil, nil]



A = Agent G = Glitter, Gold
B = Breeze OK = Safe square

(10) after {F, L, L, F, R, F, R, F, L, F}

[Stench, Breeze, Glitter, nil, nil]



P = Pit
S = Stench V = Visited
W = Wumpus

20

Summary

- **Intelligent agents need ...**
 - Knowledge about the world to make good decisions.
- **Knowledge can be ...**
 - Defined using a knowledge representation language.
 - Stored in a knowledge base in the form of sentences.
 - Inferred, using an inference mechanism and rules.

21

Thank you!



22

CZ3005 Artificial Intelligence

Week 8b – Logical Agent

Yu Han

han.yu@ntu.edu.sg

Nanyang Assistant Professor
School of Computer Science and Engineering
Nanyang Technological University



Recap

```
function KB-Agent (percept) returns action
  static KB,           // a knowledge base
         t             // a time counter, initially 0

  Tell (KB, Make-Percept-Sentence (percept, t))
  action ← Ask (KB, Make-Action-Query (percept, t))
  Tell (KB, Make-Action-Sentence (action, t))
  t ← t + 1
  return action
```

- > 3 steps: interpretation, inference, execution
- > KB: background knowledge (observed)
+ acquired information (deduced)

Recap

- **Intelligent agents need ...**
 - Knowledge about the world to make good decisions.
- **Knowledge can be ...**
 - Defined using a knowledge representation language.
 - Stored in a knowledge base in the form of sentences.
 - Inferred, using an inference mechanism and rules.


Knowledge Representations

- **Knowledge representation (KR)**
 - KB: set of sentences → need to
 - Express knowledge in a (computer-) tractable form
- **Knowledge representation language**
 - **Syntax – implementation level**
 - Possible configurations that constitute sentences
 - **Semantics – knowledge level**
 - Facts of the world the sentences refer to
 - e.g. language of arithmetics: x, y numbers
sentence: " $x \geq y$ ", semantics: "greater or equal"

Logic

Reasoning and Logic

- **Logic**
 - Representation + Inference = Logic
 - Where representation = syntax + semantics
- **Reasoning**
 - Construction of new sentences from existing ones
- **Entailment** as logical inference
 - the relationship between sentences whereby one sentence will be true if all the others are also true



Deduction and Induction




try not to use because it generalises.

Deduction and Induction

- **Mechanical reasoning**
 - Example
 - If a chord sequence is tonal, then it can be generated by a context-sensitive grammar.
 - The twelve-bar blues has a chord sequence that is tonal. \vdash
 - The twelve-bar blues has a chord sequence that can be generated by a context-sensitive grammar.
- **Deductive inference**
 - KB: Monday \Rightarrow Work, Monday \vdash Work *sound*
- **Inductive inference**
 - KB: Monday \Rightarrow Work, Work \vdash Monday *unsound!*
 - Generalization e.g., "all swans are white ..."

Entailment and Inference

- **Entailment**
 - Generate sentences that are necessarily true, given that the existing sentences are true
 - Notation: $KB \models \alpha$
 - e.g. Wumpus world:
 - $\{ \neg S(1,1), \neg B(1,1) \} \models \text{"OK}(2,1)"$
 - Arithmetics: $\{ "x \geq y", "y \geq z" \} \models "x \geq z"$
- **Inference**
 - The act or process of deriving logical conclusions from premises known or assumed to be true.
 - **Tell**, given KB: $(KB \models \alpha) !$
 - **Ask**, given KB and α : $(KB \models \alpha) ?$



Knowledge Representation Languages

- **Formal (programming) languages**

- Good at describing algorithms and data structures
 - e.g. the Wumpus world as a 4x4 array, $\text{World}[2,2] \leftarrow \text{Pit}$
- Poor at representing incomplete / uncertain information
 - e.g. “there is a pit in [2,2] or [3,1]”, or “...a wumpus somewhere”
- > *not expressive enough*

- **Natural languages**

- Very expressive (too much, thus very complex)
- More appropriate for communication than representation
- Suffer from ambiguity
 - e.g. “It’s hot!”
 - e.g. “small cats and dogs” compared to “ $-x + y$ ”.

9

Properties of Semantics

- **Interpretation (meaning)**

- *Correspondence between sentences and facts*
- **Arbitrary meaning**, fixed by the writer of the sentence
 - e.g. Natural languages: meaning fixed by usage (cf. dictionary)
exceptions: encrypted messages, codes (e.g. Morse)
- **Systematic relationship**: compositional languages
 - *The meaning of a sentence is a function of the meaning of its parts.*
- **Truth value**
 - A sentence make a claim about the world \rightarrow TRUE or FALSE
 - *Depends on the interpretation and the state of the world*
 - e.g. Wumpus world: $S(1,2)$ true if means “Stench at [1,2]” and the world has a wumpus at either [1,3] or [2,2].

11

Properties of Representations

- *KR languages should combine the advantages of both programming and natural languages.*

- **Desired properties**

- Expressive
 - Can represent everything we need to.
- Concise
- Unambiguous
 - Sentences have a unique interpretation.
- Context independent
 - Interpretation of sentences depends on semantics only.
- Effective
 - An inference procedure allows to create new sentences.

10

Properties of Inference

- **Definition**

- *Inference (reasoning) is the process by which conclusions are reached*
- Logical inference (deduction) is the process that implements entailment between sentences

- **Useful properties**

- **Valid sentence (tautology)**
 - iff TRUE under all possible interpretations in all possible worlds.
 - e.g. “ S or $\neg S$ ” is valid, “ $S(2,1)$ or $\neg S(2,1)$ ”, etc.
- **Satisfiable sentence**
 - iff there is some interpretation in some world for which it is TRUE
 - e.g. “ S and $\neg S$ ” is unsatisfiable

12

Properties of Inference

• Soundness

- A deductive system is considered sound if it never proves a false statement; it only derives conclusions that are true whenever the premises are true.

- Proof: sequence of operations of a sound inference

- Record of operations that generate a specific entailed sentence
e.g. "Smoke \Rightarrow Fire" and "Smoke" \models "Fire"
"Fire \Rightarrow Call_911" and "Fire" \models "Call_911"



• Completeness

Satisfiable rules but we don't know conditions are fulfilled and \Rightarrow Require some facts for conc

- A deductive system is complete if it can prove all true statements.
- Any statement that is true within the system can be formally derived or proven using the rules of the system.

13

Example: Sound & Complete



Premise 1: It is raining outside. (A)

Premise 2: If it is raining, then the ground is wet. (A implies B)

Using Modus Ponens, we can infer that the ground is wet. (B)

In this case, the argument is both sound and complete.

- It is sound as the conclusion (the ground is wet) logically follows from the premises.
- It is complete as all statements derived from these premises can be proven using the given rule (Modus Ponens).

14

Example: Sound, but NOT Complete





Premise 1: If it is sunny, then people are happy. (A implies B)

Not complete: we don't know right now if it is sunny or not. Doesn't tell us if part satisfied or

- In this case, the argument is sound because the conclusion follows logically from the premise.
- However, this deductive system is not complete because it cannot prove statements about the weather condition taking place.

15

Example: Complete, but NOT Sound



Premise 1: It is raining outside. (A)

Conclusion: The moon is made of cheese.

we dk where this conc comes from doesn't follow any deduction.

- The argument is complete because it tries to make a statement about the moon, which is a valid topic within the system.
- However, it is not sound because the conclusion does not logically follow from the premise;
 - raining outside does not imply anything about the composition of the moon.

16

Soundness & Completeness: Refers to properties of entire KB

Inference and Agent Programs

- **Inference in computers**

- Does not know the interpretation the agent is using for the sentences in the KB
- Does not know about the world (actual facts)
- Knows only what appears in the KB (sentences)
 - e.g. Wumpus world: doesn't know the meaning of "OK", what a wumpus or a pit is, etc. – can only see: $KB \models "[2,2] \text{ is OK}"$
- > *Cannot reason informally*
 - does not matter, however, if $KB \models "[2,2] \text{ is OK}"$ is a valid sentence

- **Formal inference**

- Can handle arbitrarily complex sentences, $KB \models P$

17

Thank you!



18