

Module 2: Intelligent Agent & Search

1. Types of Environment

Accessible? : Have access to complete state of environment

Deterministic? : Next state of environment completely determined by current stated actions selected by the agent

Episodic? : Each episode is not affected by previous actions taken

Sequential?

Static? Dynamic? : Env^b & when agent is deliberating?

discrete? continuous: Limited # of distinct percepts & actions

→ 1. Goal Form^a 4. Action Execution.

Design of Problem-Solving Agent → 2. Problem Form^a 3. Search Process

uninformed
informed

2. What is a well defined problem formulation.

Def of problem : info used by agent to decide what to do.

Specification: Initial state.

→ Action state (i.e. avail actions (successors)) → Inst Func^a: path cost, goal, 2nd action step cost along path

→ State space (i.e. states reachable from initial state)

· sol^a: seq. of action from one state to another

→ Goal state predicate

· single state enumerated list of states, abstract properties.

Solution: path (seq. of operators from initial state to state) Frontier (# candidates that satisfies goal-test).

nodes for expansion)

2. Search strategies

Completeness: does it always find sol^a if one exists

Time Complexity: How long? # of nodes expanded (Explored Nodes)

Space Complexity: Max # nodes in mem. (unexplored Nodes)

Optimality: Does it always find the best (least-cost) sol^a?

b: max branching factor of search tree
d: depth of least-cost sol^a

A) Uninformed Search

	What?	Time	Space	Complete	Optimal
BFS	Expand shallowest unexpanded node (FIFO)	$O(b^d)$	$O(b^d)$	Yes	Yes, when cost are =
UCS	FIFO w/ priority queue to path cost (min) to order elements	# nodes w/ path cost \geq cost of optimal sol ^a	=	Yes	Yes
DFS	Expand deepest unexpanded node n LIFO stack + back-track	$O(b^m)$. If sol ^a dense if sol ^a sparse, faster than BFS	$O(b^m)$	Infinite depth: No finite depth with loops: No · with repeated state check: Yes finite w/o loops	No
Depth Limited	DFS with cutoff in max depth of a path	$O(b^l)$	$O(b^l)$	Yes, if $l > d$	No

· Systematic generation of new states (\rightarrow goal-test)

· Inefficient (Exponential space & time complexity)

y

B) Informed Search Strategies

- USE problem specific Knowledge
 - To decide the order of node expansion
- Best First Search: expand the most desirable unexpanded node.
 - use an eval^{fun} to estimate desirability of each node.

Types of eval Funⁿ

1. Path cost funⁿ: $g(n)$
 - cost from initial state to search node (n)
 - No info on cost towards goal
 - Need to estimate cost toward goal

2. Heuristic $h(n)$
 - Estimated cost of cheapest path from n to goal state $h(n)$
 - Exact cost cannot be determined
 - depends only on state at that node $\cdot h(n) < \text{Real cost}$

UCS uses:

$g(n)$: cost to reach n (Past Exp)
Efficient but optimal & complete

→ cost from n to goal
(Future pred^e)

Greedy:

Expand node closest to goal
 $h(n)$ at eval func: E.g straight
line dist from n to Bucharest

A*

combines Greedy & ucs:
 $\text{Eval. fcn} = g(n) + h(n)$
Estimated total cost of path through n to goal

what?	Time	Space	complete	optimal
Expand node closest to goal $h(n)$ at eval func: E.g straight line dist from n to Bucharest	$O(b^n)$	$O(b^m)$	No	But ↓ search space ↔ No
combines Greedy & ucs: $\text{Eval. fcn} = g(n) + h(n)$ Estimated total cost of path through n to goal	Exponential in len of sol ⁿ	All generated nodes in mem Exponential in len sol ^m		

Module 3: Constraint Satisfaction & Adversarial Search

Constraint Satisfaction Problem

State: defined by variables, V_i with values from Domain, D_i

constraints - Goal test: a set of constraints specifying allowable combinations of values for subsets of variables.

Goal: discover some state that satisfies a given set of constraints.

Consistent / legal Assignment: assignment that don't violate any constraints

Solution: Assignment with every variable given a value / complete assignment satisfies all constraints.

Initial state:

* use search strategies to find solⁿ to CSP

→ Each state is a node in the tree

→ Solutions are at bottom of tree, use DFS

Possible heuristics for CSP

1. Order to assign variables

2. Order of values to try for each variable

3. Implications of current variable assignment to other unassigned variables

• (i) Forward checking: propagates info from unassigned to assigned values, keeps track of remaining legal values for unassigned var & terminates search when any var has no legal values → prevents unnecessary expansion of nodes when the node has no more possible values ↑ speed of search.

(ii) constraint propⁿ: propⁿ of implications of a constraint on one var onto other var, whereby nodes are expanded if they don't violate the constraint

WAYS to solve CSP Problem

1. Standard search - uninformed Algo. e.g. backtracking.

→ IDEA: do not waste time searching when constraints have already been violated.

· Goes back to previous node if current node violates some constraint.

2. Most constraining variable:

· ↓ branching factor of future choices by selecting var that is involved in the largest # of constraints on unassigned var

· ↓ domain of many other unassigned var.

3. Least constraining variable

· choose the value that leaves max flexibility for subsequent var assignment

4. Min Conflict Heuristics

· hire an initial random assignment, select var that violate constraint, assignⁿ of var to value than min # of violated constraints

· local heuristic search, may not lead to solⁿ

Adversarial Search-Games as Search problems.

- Abstraction : ideal repⁿ of real world
Accurate formulation: state space repⁿ
- complexity: Abstract by not simple.
- time limited: complete search difficult \rightarrow uncertainty on actions desirability.

Types

\Rightarrow Randomness in game

Types	deterministic	Chance
Perfect Info	chess, checkers, Go, othello	Backgammon, Monopoly
Imperfect Info.		Bridge, Poker, Scrabble Nuclear war

each player has complete info about his opponent posⁿ & about choices avail to him

Min-max Search Strategy.

\Rightarrow Maximise one's own utility & minimize the opponent's. Assumpⁿ: opponent does the same.

3 step.

- Calc. utility.
1. Generate entire game tree down to terminal states
 2. A. Assess the utility of each terminal state
B. Determine the best utility of parents of terminal state
C. Repeat the process for parent until root is reached

3. Select the best moves
(i.e. moves with highest utility)

Assumes perfect decisions.
No time limit

For Imperfect Decisions.

i.e. Time & Space requirements \rightarrow complete game tree search is intractable \rightarrow impractical to make perfect decisions

Modifⁿ to Min-max Algo

1. Replace utility function by Evalⁿ func (An estimated desirability of posⁿ)

(expected utility of game from given posⁿ)

Requirements of Evalⁿ func

- Computationally Efficient
- Agrees w/ utility func on terminal states
- Accurately reflects chances of winning

A tradeoff
btw accuracy &
efficiency.

2. Partial Search tree.

E.g. depth-limit

Replace terminal test by cut-off test

* Quiescent Search

\rightarrow eval func only applied to quiescent posⁿ

\hookrightarrow posⁿ that are not likely to have a large variation in eval in the near future

pretty confident w/ eval funcⁿ

Module 3: Markov Decision Process.

MDP components: $\{S, T, A, R\}$

State transition func Action Reward

Solution: policy: $\pi(s)$ mapping from states to actions

Value Function: $V(s)$ ↗ Five of multiple trajectories.

- Refers to expected sum of discounted rewards if agent executes optimal policy π starting in state s

$$V^\pi(s) = \mathbb{E}_\pi [G_t | S_t = s]$$

Action value of state-action pair: $Q^\pi(s, a) = \mathbb{E}_\pi[G_t | S_t = s, A_t = a]$

$$V^\pi(s) = \sum_{a \in A} Q^\pi(s, a) \pi(a|s)$$

Bellman Eqn.

$$Q(s, a) = \sum_{s'} P(s'|s, a) [r(s, a, s') + \gamma V(s')] \quad \begin{matrix} \text{↳ discount factor} \\ \downarrow \\ \text{Q value of taking action } a \text{ in state } s \end{matrix} \quad \begin{matrix} \text{↳ Value of state } (s') \\ \text{↳ immediate reward after taking action } a \text{ in state } s \text{ & transitioning to } s' \end{matrix}$$

To choose optimal policy, choose action that max Q-value for each state

$$\text{optimal action } \pi^*(s) = \arg \max_a Q(s, a)$$

$$V(s) = \max_{a \in A} Q(s, a)$$

Optimal value function V^* is derived from optimal Q-values by taking the max over all possible actions

Value Iteration Function

Initialise $V_0(s)$, for all s

$$\text{For } i=0, 1, \dots, V_{i+1}(s) = \max_a \sum_{s' \in S} P(s'|s, a) [r(s, a, s') + \gamma V_i(s')]$$

Repeat this process until converges.

use V^* to find optimal policy π^* in each state.

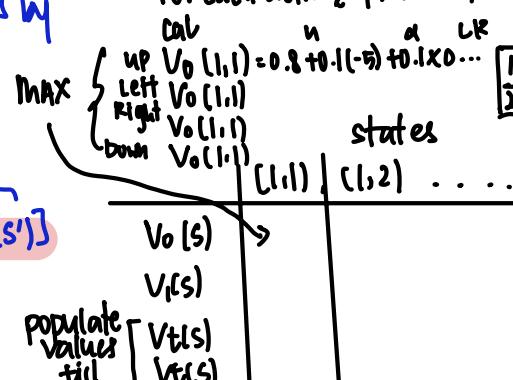
$$\pi^*(s) = \arg \max_a \sum_{s'} P(s'|s, a) [r(s, a, s') + \gamma V(s')] = \arg \max_a Q(s, a)$$

optimal policy max accumulated rewards over trajectory

$$T = \langle S_1, A_1, R_1, \dots, S_T, A_T, R_T \rangle$$

$$G_t = R_t + \gamma R_{t+1} + \gamma^2 R_{t+2} + \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k} \leq \frac{R_{\max}}{1-\gamma} \quad (0 < \gamma < 1)$$

For each action $\{up, down, left, right\}$



1	2
2	-1

Policy Iteration

Start with some initial policy, π_0 & alternate b/w:

- Policy evaluation: calculate $V^{\pi_i}(s)$ for every state, s

$$V^{\pi_i}(s) = \sum_{s'} P(s, \pi_i(s), s') [R(s, \pi_i(s), s') + \gamma V^{\pi_i}(s')]$$

- Policy improvement: calculate new policy π_{i+1} based on update utility

$$\pi_{i+1}(s) = \operatorname{argmax}_{a \in A} \sum_{s'} P(s'|s, a) [r(s, a, s') + \gamma V^{\pi_i}(s')]$$

until process converges

→ Initial policy $\pi(s_1) = \text{Right}$
 $\pi(s_2) = \text{Right}$

Policy Eval

$$V(s_1) =$$

$$V(s_2) =$$

$$\begin{aligned} \pi(s_{11}) &= \text{Right} \\ \pi(\text{Right}) &= 0.79 \\ \pi(\text{Left}) &= 0.775 \end{aligned} \quad \left[\text{more} \right]$$

* Algo

1. Initialise $V(s)$ to $\pi(s)$ for all states s

2. Policy Evaluation (calculate the V)

- a. Repeat
- b. $\Delta \leftarrow 0$
- c. For each states:
- d. $v \leftarrow V(s)$
- e. $V(s) \leftarrow \sum_{s'} P(s|s, \pi(s)) [R(s, \pi(s), s') + \gamma V(s')]$
- f. $\Delta \leftarrow \max(\Delta, |v - V(s)|)$
- g. until $\Delta < w$ (small threshold)

3. Policy Improvement (calculate new policy π)

- a. policystable $\leftarrow \text{True}$
- b. For each state s :
- c. $b \leftarrow \pi(s)$
- d. $\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s'} P(s'|s, a) [R(s, a, s') + \gamma V(s')]$
- e. if $b \neq \pi(s)$, the policystable $\leftarrow \text{False}$
- f. if policystable $\leftarrow \text{True}$, then stop, go to step 2

① calculate value

states	S_0	$V(S_0) = 8.03$
	S_1	$V(S_1) = 11.2$
	S_2	$V(S_2) = 8.9$

② calculate optimal policy

π^*	a_0	a_1	Actions
S_0	5	6	
S_1			
S_2			

optimal policy
 $\Rightarrow \pi(S_0) = a_1$
 $\pi(S_1) = 11$
 $\pi(S_2) = 11$

Module 5: Reinforcement Learning

Reinforcement Learning: learn value funcⁿ w/o the transition function $P(s'|s, a)$

→ don't know what happens after

taking an action

1. Monte Carlo (use of randomness)

- Solve of problem by generating suitable random numbers & obtaining the fraction of numbers by obeying some properties

→ only for evaluating & finding $V(s) = \text{expected cumulative}$

reward for states for given policy.

* MC for estimating $V(s)$

→ Using some policy, calculate accumulated rewards for states in each episode

→ $V(s)$ is the ave of accumulated rewards over all episodes.

First visit MC policy Evalⁿ

- Initialize

- $\pi \leftarrow$ policy to be evaluated.

- $V \leftarrow$ an arbitrary state-value fun

- $\text{Returns} \leftarrow$ an empty list for all states

- Repeat many times.

- Generate an episode with π

- For each state appearing in the episode

 - $R \leftarrow$ Return first occurrence of s

 - Append R to Returns (s)

 - $V(s) \leftarrow \text{Ave}(\text{Returns}(s))$

$i : \# \text{ of episodes}$
for state s_j

$$\sum_i^L G_t \rightarrow$$

$$G_t = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_n \\ = \sum_i^n \gamma^i r_{t+i}$$

* MC for estimating Q -values.

→ objective: find optimal policy by

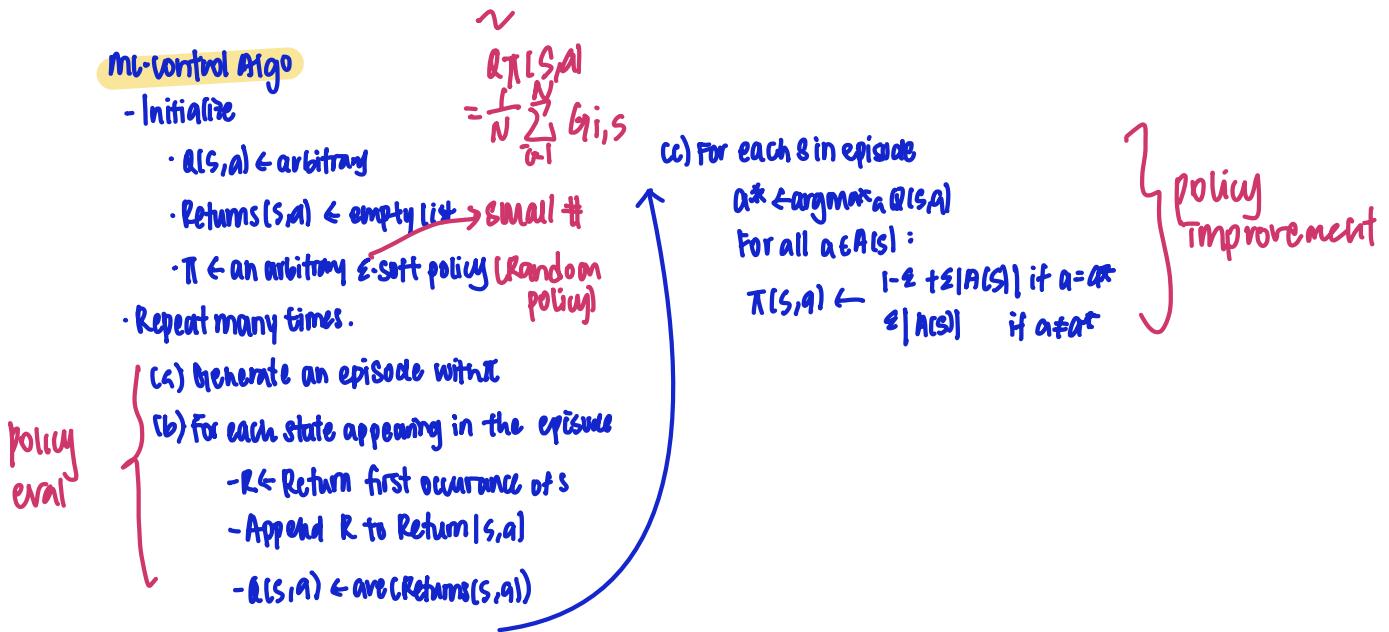
estimating action-value func

→ calculate accumulated reward for state-action pairs instead of state. Ave a_t rewards across all episodes.

s_1	s_2	s_3
a_1	a_1	a_1
a_2	a_2	a_2
a_3	a_3	a_3

Cell 1	Cell 2	Cell 3	✓	At cell 1, 2, 3 choose R_t .
R_t	8.1	9.3	9.9	
L_t	1	5	4	

→ Selecting an action: $\pi'(s,a) = \underset{a \in A}{\text{argmax}} Q^\pi(s,a)$



2. Q-Learning.

- MC needs the whole trajectory.
- Q-learning uses only one step trajectory: (s, a, r, s') to estimate Q-value.
- Update Q-value of s, a pair immediately after every observation.

$$\begin{aligned} Q_{\text{new}}(s_t, a_t) &\leftarrow Q_{\text{old}}(s_t, a_t) + \alpha [r_{t+1} + \gamma \max_a Q_{\text{old}}(s_{t+1}, a) - Q_{\text{old}}(s_t, a_t)] \\ &= (1 - \alpha) Q_{\text{old}}(s_t, a_t) + \alpha [r_{t+1} + \gamma \max_a Q_{\text{old}}(s_{t+1}, a)] \end{aligned}$$

↳ learning ↳ discount
Rate Factor

Algo params: step size $\alpha \in [0, 1]$, small $\epsilon > 0$

Initialise $Q(s, a)$ for all $s \in S^+, a \in A(s)$, except $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize s

Loop for each step of episode: $\max_{a \in A(s)}$
 (ε-Greedy)

Choose a from s using policy derived from

Take action a , observe r, s'

update $Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_a Q(s', a) - Q(s, a)]$

$s \leftarrow s'$

until s is terminal

→ Maintain Q table:

$Q = S \left[\begin{array}{c} \text{Action} \\ \vdots \end{array} \right] \text{ State-action pairs.}$

3. DQN

→ Q-learning needs to maintain a table, not feasible w/ large state / action space

→ use N.N as black box to replace table

$$s \xrightarrow{q} \boxed{w_j} \xrightarrow{?} (s, a, w_j)$$

Q-value

Module 6: Game Theory.

Nash Equilibrium.

- Each agent is selfish
- Each agent makes decision based on what he thinks others will do
- No one can do better by changing strategy alone

Eg Prisoners Dilemma

		confess	silent (B)
		A	B
(A) confess	confess	5 yrs	5 yrs
	silent	20 yrs	0 yrs

- confessing always has a higher payoff

cooperation is good for both to
But since agent is selfish →
will always choose to defect

Pure Strategy N.E

- For two strategies S_1 (for i) & S_2 (for j) in N.E
- Assumes: 1. Agent i plays s_1 and agent j can do no better than play s_2
2. Agent j plays s_2 & agent i can do no better than play s_1
- Neither agent has any incentive to deviate from N.E

"specific unchanging But: 1) Not every interaction has pure strategy N.E
choices throughout game" 2) Some scenarios may have >1 pure strategy N.E

Mixed Strategy N.E

- No pair of strategies forms a pure strategy N.E:

- Whatever pair of strategy chosen, somebody wishes they had done something else

∴ $\exists S_i^m$: Allow mixed strategy

- play 'heads' with prob 0.5
- play 'tails' with prob 0.5

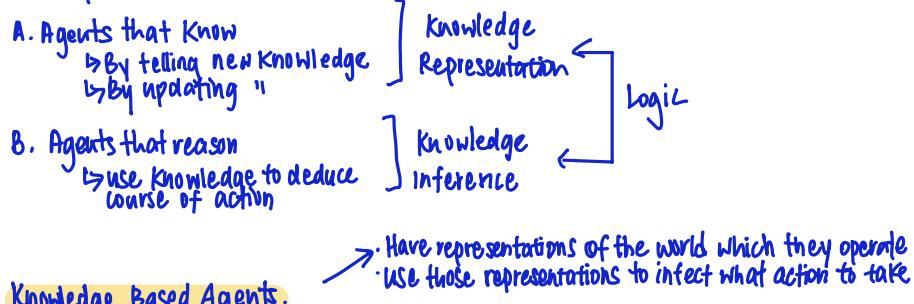
e.g. Matching Pennies Payoff Matrix

		i heads	j tails
		1	-1
j heads	i heads	-1	1
	i tails	1	-1

same face $\Rightarrow i$ wins
different face $\Rightarrow j$ wins

Week 8: Logical Agents.

KB Approach



KB

- Set of sentences / Rep² of facts
- KR language
 - tell
 - ASK
- Also infer from KB.

KB-System

- States: Instances of KB
 - Build by using tell
- Operators
 - Add/infer new sentence
- Goal: Answer & query
 - use ASK

Levels of Knowledge in KB

Range from background (observed)

& deduced

1) Epistemological - Declarative Description of Knowledge
e.g. facts: smoke in kitchens

2) Logical - Logical encoding of knowledge
e.g. facts: smoke
rules: implies (smoke, fire)

3) Implementation - Physical representation of knowledge (sentences)
e.g. implies (smoke, fire)

Knowledge Representations.

KR language = Syntax + Semantics
(Implementation) level (Knowledge) level

* Logic = Inference + Representation
Syntax + Semantics

A) [Inference]

+ Reasoning: construction of new knowledge from existing ones. Act of deriving logical conclusions from known premise (assumed to be true)

Mechanical Reasoning.

A) Deductive Inference (sound)

KB: Monday → work, Monday ⊢ work

Tell: give KB : (KB ⊢ x)!

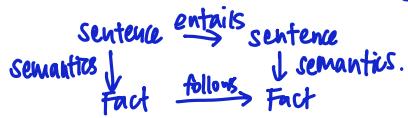
ASK: given KB & x : (KB ⊢ x) ?

B) Inductive Inference (unsound)

KB: Monday → work, work ⊢ Monday

(Form of logical Inference)

Entailment : Returns true given existing sentence return true



* USEFUL PROPERTIES

- sentence is T/F under interpretation of world
- Refers to inference procedure
- A) valid Sentence (tautology)
iff TRUE under all interpretations e.g. $s \lor \neg s$
 - B) satisfiable
iff TRUE under some interpretations e.g. s and $\neg s$ in some world
 - C) soundness (e.g. MP)
- sound if it never proves a false statement; only derives conclusions that are true whenever premise are true
 - d) completeness
- complete if it can prove all true statements valid

B) Representations (Syntax + Semantics)

* Desired properties of representations

- Expressive
- context independent
- unambiguous
- concise
- effective

* Desired properties of semantics

- correspondence b/w sentences & facts
- arbitrary meaning, fixed by writer of sentence.
- systematic rules - Meaning of sentence is a function of the meaning of its parts
- makes claim about world \Rightarrow Returns T/F

Week 9: propositional logic

* Syntax: of the representation language specifies all sentences that are well formed.

* Semantics: defines truth of each sentence w.r.t. each possible world

(1) Syntax (Backus-Naur Form)

Syntax of Propositional Logic (Backus-Naur Form)

Sentence	\rightarrow	AtomicSentence ComplexSentence
AtomicSentence	\rightarrow	LogicalConstant PropositionalSymbol
ComplexSentence	\rightarrow	(Sentence) Sentence LogicalConnective Sentence -Sentence
LogicalConstant	\rightarrow	TRUE FALSE
PropositionalSymbol	\rightarrow	P Q R ...
LogicalConnective	\rightarrow	\wedge \vee \Rightarrow \Leftrightarrow \neg

Precedence (from highest to lowest): \neg , \wedge , \vee , \Rightarrow , \Leftrightarrow

e.g.: $\neg P \wedge Q \vee R \Rightarrow S$ (not ambiguous), equal to: $((\neg P) \wedge Q) \vee R \Rightarrow S$

7

(2) validity

Semantics of Propositional Logic

• Validity

- A sentence is valid if it is true in all models.
- Valid sentences are known as **tautologies**.
- Every valid sentence is logically equivalent to **True**.

Semantics of Propositional Logic

• Interpretation of symbols

- Logical constants have fixed meaning
 - True: always means the fact is the case; valid
 - False: always means the fact is not the case; unsatisfiable
- Propositional symbols mean “whatever they mean”
 - e.g.: P “we are in a pit”, etc.
 - Satisfiable, but not valid (true only when the fact is the case)

• Interpretation of sentences

- Meaning derived from the meaning of its parts
 - Sentence as a combination of sentences using connectives
- Logical connectives as (boolean) functions:

TruthValue $f(\text{TruthValue}, \text{TruthValue})$

10

Semantics of Propositional Logic

• Satisfiability

- A sentence is satisfiable if it is true in some models.
- Satisfiability can be checked by enumerating the possible models until one is found that satisfies the sentence.
- Most problems in computer sciences are satisfiability problems.
 - E.g., Constraint satisfaction problem, Search problems.

* Types

Literal: single proposition / negative e.g. p or $\neg p$

Clauses: propositions formed from finite collection of literals & logical connectives e.g. pvevR

Classic Rules of Inference

1. MP/Implicit Elimination: $\frac{\alpha = \beta, \alpha}{\beta}$

2. And Elimination $\frac{\alpha_1 \wedge \alpha_2 \dots \wedge \alpha_n}{\alpha_i}$

3. And Introduction: $\frac{\alpha_1, \alpha_2 \dots \alpha_n}{\alpha_1 \wedge \alpha_2 \dots \wedge \alpha_n}$

4. Or Introduction $\frac{\alpha_1}{\alpha_1 \vee \alpha_2 \dots \vee \alpha_n}$

5 Double Negation $\frac{}{\neg \neg \alpha}$

Resolution

1. Unit Resolution $\frac{\alpha}{\neg \alpha}$

2. Full Resolution: $\frac{\alpha \vee \beta}{\neg \alpha \vee \beta}$

Equivalence Rules

- Equivalent notations

e.g., MP:

$$\frac{\alpha \Rightarrow \beta, \alpha}{\beta}$$

1) $\alpha \Rightarrow \beta, \alpha \vdash \beta$
 2) $\alpha \Rightarrow \beta, \alpha \models \beta$
 3) $\frac{\alpha}{\beta}$
 4) $((\alpha \Rightarrow \beta) \wedge \alpha) \Rightarrow \beta$

Equivalence Rules

- Equivalence rules

- Associativity:

$$\begin{aligned} \alpha \wedge (\beta \wedge \gamma) &\Leftrightarrow (\alpha \wedge \beta) \wedge \gamma \\ \alpha \vee (\beta \vee \gamma) &\Leftrightarrow (\alpha \vee \beta) \vee \gamma \end{aligned}$$

- Distributivity:

$$\begin{aligned} \alpha \wedge (\beta \vee \gamma) &\Leftrightarrow (\alpha \wedge \beta) \vee (\alpha \wedge \gamma) \\ \alpha \vee (\beta \wedge \gamma) &\Leftrightarrow (\alpha \vee \beta) \wedge (\alpha \vee \gamma) \end{aligned}$$

- De Morgan's Law:

$$\begin{aligned} \neg(\alpha \vee \beta) &\Leftrightarrow \neg\alpha \wedge \neg\beta \\ \neg(\alpha \wedge \beta) &\Leftrightarrow \neg\alpha \vee \neg\beta \end{aligned}$$

Limitations of propositional logic

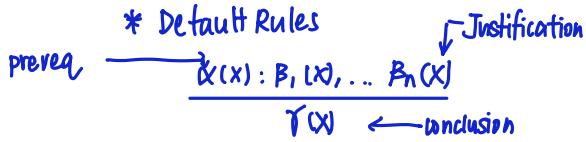
Given N objects are used in KB, there will be 2^N interpretations/worlds.

Week 10: First order Logic

Week 11: Default Logic

- Standard logic only express something is T/F based on explicitly given rules in KB.
- Default Logic is used to formalize reasoning with default assumptions. i.e. common sense \Rightarrow if nothing in KB conflicts with common sense, then infer something else)

* Default Theory = Default Rules + KB.



Applied by substituting x with c (ground instance) to $\alpha \& \beta$ to infer γ

$\alpha(c) \Rightarrow$ belongs to our set of beliefs \Rightarrow infer $\gamma(c)$ as true

Each $\beta(c) \Rightarrow$ consistent with our set of beliefs

Types of Default Rules

- Normal Defaults: $\frac{\alpha(x) : \gamma(x)}{\gamma(x)}$
 - Semi-Normal Defaults: $\frac{\alpha(x) : \beta(x)}{\gamma(x)}$, where $\beta(x) \vdash \gamma(x)$
- E.g., $\frac{\text{bird}(x) : \text{has_wings}(x)}{\text{flies}(x)}$,
where $\text{has_wings}(x) \vdash \text{flies}(x)$

Types of Default Rules

- Open Defaults (Default Schemas) have unbounded variables, e.g., x
- $$\frac{\alpha(x) : \beta_1(x), \dots, \beta_n(x)}{\gamma(x)}$$
- Closed (Grounded) Defaults use ground terms, e.g., $x=c$
- $$\frac{\alpha(c) : \beta_1(c), \dots, \beta_n(c)}{\gamma(c)}$$

* Set of Default Rules

Default Theory $\Delta, \phi \models \Phi$: set of facts given KB

e.g. $\Delta, \phi \models \{ \frac{\text{bird}(x) : \text{flies}(x)}{\text{flies}(x)}, \{(\text{bird}(\text{Tweety}), \text{cat}(\text{Sylvester})) \}$

* Types of Default Logic Inference

(1) Reiter Default Logic Inference (RDL)

* Algo

* Example

Reiter Default Logic (RDL) Inference

- Guess the extension Ξ (pronounced as "Xi")
 - Initialise beliefs $\Xi^* = \Phi \Rightarrow$ initialise initial belief to fact part in KB.
 - (loop over) c-ground instance of an (unused) default $\frac{\alpha(x) : \beta(x)}{\gamma(x)}$:
 - Check two conditions \rightarrow is α a member of class or
 - Triggered?: $\Xi^* \vdash \alpha(c)$
 - Justified?: $\beta(c)$ is consistent with Ξ
 - If yes: update beliefs $\Xi^* \leftarrow \Xi^* \cup \{\gamma(c)\}$ \rightarrow look through KB to see if any fact contradicts $\beta(c)$
 - (end loop)
 - If $\Xi = \Xi^*$ then extension found/confirmed
- The extension is added to our KB as new knowledge
- Add new sentence to KB that $\beta(c)$ is true

Example

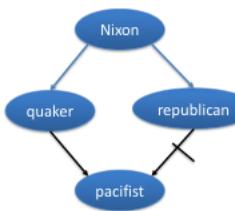
Given Theory: $T = \left\langle \Delta = \left\{ \frac{\text{bird}(x) : \text{flies}(x)}{\text{flies}(x)} \right\}, \Phi = \{(\text{bird}(\text{Tweety}), \text{cat}(\text{Sylvester})) \} \right\rangle$

- Guess the extension = $Cn(\{\text{flies}(\text{Tweety})\} \cup \Phi)$
- Our initial knowledge is $F = \Phi$
- Sylvester-instance of default not applicable:
 - not hold $\Phi \vdash \text{bird}(\text{Sylvester})$
- $\Phi \vdash \text{bird}(\text{Tweety})$ and $\text{flies}(\text{Tweety})$ is consistent with F
- $F = \Phi \cup \{\text{flies}(\text{Tweety})\}$
- No more default rules to apply
- An extension is reached

*Limitations

RDL Limitation - Nixon Diamond

The default rules may be applied in different orders, and this may lead to different extensions. E.g.:



Let Theory $T = \langle \Delta, \Phi \rangle$

$$\begin{aligned} \bullet &= \left\{ \frac{\text{quaker}(x) : \text{pacifist}(x)}{\text{pacifist}(x)}, \frac{\text{republican}(x) : \neg \text{pacifist}(x)}{\neg \text{pacifist}(x)} \right\} \\ \bullet \Phi &= \{ \text{quaker}(\text{Nixon}), \text{republican}(\text{Nixon}) \} \end{aligned}$$

16

RDL Limitation - Nixon Diamond

Given $\langle \Delta, \Phi \rangle$

- $\Delta = \left\{ \frac{\text{quaker}(x) : \text{pacifist}(x)}{\text{pacifist}(x)}, \frac{\text{republican}(x) : \neg \text{pacifist}(x)}{\neg \text{pacifist}(x)} \right\}$
- $\Phi = \{ \text{quaker}(\text{Nixon}), \text{republican}(\text{Nixon}) \}$

There are **two** extensions:

1. One that contains: **pacifist(Nixon)**
2. ... and the one that contains: **$\neg \text{pacifist}(\text{Nixon})$**

if analyse Republican first
the conclusion will be
diff. : conflicts

Conflicting conclusions

17

Addressing the RDL Limitation

- A default theory can have 0, 1 or more extensions.
- Entailment of a formula from a default theory can be defined in one of two ways:

decision
maker's
properties
IR

- Skeptical:
 - a formula is entailed by a default theory if it is entailed by all its extensions.
- Credulous: (relaxed)
 - a formula is entailed by a default theory if it is entailed by at least one of its extensions.

18

Example

- The Nixon diamond example theory has two extensions:
 - one in which Nixon is a **pacifist**; and
 - one in which Nixon is **not a pacifist**.
- Thus, we have:
 - Neither **Pacifist(Nixon)** nor **$\neg \text{Pacifist}(\text{Nixon})$** are skeptically entailed.
 - Both **Pacifist(Nixon)** and **$\neg \text{Pacifist}(\text{Nixon})$** are credulously entailed.
- The **credulous extensions** of a default theory can be **inconsistent with each other**.

(2) Makinson Approach

Makinson Approach

- Order **ground** instances of defaults in $\Delta: d_1, d_2, \dots$
- Initialize beliefs $\Xi_0 = \Phi$ and used defaults set $\Delta_0 = \emptyset$
- Define Ξ_{n+1} from Ξ_n
 - Find $d = \frac{\alpha(c): \beta_1(c), \dots, \beta_n(c)}{\gamma(c)} \notin \Delta_n$ such that
 - Triggered?: $\Xi_n \vdash \alpha(c)$
 - Justified?: Ξ_n is consistent with $\beta_1(c), \dots, \beta_m(c)$
 - If $\Xi_n \cup \{\gamma(c)\}$ is consistent with each $\beta'(c')$ in $\Delta_n \cup \{d\}$
 - $\Xi_{n+1} = \Xi_n \cup \{\gamma(c)\}$, and $\Delta_{n+1} = \Delta_n \cup \{d\}$
 - else abort -- no extension for this order of defaults
- The extension is $\Xi = \bigcup_{i \geq 0} \Xi_i$

Operational Semantics

Given a default theory $T = \langle \Delta, \Phi \rangle$, let $\Pi = (\delta_0, \delta_1, \dots)$ be (a finite or infinite) sequence of (closed) defaults from Δ without multiple occurrences.

- $\Pi[k]$ denotes the initial segment of sequence Π with length k .
- Each sequence Π is associated with two sets:
- $In(\Pi) = Cn(\Phi \cup \{\text{consequence}(\delta) \mid \delta \text{ occurs in } \Pi\})$
 - $Out(\Pi) = \{\neg \phi \mid \phi \in \text{justifications}(\delta) \text{ for some } \delta \text{ in } \Pi\}$
- Model of the world
initial set of facts
union in conclusions for
- Everything I have denied the existence of as a result of my assumptions (i.e., the negation of all the justifications)
- For any default rule that we have considered \Rightarrow add negation of justification part of those rules

Process, Successful, Closed

Π is a process of $T = \langle \Delta, \Phi \rangle$ iff default δ_k is applicable to $In(\Pi[k])$ for every k such that δ_k occurs in Π .

Let Π be a process. We define:

- Π is **successful** iff $In(\Pi) \cap Out(\Pi) = \emptyset$ (Nothing in the out set can be inferred from the in set); Otherwise, it **fails**.
- Π is **closed** iff every $\delta \in \Delta$ that is applicable to $In(\Pi)$ already occurs in Π .

If every default rule in the current default theory that's applicable has already occurred and been extended

Extension

Makinson Approach

- No extension guessing
- Choose the order of defaults in $\Delta: d_1, d_2, \dots$
- There still may be more than one possible extension
 - Different orders of defaults can lead to different Ξ
- We get the same extensions as in Reiter's approach
 - If they exist at all

Example

Consider $T = \langle \Delta, \Phi \rangle$ with $\Phi = \{\alpha\}$ and defaults from Δ :

$$\delta_1 = \frac{\alpha : \neg \beta}{\neg \beta}, \quad \delta_2 = \frac{\beta : \gamma}{\gamma}$$

if $\alpha \vdash T$ $\neg \beta \vdash T$ $\beta \vdash T$ $\gamma \vdash T$

nothing conflicts then

For $\Pi_a = (\delta_1)$ we have $In(\Pi_a) = Cn(\{\alpha, \neg \beta\})$, $Out(\Pi_a) = \{\beta\}$

For $\Pi_b = (\delta_2, \delta_1)$ we have $In(\Pi_b) = Cn(\{\alpha, \neg \beta\})$, $Out(\Pi_b) = \{\beta\}$

$\neg R_2$ triggered first but not supported by KCB
 \therefore Apply R_1 : same as Π_a

Example

- Consider $T = \langle \Delta, \Phi \rangle$ with $\Phi = \{\alpha\}$ and defaults from Δ :
- $$\delta_1 = \frac{\alpha : \neg \beta}{\eta}, \quad \delta_2 = \frac{\text{true} : \gamma}{\beta}$$
- $\Pi_1 = (\delta_1)$ is **successful**.
 $In(\Pi_1) = Cn(\alpha, \eta)$ and $Out(\Pi_1) = \{\beta\}$
 but **not closed**, since δ_2 is applicable, too.
- $\Pi_2 = (\delta_1, \delta_2)$ is **closed**, but **not successful**.
 $In(\Pi_2) = Cn(\alpha, \eta, \beta)$ and $Out(\Pi_2) = \{\beta, \neg \gamma\}$,
 $In(\Pi_2) \cap Out(\Pi_2) = \beta$
- $\Pi_3 = (\delta_2)$ is a **closed** and **successful** process.
 $In(\Pi_3) = Cn(\alpha, \beta)$ and $Out(\Pi_3) = \{\neg \gamma\}$,
 $In(\Pi_3) \cap Out(\Pi_3) = \emptyset$
- δ_1 applies but
 - justification $\beta : T$ conflicts w/ δ_2
 $\neg \beta : \neg T$ non cumulative
 $\neg \beta_1$ triggered but not successfully extended: N.A.
- Nothing concludes $T \Rightarrow$
 $\beta \rightarrow T \rightarrow$ Add to in & negation of $\neg \beta$ out

Process Tree

Extension

- Let $T = \langle \Delta, \Phi \rangle$ be a default theory. A set of formulae Σ is an **extension** of T iff there is some **closed and successful** Π such that $\Sigma = In(\Pi)$.
- To **find a successful process**: generate a process Π , test whether $in(\Pi) \cap Out(\Pi) = \emptyset$. If not, then backtrack (try another process).

(3) Process Tree

extended:

Process Tree

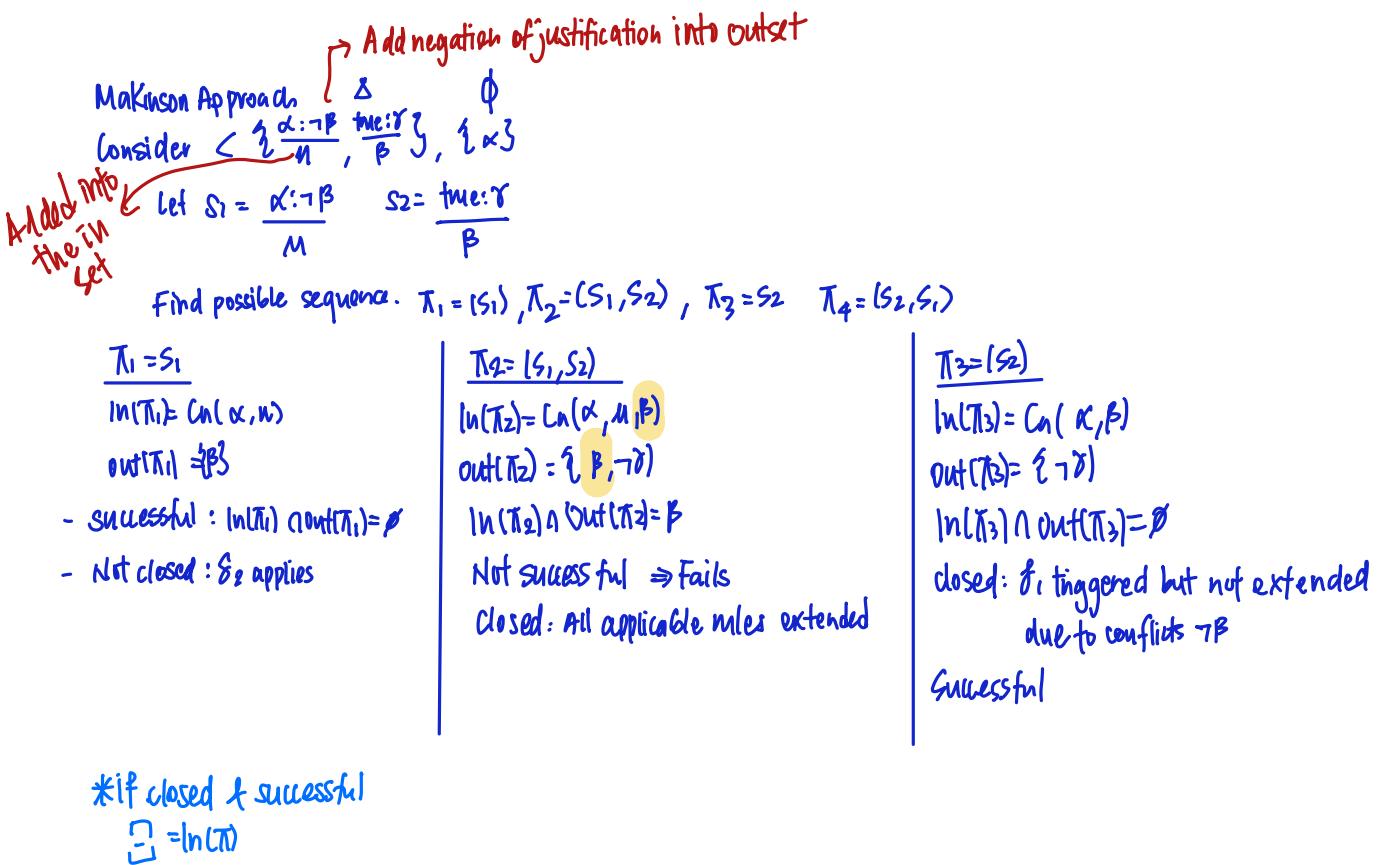
$T = \langle \Delta, \Phi \rangle$ be a default theory. A **process tree** is a tree $G = (V, E)$ such that all nodes $v \in V$ are labelled with two sets of formulae:

- an In-set $In(v)$ and
- an Out-set $Out(v)$.

The root of G is labelled with $Cn(\Phi)$ as the In-set and \emptyset as the Out-set. Every $e \in E$ denotes a default application and is labelled by it.

A process is thus a path in G starting from the root.

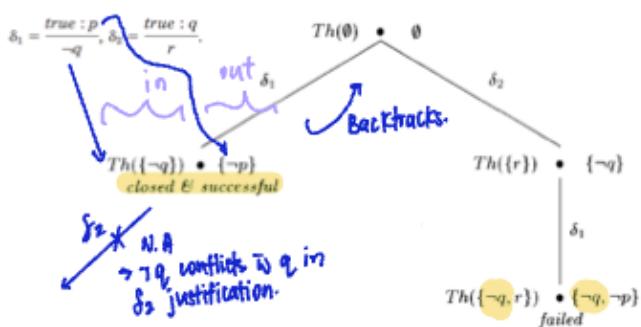
A node $v \in V$ is **expanded** if $In(v) \cap Out(v) = \emptyset$. Otherwise, it is a "failed" leaf of the tree.



Process tree

Process Tree Example

Let $T = (W, D)$ be the default theory with $W = \emptyset$ and $D = [\delta_1, \delta_2]$ with



Process Tree: Properties

- A process is thus a path in G starting from root.
- A node $v \in V$ is **expanded** if $\text{In}(v) \cap \text{Out}(v) = \emptyset$.
- Otherwise, it is a "failed" leaf of the tree.
- Expanded $v \in V$ has a child node, w_δ , for every $\delta = \frac{\alpha : \beta_1, \dots, \beta_n}{\gamma}$
 - w_δ does not appear on the path from the root to v
 - δ is applicable to $\text{In}(v)$
 - w_δ connected to v by an edge labelled with δ
 - w_δ is labelled with $\text{In}(w_\delta) = Cn(\text{In}(v) \cup \{\gamma\})$ and $\text{Out}(w_\delta) = \text{Out}(v) \cup \{\neg\beta_1, \dots, \neg\beta_n\}$