



# Lecture 3: Buffer Overflows

*presented by*

**Li Yi**

*Assistant Professor*  
SCSE

N4-02b-64

[yi\\_li@ntu.edu.sg](mailto:yi_li@ntu.edu.sg)

# Introduction

- The course will work its way up from the operating system to the application layer
- We will start by looking for vulnerabilities in the way programs get executed
- The discussion will focus on general principles
  - For constructing a successful attack, details of the platform attacked need to be taken into account
  - “Old” attacks rarely work on today’s platforms

# Agenda

- Background
  - Variables and buffers
  - Call stack
- Buffer overflow attacks
- Defences
  - Safe functions
  - Canaries
  - Split control and data stack
  - Data Execution Prevention
  - Address Space Layout Randomization
- Return-oriented programming

# Background: C/C++

- Implementations of network protocols are often written in C/C++
  - Typical task: “serializing” of some composite data structure, i.e., package the data structure into a string, send it, unpack it at the receiver
  - Makes use of **string operations**
- Trade-off: **performance – robustness**
  - Management of memory objects is often intentionally left to the programmer for performance optimizations
  - You have to know what you are doing and be circumspect to get your code right

# A very simple Web server

```
#include <stdio.h>

int read_req(void) {
    char buf[128];
    int i;
    gets(buf);
    i = atoi(buf);
    return i;
}

int main(int ac, char **av) {
    int x = read_req();
    printf("x=%d\n", x);
}
```

```
$ ./readreq
```

```
123
```

```
x=123
```

```
$ ./readreq
```

```
148214899412412841241241
```

```
x=2147483647
```

```
$ ./readreq
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
AAAAAAAAAAAAAAAAAAAAAAAAAAAA
```

```
Segmentation fault (core  
dumped)
```

What do you think has happened?

# Buffer Overflows (1980s)

- Log-in in a version of Digital's VMS operating system: to log in to a particular machine, enter

`username/DEVICE =<machine>`

- Length of the argument “`machine`” was not checked; a device name of more than 132 bytes overwrote the privilege mask of the process started by login
- Users could thus set their own privileges

# Memory Access in C/C++

- **Strings** are written to / read from memory
- Memory is accessed via **pointers**
  - Pointers are variables that hold **addresses** as values
- **“Unsafe”** write: given a pointer and a string, **write to memory** starting at the address pointed to **until the end of the string**
  - **gets, strcpy, sprintf, ...**
  - **NUL** character terminates strings
  - **No warning when too many characters are written!**
- **“Safe”** write: given a pointer, a string, and a **bound**, write to memory starting at the address pointed to until end of string **or the bound** is reached
  - **fgets, strncpy, snprintf, ...**

# Variables & Buffers

- **Abstract view:** programs store data in **variables**
- **Implementation:** allocate a region of memory (a.k.a. **buffer**) to store the value assigned to a variable
- What can go wrong when assigning a value to a variable?
  - Miscalculate position of buffer and write to a wrong location
  - Write more data than the buffer can hold

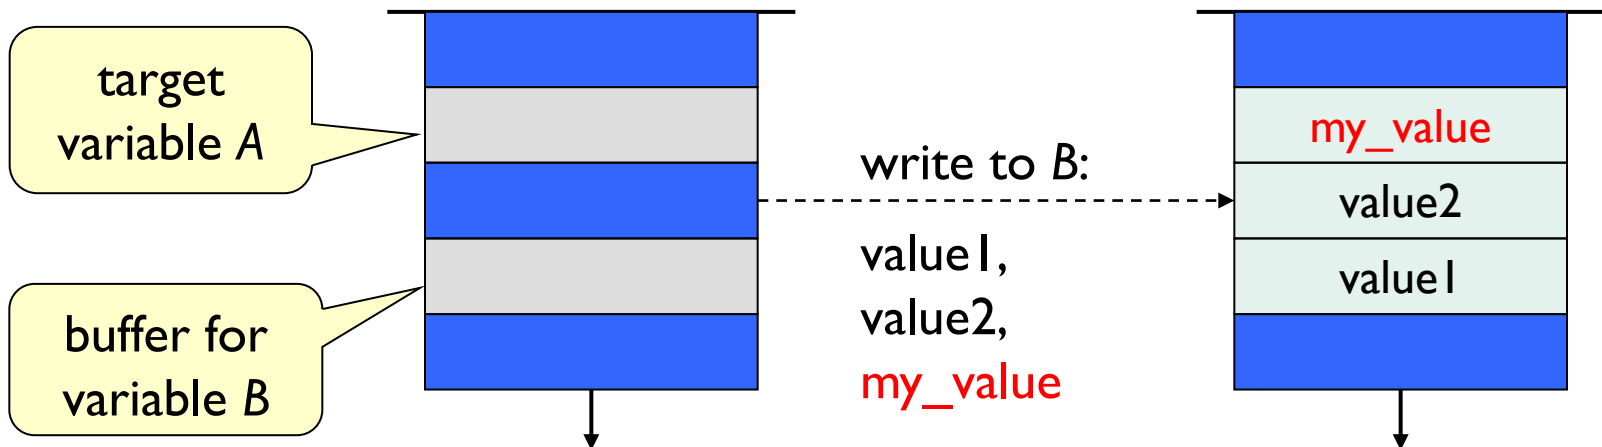


# Buffer Overflows

- If the value assigned to a variable exceeds the size of the allocated buffer, memory not allocated to this variable is overwritten
- This is called a **buffer overflow (overrun)**
- If the memory location overwritten had been allocated to another variable, the value of that variable is changed
  - “Access via the layer below”
    - Value of a sensitive variable **A** could be changed by assigning a (deliberately) malformed value to some other variable **B**
- Depending on the location of the buffer, there are
  - **Stack-based buffer overflow** (covered in this lecture)
  - **Heap-based buffer overflow** (e.g., [CVE-2021-3156](#): the “sudo” vulnerability)

# Buffer Overflows

- Assign a value to variable *A* by writing to variable *B*
- Data written to a buffer is **written upwards** ↑ (towards higher addresses) from address of buffer



# Buffer Overflows

- **Unintentional buffer overflows** crash software and have been a focus for reliability testing
- **Intentional buffer overflows** are a concern if an attacker can modify security relevant data
- Attractive targets are **return address** (specifies next piece of code to be executed) and **security settings**
- **Type-safe** languages like Java guarantee that memory management is “error-free” (more later)

# The Call Stack


Calling and returning from functions

# Function Calls

- Structured programs use **functions** (subroutines, methods, etc.) which may in turn call further functions
- When the runtime environment executing a program encounters a function call, it collects the data needed for executing the function in a **frame**, stores the frame, and starts executing the function
- Runtime environment returns to caller when the execution of the function is completed
- Function calls can be nested so frames are stored on a **call stack** (system stack)

# Stack & Heap

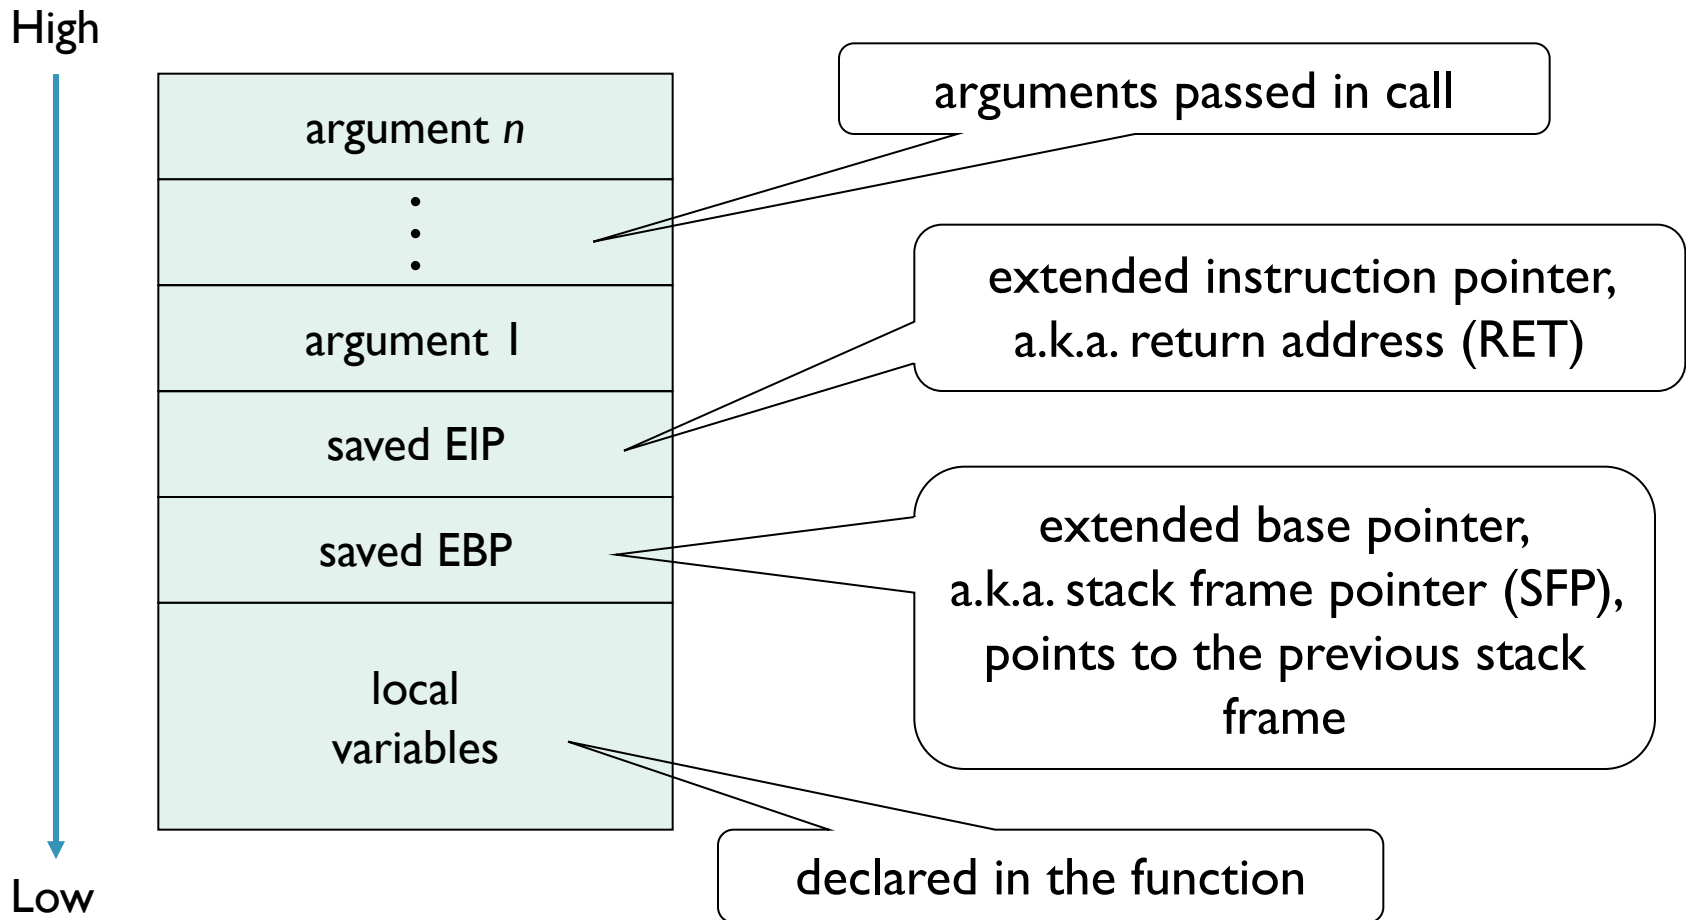
- Stack: stack frames contain **return address**, **local variables** and **function arguments**; relatively easy to figure out in advance where a given buffer will be placed on the stack
- Heap: dynamically allocated memory; difficulty of guessing where a given buffer will be taken from the heap depends on memory allocation scheme



# Stack Frames

- Stack frame contains function arguments, return address, statically allocated buffers, and more
- When the call returns, execution continues at the return address specified
- Call stack by convention starts at the top of memory and grows downwards ↓
- Layout of stack frames is reasonably predictable (depending on compiler, operating system, ...)
- Stack frame contains both user data and control data!

# Stack Frame – Typical Layout

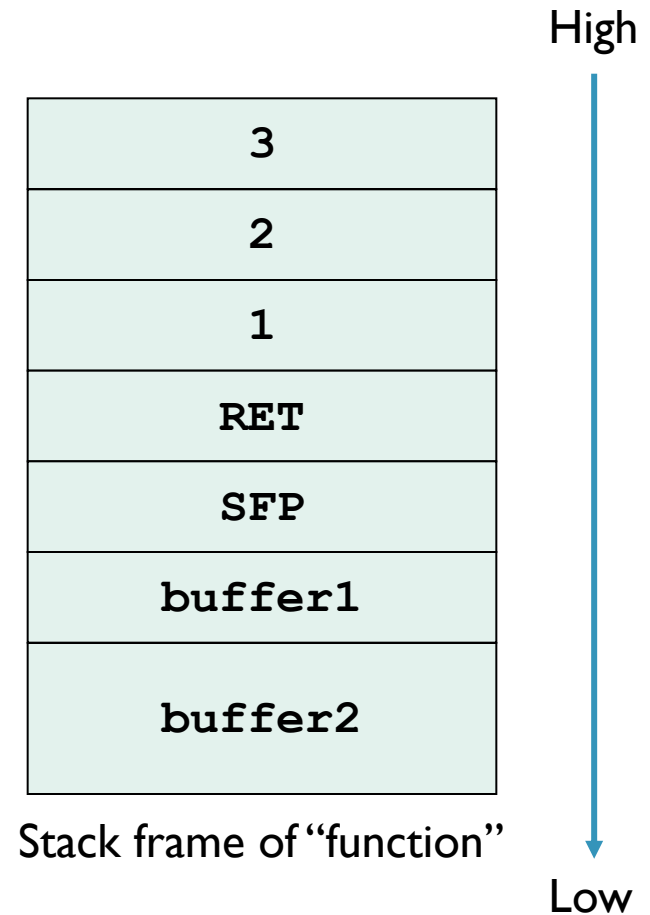




## Stack Frame – Example

```
void function(int a, int b,  
             int c)  
{  
    char buffer1[5];  
    char buffer2[10];  
}  
  
void main()  
{  
    function(1,2,3);  
}
```

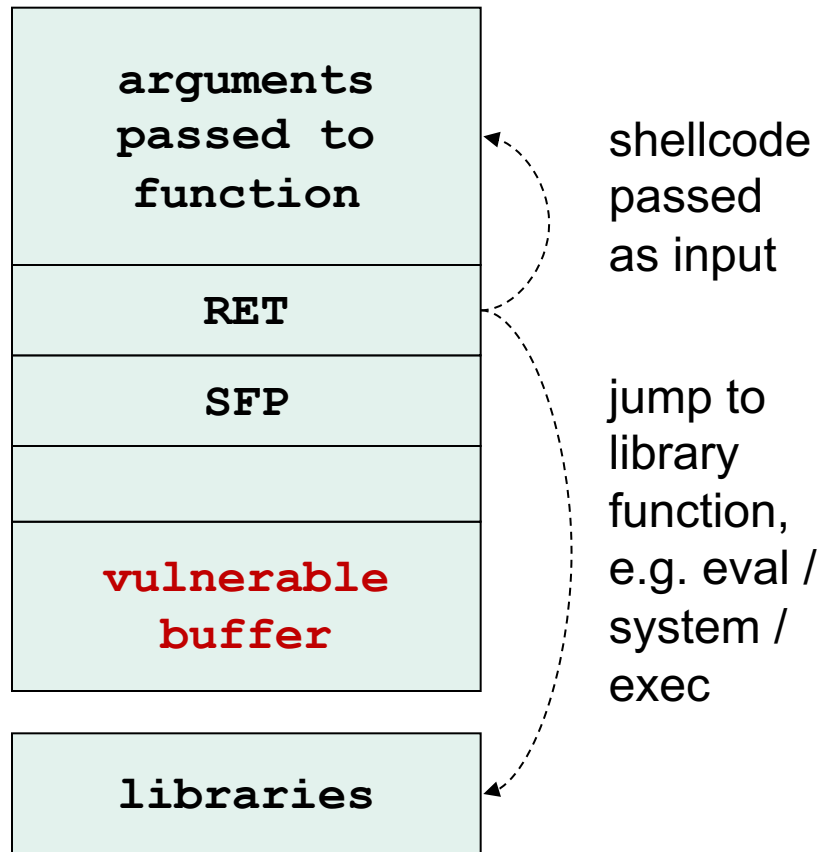
More examples in tutorial ...



# Buffer Overflow Attacks – Pattern

- Find a function that contains an “unsafe” **write of user defined input** to a local variable
- Attacker supplies specially crafted input to function
- Attacker’s input overflows buffer allocated to that local variable until it overwrites return address to make return address point to the attacker’s code
- Attacker’s code commonly known as “**shellcode**”

# Stack-Based Buffer Overflows



1. Find a buffer on the stack of a privileged program that can be overflowed
2. Overflow buffer to overwrite return address with start address of the code you want to execute
3. Jump to your code; your code is now privileged too

# “Classic” Code Example

- Declare a local short string variable

```
char buffer[80]  
gets(buffer);
```

and use standard C library routine call to read single text line from standard input and save it into `buffer`

- Works for a line of a typical character-based terminal; corrupts the stack if input is longer than 79 characters
- Attacker loads malign code into buffer and redirects return address to the start of shellcode

# Details (not discussed in detail)

More technicalities need to be considered in this and related attacks


- Attacker's malign input must not contain terminating characters; strings in C are null-terminated
- Some tricks to enter non-printable characters
- Some issues concerning little Endian and big Endian memory structures
- Attacker might not quite know where shellcode will be put; use **landing pad** (sledge) of NOP operations

# Where to put the Shellcode?

- When attack starts, the system isn't corrupted yet
- **argv[]-method**: put shellcode on the stack as part of the malign input
  - Attacker has to guess distance between return address and address of the argument containing the shellcode
  - Detailed examples in “Smashing the Stack for Fun and Profit”
- **return-to-libc-method**: jump to **eval** library function that will execute commands provided as user input
- **Return-oriented programming (ROP)**: jump to code segment in some existing executable

# Detour – Relative Addressing

- Note: attacker does not have to guess the absolute address of the shellcode, only its **relative address**
- Relative addressing: compute  
$$actual\ address = base + offset \times size$$
- When addressing an element in an array, **base** is the start address of the array, **offset** is the array index (multiplied by the size of array elements)
- Used by compilers as it is usually not known at compile time where code will be loaded in memory



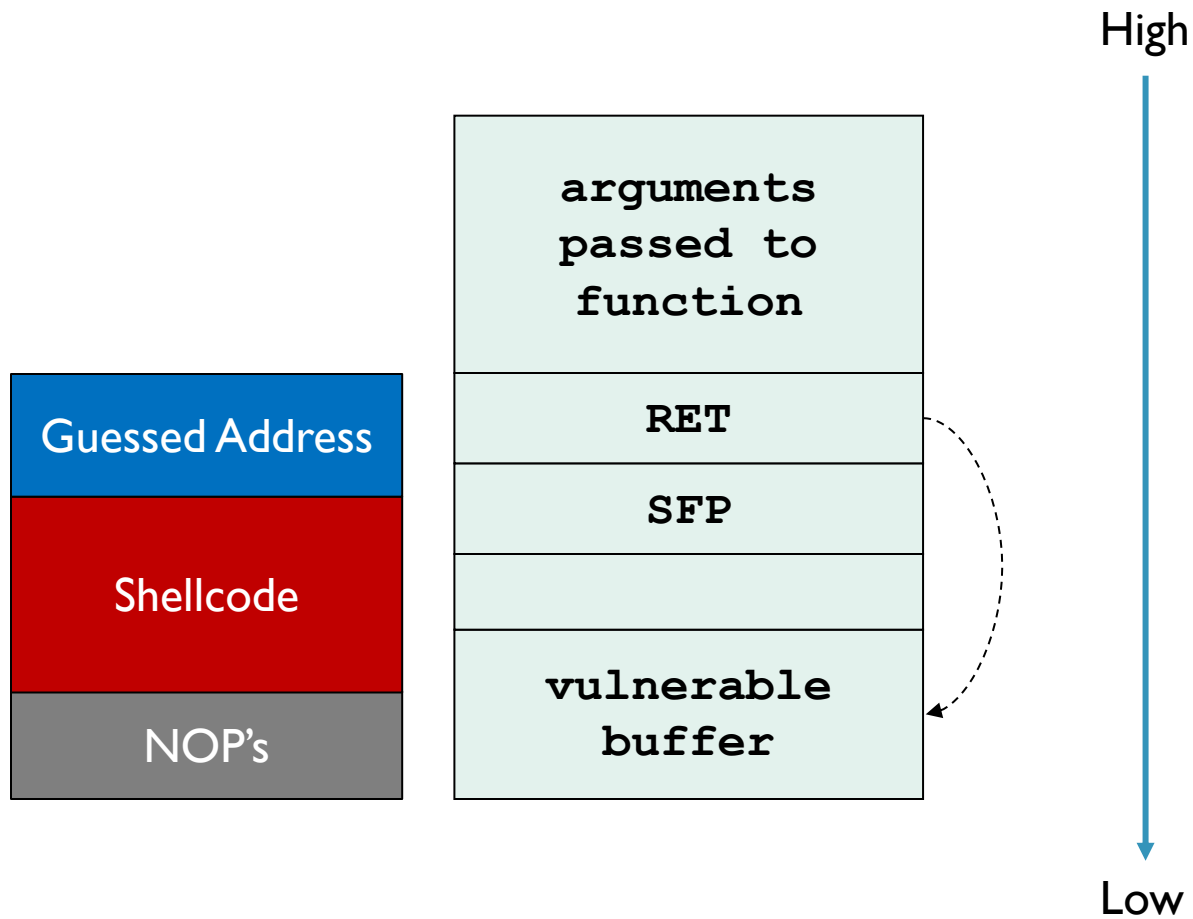
# Putting it all together

- Now assume we (the attacker) have a program that does a string copy from an input buffer that we control, how do we use this shellcode?
  - We can control where the program will return to by overwriting the return address, but we don't know where the shellcode will sit in memory, so we have to guess
  - We put our guesses at the end since it's the end of the buffer that will overwrite the return address
  - We pad the front of the shell code with NOP's (landing sled). This way if we jump into any address in that region, we will eventually execute the shellcode





# Putting it all together



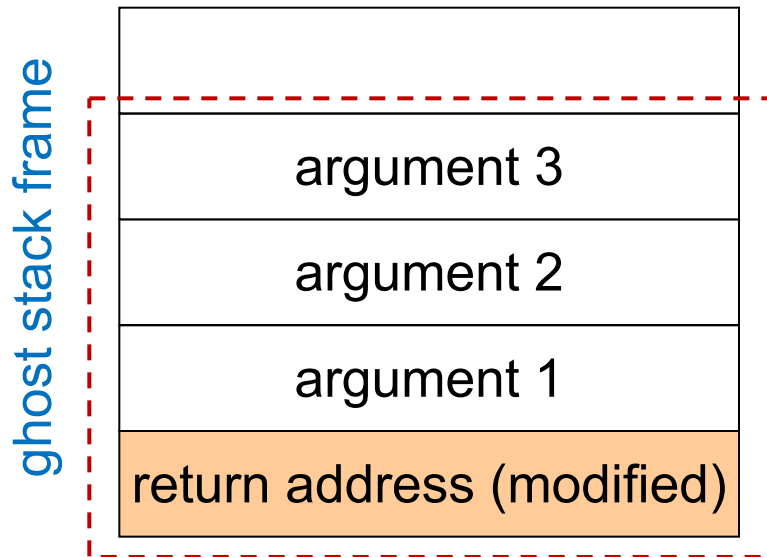
# Internet Worm of 1988 (Morris)

- Sent special 536 byte message to overflow a buffer in the **fingerd** daemon and overwrite the **system stack**:

<code>pushl</code>	<code>\$68732f</code>	push <code>‘/sh &lt;NUL&gt;’</code>
<code>pushl</code>	<code>\$6e69622f</code>	push <code>‘/bin’</code>
<code>movl</code>	<code>sp, r10</code>	save address of start of string
<code>pushl</code>	<code>\$0</code>	push 0 (arg <b>3</b> to <code>execve</code> )
<code>pushl</code>	<code>\$0</code>	push 0 (arg <b>2</b> to <code>execve</code> )
<code>pushl</code>	<code>r10</code>	push string addr (arg <b>1</b> to <code>execve</code> )
<code>pushl</code>	<code>\$3</code>	push argument count
<code>movl</code>	<code>sp, ap</code>	set argument pointer
<code>chmk</code>	<code>\$3b</code>	do “ <code>execve</code> ” kernel call

- On return to **main**
  - `execve (“/bin/sh”, 0, 0)` is executed, opening a connection to a remote shell via TCP
- Video: [https://youtu.be/xdnwR\\_T-qx0?t=262](https://youtu.be/xdnwR_T-qx0?t=262)

# return-to-libc



- Return address changed so that execution returns to a library function
  - Library function expects its arguments on the stack
- Attacker puts arguments to library function in a ghost stack frame
- Defence: library functions that take arguments only from CPU registers

# Defending Against Buffer Overflow

Closing the loopholes

# Countermeasures

- Can be classified in terms of their “locations”
  - Programming language
  - Libraries
  - Compiler
  - Hardware
- Can be classified in terms of security strategies
  - Prevention: stack overflows cannot occur
  - Detection: stop execution when a stack overflow is detected
  - Mitigation: consequences of a stack overflow are made less serious

# Countermeasures

- We will now systematically analyze how a buffer overflow attack can be stopped
- Each defence can help if the previous ones had not been effective:
  1. Check input arguments for variables / arrays so that input sizes fit the memory allocated
  2. Protect return address
  3. Do not allow executable inputs
  4. Make it difficult to guess location where shellcode / system libraries will be placed

# Prevention – Programmer

- When programming in C or C++, the first line of defence against buffer overflow is the programmer
- C is infamous for its **unsafe string handling functions**: **strcpy**, **sprintf**, **gets**, ...

- Example: **strcpy**

```
char *strcpy( char *strDest,  
              const char *strSource );
```

- Throws exception if source or destination buffer are null
- **Undefined if strings are not null-terminated**
- No check whether destination buffer is large enough
- <http://www.cplusplus.com/reference/cstring/strcpy/>

# ‘Safe’ Functions

- Replace unsafe string functions by functions where the number of bytes/characters to be handled are specified: `strncpy`, `_snprintf`, `fgets`, ...

- Example: `strncpy`

```
char *strncpy(char *strDest, const char  
*strSource, size_t count );
```

- You still have to get the byte count right

- <http://www.cplusplus.com/reference/cstring/strncpy/>



# Using 'Safe' Functions

- You have to **get your arithmetic right**
  - Problem will be discussed in next lecture
- You have to know the correct **maximal size** of your data structures
  - Straightforward for data which are used within a single function
  - May be tricky for data structures shared between programs
  - If you underestimate the required length of the buffer your code may crash

## Guards – Example

```
bool HandleInput_Strncpy1( const char* input)
{
    char buf [80];
    if (input == NULL) {
        assert(false);
        return false; }

    strncpy (buf, input, sizeof(buf) - 1);
    buf[sizeof(buf) - 1] = '\0';
    // more processing ...
    return true;
}
```

Problem: if input is too long it will be truncated; this might cause problems elsewhere

## Guards – Example

```
bool HandleInput_Strncpy2( const char* input)
{
    char buf [80]
    if (input == NULL) {
        assert(false);
        return false; }

    buf[sizeof(buf) - 1] = '\\0';

    strncpy (buf, input, sizeof(buf));

    if(buf[sizeof(buf) - 1] != '\\0';
        { return false;} //Overflow!


    // more processing ...
    return true;
}
```

# strsafe.h

- **strsafe** header for un-defining unsafe functions, e.g.  
`#define strcpy unsafe_strcpy`
  - Errors raised at compile time when code contains this unsafe function
- String handling library that is true to the abstraction
  - No overflow of destination buffer
  - Buffers guaranteed to be null-terminated
- Question: why can't we get rid of all buffer overflows by replacing **strcpy**, **sprintf**, **gets** and the like?

# Type Safe Languages

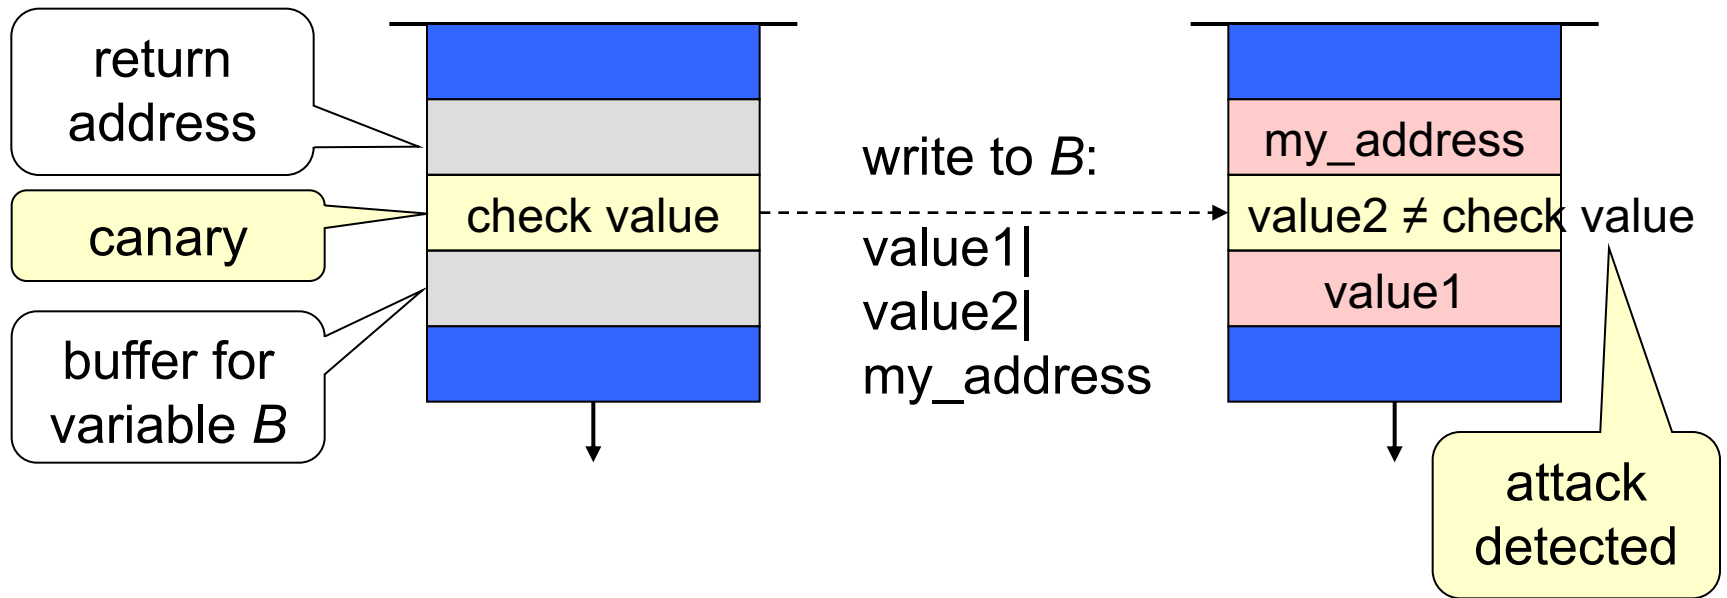
- Defence: use a programming language where buffer overflows cannot happen “by design”
- Safety guaranteed by static checks and by runtime checks
  - Automatic array bounds checking
  - Automatic garbage collection
  - Programmers can't make mistakes when managing memory
  - Programmers cannot optimize memory usage
  - Programmers cannot themselves take care of multiple copies of sensitive data (e.g. a password)
- Examples: Java, Ada, C#, Perl, Python, etc.
- More on type safety later in the course



# Detection – Compiler

- Detect attempts at overwriting return address
- Compiler places a **check value** ('canary') in the memory location just below the return address
  - The term canary is borrowed from coal mining
- Before returning, check that canary hasn't changed
- Stackguard: random canaries
  - Alternatives: null canary, terminator canary
- **Source code has to be recompiled to insert placing and checking of the canary**

# Canaries



# Remark on Check Values

- (Integrity) check values are a generic defence
- Security requirements:
  - Check values must not be predictable
  - Reference values must be integrity protected
- Two options for integrity protection:
  - Write to protected memory
  - Calculate checksums using a secret cryptographic key (magic value); only the key has to be protected; no need to write the key when compiling a program



# Non-Executable Memory

- Support at hardware level for marking memory as non-executable
  - AMD: NX (no execute) bit since Athlon 64
  - ARM: XN (eXecute never)
  - Intel: XD (execute disable bit, EDB)

# Data Execution Prevention

- **Data Execution Prevention** (DEP): mark memory location a process has written to as non-executable
- **W $\oplus$ X** (W $\wedge$ X if you are a C programmer) protection
  - Memory can be writeable or executable, but not both
  - Shellcode placed by attacker will not be executed
- This approach cannot be applied to components creating executable code, e.g. compilers
  - Exploited by **JIT spraying attacks**

# Randomizing Memory Allocation

- Stack buffer overflow attacks have to know the position of a target buffer in relation to return address, and the approximate location of attack code
- Random changes to memory allocation can reduce impact of buffer overflow attacks
- Example: Address Space Layout Randomization (ASLR) in Linux, Windows, ...
  - Randomizes memory location, e.g. of stack, heap, libraries
  - Defence against [return-to-libc](#) attacks

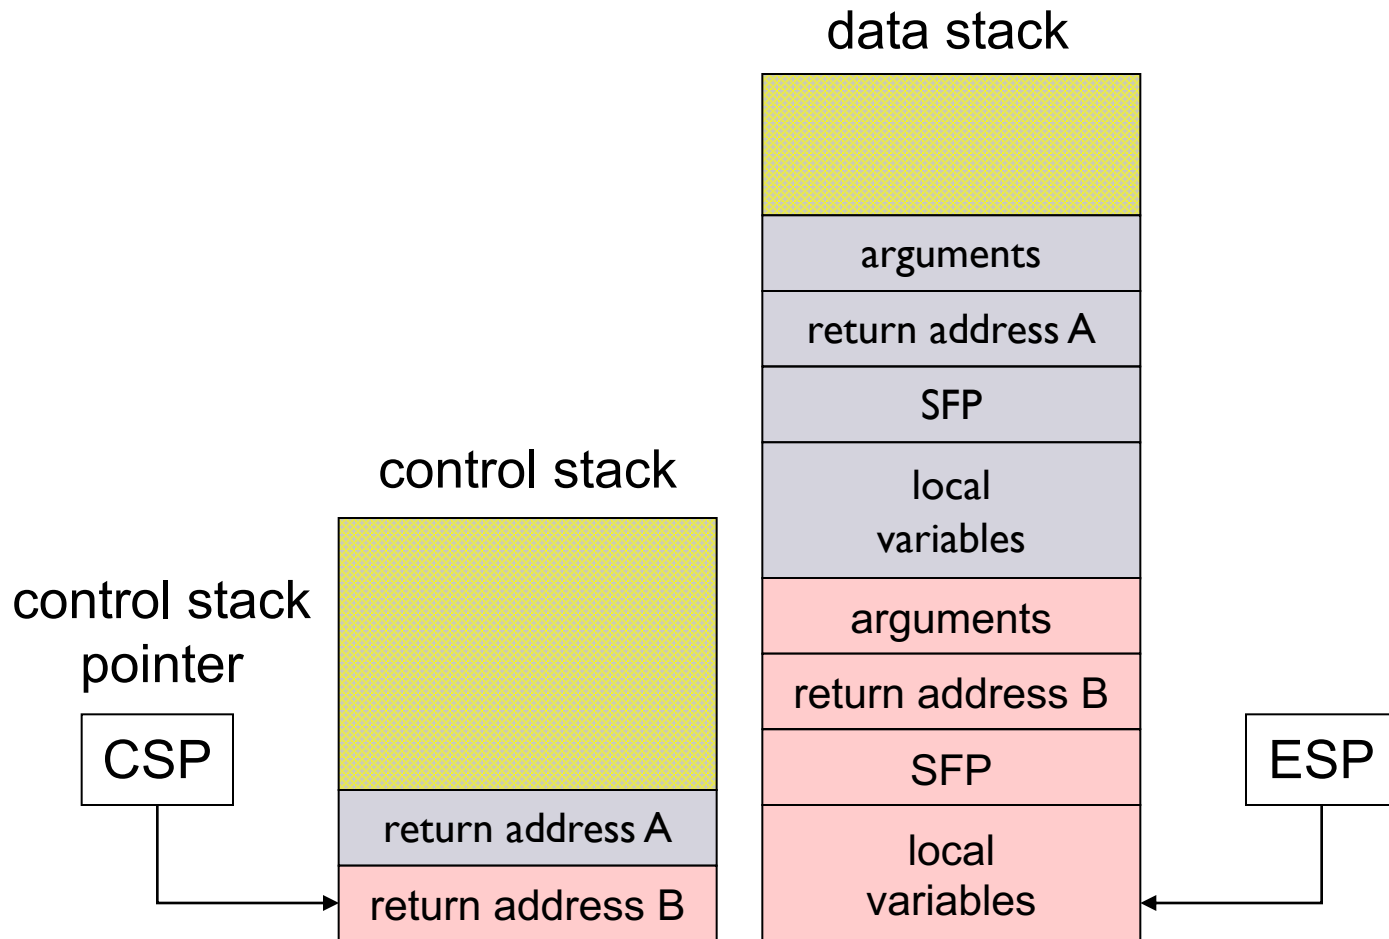
# Lesson

Randomness is  
important for  
security

# Split User Data from Control Data

- Stack-based buffer overflow attacks overwrite return address
- Such attacks would be blocked if the return address is not taken from the stack but from a location the user (input) cannot touch
- Abstraction: separation of control and user data
- Control stack: memory area separate from data stack
- Implementations in hardware and software have been proposed in the research literature

# Split Control and Data Stack



# Return-Oriented Programming

A smart way around Data Execution Prevention

# Bypassing DEP

- Shellcode need not be inserted; the attacker may use existing executables, e.g. from system libraries, in **return-to-libc** attacks
- Attacker is limited to given executables
- Can one use executables in unintended ways, e.g. by selectively using instruction blocks from executables?



# Return-Oriented Programming

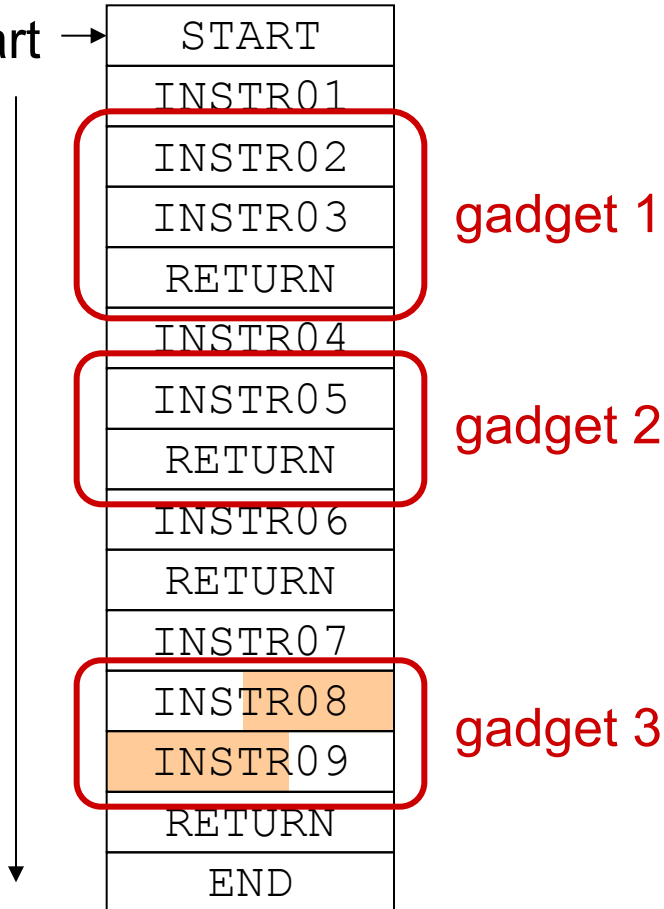
- Attacker uses short code segments in executables, so-called **gadgets**, that end with a return command
  - When **misaligned memory access** is permitted the attacker may find byte sequences starting in the middle of a word that constitute valid machine instructions
  - Gadgets serve as building blocks for writing shellcode
  - If enough gadgets are found in the code base, shellcode with **arbitrary functionality** can be built (→ **weird machine**)
- Stack buffer overflow attack puts gadget addresses on stack; after returning from one gadget, the attack jumps to the next (call stack used as a **trampoline**)

# Return-to-libc and ROP

Return-to-libc      executable

jump to start →

execute  
entire  
function




ROP: on the stack

[gadget 2]  
[gadget 1]  
[gadget 3]  
[gadget 1]



ROP shellcode executed



# Return-Oriented Programming

- ROP more powerful than return-to-libc attacks
- Initially assumed to be only an issue in architectures with variable length instructions
- Later shown that such vulnerabilities can also exist in RISC architectures (fixed length instructions)
- Real exploits have been documented
  - AVC Advantage voting machine
  - Adobe Reader 9.3 with DEP enabled
- Similar attack patterns exist where gadgets need not end with a return (**jump-oriented programming**)
- Tutorial Video: <https://www.youtube.com/watch?v=zaQVNM3or7k>

# Resources

- Hovav Shacham. [The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls](#), Proceedings of CCS 2007, pages 552–561
  - <http://cseweb.ucsd.edu/~hovav/papers/s07.html>
- Erik Buchanan, Ryan Roemer, Hovav Shacham, Stefan Savage. [When Good Instructions Go Bad: Generalizing Return-Oriented Programming to RISC](#), Proceedings of CCS 2008, pages 27–38
- Stephen Checkoway, Ariel J. Feldman, Brian Kantor, J. Alex Halderman, Edward W. Felten, Hovav Shacham. [Can DREs Provide Long-Lasting Security?](#) USENIX 2009
- Jduck. The Latest Adobe Exploit and Session Upgrading, 2010
  - <https://blog.rapid7.com/2010/03/18/the-latest-adobe-exploit-and-session-upgrading/>

# Comments and Further Countermeasures

# Growing the Stack

- Comment in Paul A. Karger & Roger R. Schell: [Thirty Years Later: Lessons from the Multics Security Evaluation](#)
- Stack buffer overflows occur when stack grows **downwards** but data are entered **upwards** into a buffer
- Such problems would not occur if the stack starts at the bottom of memory and grows upwards

# Targets

- Overwriting the return address is a common form of a buffer overflow attack
- If return address cannot be reached, alternative targets include:
  - overwrite a function pointer variable on the stack
  - overwrite security-critical variable value on stack
  - overwrite previous frame pointer
  - **overwrite arguments with ghost stack frames**

# No Silver Bullet

- None of the countermeasures are perfect
- The earlier buffer overflows are addressed in the design process the better
- Systematic work on removing security relevant buffer overflows started 15-20 years ago
- Prediction (Jon Pincus, 2004): buffer overflows will disappear as an issue in the next 1-2 years, but there will be other software security issues
  - This prediction has by and large held up; major problems today are cross-site scripting, SQL injection, ...



# Conclusions

- Aleph One's paper on buffer overflow attacks focused general attention on software security
- Main design problems exploited:
  - Unsafe write operations
  - Stack frame contains both user and control data
- Significant progress has been made since, both in research and in the defences deployed in the wild
  - Stack buffer overflow attacks hardly work in today's operating systems
- Attackers had to become much more sophisticated
  - More on this in the coming lectures

# Tutorial

# ASLR & Jump-Oriented Programming

- Collect information on ASLR in Linux and in Windows
  - What is being randomized? How random is randomization?
- **Jump-oriented programming** creates its own trampoline for linking gadgets instead of using the call stack
  - **Jump-Oriented Programming: A New Class of Code-Reuse Attack**
- Read up on ASLR and JOP; you should be able to explain the fundamentals of these attack methods