



# Lecture 12: Type Safety

*presented by*

**Li Yi**

*Assistant Professor*  
SCSE

N4-02b-64

[yi\\_li@ntu.edu.sg](mailto:yi_li@ntu.edu.sg)

# Motivation

- Taint analysis tools (and other security analysis tools) are applied after vulnerable code has been written
- Why not write bug-free code in the first place?
- Sentiment after *Smashing the stack for fun and profit*: blame yourself for writing code in C/C++; use a type safe language like Java and your problems go away
- This lecture will examine what is in the claim that type safe languages will guarantee secure software

# Type Safety

- Guarantees the absence of **untrapped errors**
- The precise meaning of type safety thus depends on the definition of error
  - A language is unsafe if you are not told when an error has occurred during execution
- Luca Cardelli: practitioners who invented type safety often meant just “**memory integrity**”, while theoreticians always meant “**execution integrity**”, and it is the latter that seems more relevant now
- Languages do not have to be typed to be safe: LISP

# Memory and Type Safety

**Memory safety:** each dereference can only access “intended target”

- **Spatial** access errors (e.g., out-of-bounds array access)
  - E.g., if `i` is negative or is too large
- **Temporal** access errors (e.g., stale pointer dereference)
  - E.g., if `a` had been freed prior to `* (p+i)`

```
p = &a
```

```
// p's valid target is a
```

```
* (p+i)
```

```
// valid only if it accesses a
```

# Memory and Type Safety

- **Type safety**: operations on the object are always **compatible** with the object's type
  - Each object is ascribed a **type** (`int`, pointer to `int`, pointer to function)
  - Type safe programs do not “go wrong” at runtime
- Type safety is **stronger** than **memory safety**

```
int (*cmp)(char*,char*);  
int *p = (int*)malloc(sizeof(int));  
*p = 1;  
// memory safe but not type safe  
cmp = (int (*)(char*,char*))p;  
cmp("hello", "bye"); // crash!
```

# Type Safe Languages

- Useful for improving quality of software
- Safety guaranteed by **static checks** and by **runtime checks**
- Language design: trade-off between efficiency and flexibility
  - If too general, incurs a high runtime overhead to enforce
  - If too restrictive, limits expressiveness and utility of language
- Examples: Java, Ada, C#, Rust, Go, Perl, Python, etc.
  - **Dynamically typed languages**, like Python, which do not require declarations that identify types, can be viewed as **type safe** as well
  - Each operation on a dynamic object is permitted, but may be **unimplemented**
  - In this case, it throws an **exception** (unfortunate, but still well-defined)

# Type Safety – Limitations

- Typical enforcement of type safety is **expensive**
  - **Garbage collection** avoids temporal violations
  - **Bounds** and **null-pointer checks** avoid spatial violations
  - Hiding representation (**abstract types**) may inhibit optimization
    - Many C-style casts, pointer arithmetic, & operator, are not allowed
- Type system may not deal with all security problems
  - E.g., **race conditions** are possible in Java code
  - **Erasure of sensitive data**: consider a login program accepting a user password; how to make sure that you know where copies of the password are held and that all copies are erased once the login process has finished?

# Execution Integrity



# Execution Integrity

- Execution integrity is at the core of software security
- Separate data from control information; we have seen examples before
- System-environment integrity: certain aspects of the processor, e.g., condition flags, cannot be changed
- Memory access: module can only access (read/write/execute) certain memory regions; e.g., module cannot overwrite its own code
- Interface: calls from a module only to predefined routines, returns only at defined entry points


# Control-Flow Integrity

- **Control-flow integrity**: execution must follow a **control-flow graph** (CFG) determined in advance
  - E.g., function must return to its caller
- CFG can be defined by analysis, by execution profiling, or by an explicit security policy, e.g., written as a security automaton
- Machine-code rewriting can be used to insert **runtime checks** that dynamically ensure that control flow adheres to a given CFG

# Call Graph

```
sort2(int a[], int b[], int len){  
    sort(a, len, lt);  
    sort(b, len, gt);  
}
```

```
bool lt(int x, int y){  
    return x < y;  
}  
bool gt(int x, int y){  
    return x > y;  
}
```




Which functions call other functions

# Control Flow Graph

```
sort2(int a[], int b[], int len){  
    sort(a, len, lt);  
    sort(b, len, gt);  
}
```

```
bool lt(int x, int y){  
    return x < y;  
}  
bool gt(int x, int y){  
    return x > y;  
}
```



Break into **basic blocks**  
Distinguish **calls** from **return**


# Control-Flow Integrity

- Observation: **direct calls** need not be monitored
  - Assuming the code is immutable, the target address cannot be changed
- Therefore, monitor only **indirect calls**
  - `jmp, call, ret` via registers
- **Forward-edge CFI** for indirect calls and jump instructions
- **Backward-edge CFI** for return instructions
  - **Shadow stack** for return instructions: separate isolated stack for storing return addresses
  - Enforces natural semantics of call/return: return is supposed to transfer the control flow to the next instruction after the call instruction that invoked the current function

# Control Flow Graph

```
sort2(int a[], int b[], int len){  
    sort(a, len, lt);  
    sort(b, len, gt);  
}
```

```
bool lt(int x, int y){  
    return x < y;  
}  
bool gt(int x, int y){  
    return x > y;  
}
```




Indirect transfer (call via register, or ret)

# Control-Flow Integrity

- During execution, each machine instruction that transfers control must target a valid **destination**, as determined by the given CFG
- For instructions that target a **constant destination**, this can be done by a **static check**
- **Computed control-flow transfer needs dynamic check**
- At each destination, insert an ID (label) identifying an **equivalence class of destinations**
  - Design challenge: how to define equivalence?
- Before each source, insert an ID-check that validates that the destination has the right ID

# Control-Flow Integrity




Equivalence class I: use the same label at all targets



# Control-Flow Integrity

- Design question: how many different IDs?
  - Unique IDs?
  - IDs chosen from a given set and reused?
- Attacker who does not know the IDs has to be lucky when deviating from the CFG to reach a destination with a valid ID
- Attacker who knows the IDs may try to construct an attack that avoids detection
  - The more IDs are reused, the more options for attacks
- CFI enforcement is not a panacea: exploits within the bounds of the allowed CFG are not prevented

# Control-Flow Integrity



Equivalence class 2:

- Return sites from calls to `sort` must share a label (L)
- Call targets `gt` and `lt` must share a label (M)
- Remaining label unconstrained (N)

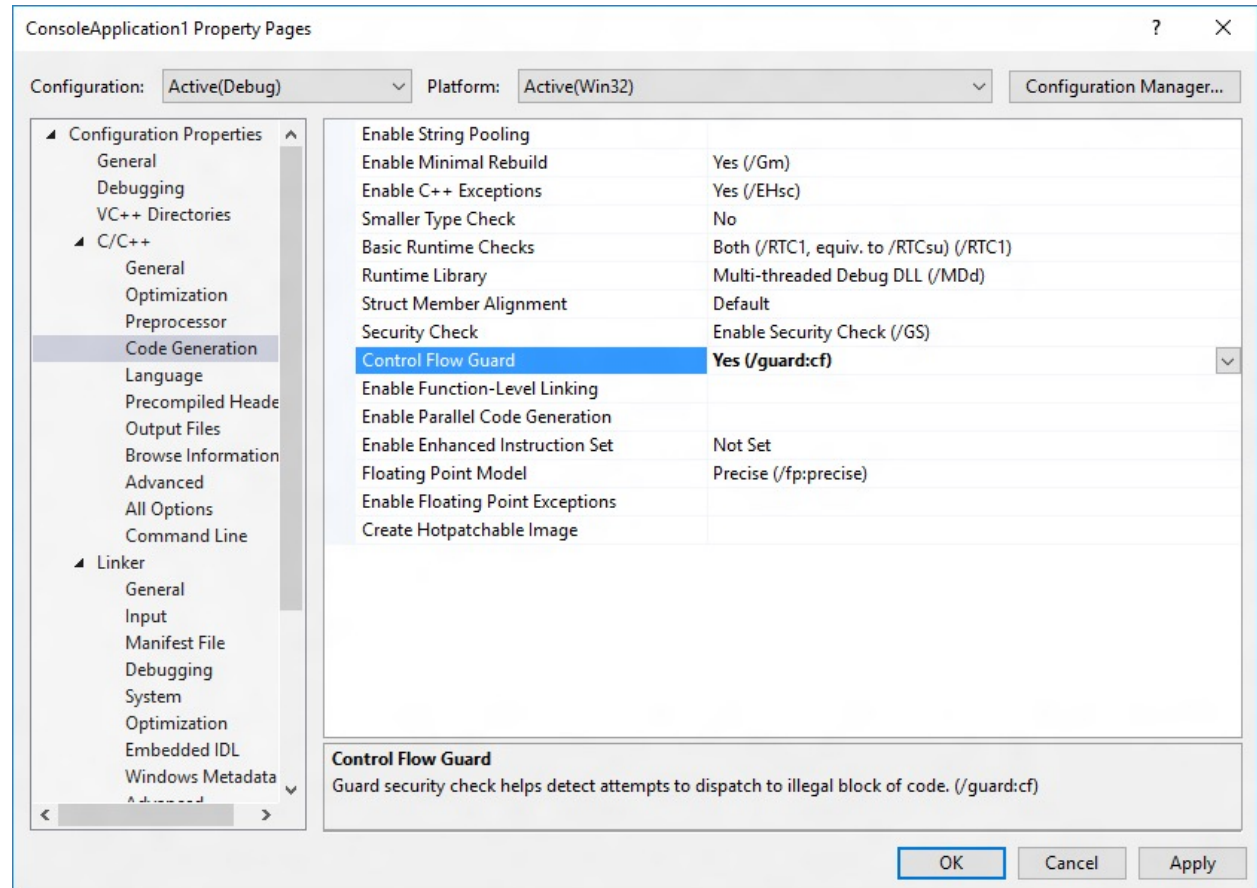
Still permits call from site A to return to site B

# Backward-Edge CFI & Exceptions

- Note that there are benign cases in which the call/return semantic is broken
- E.g., if the currently executed function generates an **exception**, it need not be the calling function that catches the exception
- Another function in the call hierarchy or a default exception handler can be responsible
  - In such a case, the current function **returns to the function in the call stack that implements the exception handler**

# Microsoft Control Flow Guard (CFG)

- CFG first released for Windows 8.1 Update 3 in 2014
- Compiler, operating system, user mode library, and kernel mode module are cooperating



# Type Confusion

Escaping out of the type system

# Type Confusion

- **Static type checking:** check all possible executions of a program to see whether a type violation could occur
- **Dynamic type checking:** check the class tag when access is requested
- **Type confusion attack:** exploit vulnerability in type checking procedure to create two pointers to the same object with incompatible type tags
- Attack code passes static type checking, but actual access is to object of the wrong type

# Type Checking in Java

- Type safety (memory safety): programs cannot access memory in inappropriate ways
- Java: each object has a class; only certain operations are allowed to manipulate objects of that class
- Every object in memory is labelled with a **class tag**
- When a Java program has a reference to an object, it has internally a **pointer** to the memory address storing the object
- The pointer can be thought of as **tagged with a type** that says what kind of object the pointer is pointing to

# Type Confusion

- Assume the attacker manages to let two pointers point to the same location

```
class T {  
    SecurityManager x;  
}  
  
class U {  
    MyObject x;  
}
```


class definitions

```
T t = the pointer tagged T;  
U u = the pointer tagged U;  
t.x = System.getSecurity() ;  
MyObject m = u.x;
```

malign applet



# Type Confusion



# Type Confusion

- The `SecurityManager` field can now also be manipulated from `MyObject`
- We sketch a type confusion attack in Netscape Navigator 3.0β5 (discovered by Drew Dean), fixed in version 3.0β6
- Source: Gary McGraw & Edward W. Felten: [Java Security](#), John Wiley & Sons, 1997

# Netscape Vulnerability

- Java allows a program that uses type `T` also to use type *array of* `T`
- Array types defined by the virtual machine for internal use; array names begin with the character `[`
  - Array of float: `[F`
  - Array of array of double: `[[D`
  - Array of objects: `[Ljava/lang/Object`
  - A programmer defined *classname* is not allowed to start with this character; hence, there should be no danger of conflict
- However, a Java *bytecode file* could declare its own name to be a special array types name, thus redefining one of Java's array types
  - Attempting to load such a class would generate an error, but the Java VM would install the name in its internal table anyway
  - Access to the layer below: from Java code down to bytecode

# SUN Security Bulletin #00218

- Relates to the byte code verifier of the Java Runtime Environment (JRE) [March 18, 2002]
- Cause: flaw in the checking of type casts
- Applets could exploit this flaw to increase their privileges and get out of the sandbox
- If the user loads an applet designed to exploit this weakness, the attacker can execute arbitrary code on the local machine with the user's privileges

# Summary

- Type confusion attacks are an example for attacks below the level that performs access control
- Type confusion attacks often result in complete system penetration
- It does not seem to be that easy to spot such vulnerabilities; relatively few are reported
- It is not easy either to prove that a system is type safe, or to achieve type safety first time round when a new system is being designed

# Breaking Type Safety from Hardware

\* this will not be tested \*

# Breaking Type-Safety with a Torch

- Type-confusion attack using **random memory error**
- Create a data structure where a random memory error is likely to change a reference so that it points to an object of the wrong type
- Then exploit the fact that pointers of different types point to the same object in memory



# The Attack: Classes


```
class A {  
    A a1;  
    A a2;  
    B b;  
    A a4;  
    A a5;  
    int i;  
    A a7;  
};
```

```
class B {  
    A a1;  
    A a2;  
    A a3;  
    A a4;  
    A a5;  
    A a6;  
    A a7;  
};
```

Size of the above classes `A` and `B` is a power of 2




# Attack: Flipping a Bit (Case I)



bit flip in position  $k$   
when  $2^k$  is larger  
than the object size

Precondition: memory filled  
mainly with type B objects

# Attack: Flipping a Bit (Case 2)




bit flip in position  $k$   
when  $2^k$  is smaller  
than the object  
size

$p$  now likely to point to  
an object of type  $A$ , but  
start address of this  
object is **misaligned**

# Misaligned Pointer References

- Assume a bit is flipped in a pointer variable  $p$  of class  $A$  at address  $x$
- $offset$ : distance between base of object and beginning of  $b$  field
- Dereferencing the  $b$  field of pointer  $p$  into a pointer  $s$  (of type  $B$ ) fetches value of  $s$  from address  $x + offset$
- When the  $k$ -th bit of  $p$  has flipped, the fetch would be from address  $(x \oplus 2^k) + offset$
- Likely that  $s$  now points to an object of type  $A$

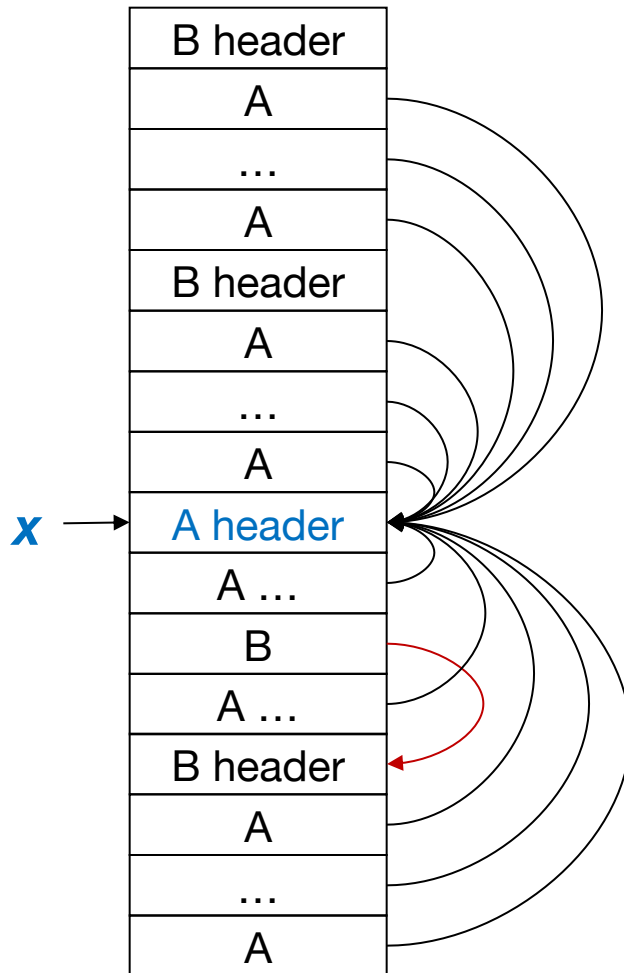
# The Attack: Pointer References



# Attack – Memory Layout

- Fill heap with lots of objects of type *B* and one object of type *A* at location *x*
- All *A* fields in the objects of type *B* point to this object (contain address *x*)
- The *B* field in the *A* object points to one of the objects of type *B*
- Induce – or wait for – a **random bit error**
- Keep checking the objects in memory until a position with a bit error has been found

# Type-Confusion Attack: Data Structure



- Suppose a bit flips in a  $B$  object
- Original position contained an  $A$  object pointing to address  $x$
- Modified pointer  $x \oplus 2^k$  is likely to point to an  $A$  field
- When dereferencing its  $b$  field one is likely to hit a field of type  $A$

# Type Safe Attack Code

```
A p;  
B q;  
int offset = 6*4;  
void write(int target,  
           int value)  
{  
    p.i = target - offset;  
    q.a6.i = value;  
}
```

let *p* and *q* point to  
same memory  
location *x*

writes *value* into  
location *target*

```
class A {  
    A a1;  
    A a2;  
    B b;  
    A a4;  
    A a5;  
    int i;  
    A a7;  
};
```

```
class B {  
    A a1;  
    A a2;  
    A a3;  
    A a4;  
    A a5;  
    A a6;  
    A a7;  
};
```

writes *target - offset* into  
field *q.a6*

# Mounting the Attack

- How to induce a memory error?
- If you have physical access to the processor, shine a light bulb at it
- Demonstrated at the 2003 IEEE Symposium on Security & Privacy on IBM and Sun's JVM
  - Sudhakar Govindavajhala and Andrew W. Appel: [Using Memory Errors to Attack a Virtual Machine](#), Proceedings of the 2003 IEEE Symposium on Security and Privacy
- Can overwrite arbitrary memory locations (as in the double free vulnerability)



# Discussion

- Assessment of attack
  - To use a light bulb, you need physical access
  - You need to be able to fill a large portion of memory to increase the chance that a random memory flip has the desired effect
  - Summary: nice demo, but not particularly serious
- Countermeasures
  - Error-correcting memory would repair the bit flip
  - Error-detecting memory would flag the bit flip
  - Cost: a few bits of each memory word

# Rowhammer

\* this will not be tested \*


# DRAM – Coupling Effects

- Can we do better than relying on random bit flips
- Can we launch attacks remotely?
- DRAM – Dynamic Access Memory
- Electrical couplings in DRAM is not new
  - May cause disturbances during operation
  - Major hurdle in DRAM scaling
- Attack vector at physical layer: change in one memory cell can flip neighbouring bits

# Rowhammer Attack

- It depends on the given DRAM chip which bits in which rows can be flipped by hammering a neighbour
  - These effects are repeatable!
- Arrange memory so that
  - Vulnerable row (victim) holds a data structure where a bit flip in the vulnerable position is of use to the attacker
  - Attacker has write access to a row (aggressor) next to victim
- Write to aggressor row repeatedly (“hammering”)
- Induce an error in the victim row; the change in the data structure creates an exploitable vulnerability

# Rowhammer Attack



- **Rowhammer:** purple row hammers on yellow rows
- **Double-sided rowhammer:** yellow rows hammer on purple row
  - More effective at physical level
  - More of a challenge to get write access to the two rows next to your target

# Rowhammer – Defences

- Change layout of DRAM
  - Addresses the issue at the level of physical structures
- Refresh memory cells more frequently
  - Addresses the issue at the “physical operational” level; increases power consumption
- Apply error correcting codes to memory words
  - Adds a defence in memory at a logical level
- Access counters to detect hammering
- After closing a row, open neighbouring rows with low probability (**randomize allocation of memory rows**)
  - Addresses the issue at a “logical operational” level

# Drammer

- Rowhammer attack on Android devices
- Method: modify page table entries (PTEs)
  - PTEs map virtual memory to physical memory
  - Layer of indirection!
- Preparation (easily said, more difficult to do):
  - Find out which memory rows are susceptible to hammering [on the given device](#)
  - Arrange memory in a way that a PTE for a file where you have R/W access occupies a vulnerable memory row
  - Fill physical memory with PTEs for files with R/W access

# Drammer – Execute Attack

- Hammer the aggressor row to induce a bit flip in victim row so that PTE points to another PTE in its own page table
- You thus get write access to a PTE and can scan kernel memory by writing new PTEs
- Once you find your credentials data structure you can overwrite it with zeros and become root



# Drammer – Ressources

- <https://www.vusec.net/projects/drammer/>
- <http://thehackernews.com/2016/10/root-android-phone-exploit.html>
- <https://vvdveen.com/publications/drammer.pdf>
- <https://www.blackhat.com/docs/us-15/materials/us-15-Seaborn-Exploiting-The-DRAM-Rowhammer-Bug-To-Gain-Kernel-Privileges.pdf>

# Conclusion

# Summary

- Type safety helps programmers avoiding certain classes of vulnerabilities
  - Programmers still have to catch and handle exceptions
- Memory integrity and control-flow integrity have raised the bar for attackers over the past 15 years
- Access to the layer below can undermine type safety
  - Direct manipulation of Java bytecode
  - Direct manipulation of memory
- You then need further defences at the layer below to preserve the security guarantees from a higher layer