

A Very Short Introduction to CTF - Part I

Wu Xiuheng

February 22, 2022

Overview

Some Basics

Reverse Engineering

Forensics

Appendix

Overview

- The purpose is to get you familiar with Jeopardy! style of CTF competitions.
- There are lots of links (in blue color) to online resources within the slides.
- Although lots of tools and challenges are mentioned, you do not need to understand or solve all of them.
- Think and try before rushing to provided hint or solutions.

The Name of the Game

Capture the Flag or CTF comes from an outdoor game where two teams compete for capturing the opponent's flag and keeping their own flag safe.



(Capture the Flag, Fahne im Spiel – Ronnie Berzins – CC BY-SA 3.0)

When it comes to the context of cyber security, two forms exist:

- Attack-defense style
- *Jeopardy!* style

When it comes to the context of cyber security, two forms exist:

- Attack-defense style
- *Jeopardy!* style

Discover flags (usually a text string) by using techniques including:

- reverse-engineering,
- forensics,
- web exploitation,
- cryptanalysis. . .

When it comes to the context of cyber security, two forms exist:

- Attack-defense style
- *Jeopardy!* style

Discover flags (usually a text string) by using techniques including:

- reverse-engineering,
- forensics,
- web exploitation,
- cryptanalysis. . .

You are expected to

- have some basic knowledge,
- be able to program,
- learn new things quickly

A Simple Example

Puzzles in CTF competitions are usually categorized and arranged by a progression of difficulty.

Warmup challenge of PicoCTF 2018

Cryptography doesn't have to be complicated, have you ever heard of something called rot13? `cvpbPGS{guvf_vf_pe1cgb!}`

A Simple Example

Puzzles in CTF competitions are usually categorized and arranged by a progression of difficulty.

Warmup challenge of PicoCTF 2018

Cryptography doesn't have to be complicated, have you ever heard of something called rot13? `cvpbPGS{guvf_vf_pelcgb!}`

Solution: Wikipedia tells us that [ROT13](#) is a simple substitution cipher similar to Caesar cipher.

Input	ABCDEFGHIJKLMNOPQRSTUVWXYZabcdefghijklmnopqrstuvwxyz
Output	NOPQRSTUVWXYZABCDEFGHIJKLMnopqrstuvwxyzabcdefghijklm

(solution continues on next page)

A Simple Example (cont'd)

Using `tr` (a UNIX utility which can be found on Linux/macOS):

```
tr 'A-Za-z' 'N-ZA-Mn-za-m' <<< "cvpbPGS{guvf_vf_pelcgb!}"  
=> picoCTF{this_is_crypto!}
```

Or python:

```
u = ''.join([chr(x) for x in range(65,91)])  
l = ''.join([chr(x) for x in range(97,123)])  
trpair = str.maketrans(u[13:]+u[:13]+l[13:]+l[:13], u+l)  
print('cvpbPGS{guvf_vf_pelcgb!}'.translate(trpair))
```

A Simple Example (cont'd)

Using `tr` (a UNIX utility which can be found on Linux/macOS):

```
tr 'A-Za-z' 'N-ZA-Mn-za-m' <<< "cvpbPGS{guvf_vf_pelcgb!}"  
=> picoCTF{this_is_crypto!}
```

Or python:

```
u = ''.join([chr(x) for x in range(65,91)])  
l = ''.join([chr(x) for x in range(97,123)])  
trpair = str.maketrans(u[13:]+u[:13]+l[13:]+l[:13], u+l)  
print('cvpbPGS{guvf_vf_pelcgb!}'.translate(trpair))
```

ROT_{13} is self-inverse, i.e., $\text{ROT}_{13}(\text{ROT}_{13}(x)) = x$

Some Basics

Some [binary-to-text encoding](#) are very common in various ctf problems. Be familiar with them so that you can recognize them when they appear in the problem.

- [Base32](#): strings encoded in base32 contains capital alphabets A-Z, numerals 2-7 (variations exists) and = for padding.

Try decode this NFVW433XL5UG6527MJQXGZJTGJPX033SNNZQ==== (answer in white below)

```
echo -n "NFVW433XL5UG6527MJQXGZJTGJPX033SNNZQ====" | base32 -d
```

Some [binary-to-text encoding](#) are very common in various ctf problems. Be familiar with them so that you can recognize them when they appear in the problem.

- [Base32](#): strings encoded in base32 contains capital alphabets A-Z, numerals 2-7 (variations exists) and = for padding.
- [Base64](#): strings encoded in base64 contains alphabets A-Za-z, numerals 0-9, 2 more special characters +/ (variations exists) and = for padding.

Try decode this YmFzZTY0X21zX3Nob3J0ZXI= (answer in white below)

echo -n "YmFzZTY0X21zX3Nob3J0ZXI=" | base64 -d

Some **binary-to-text encoding** are very common in various ctf problems. Be familiar with them so that you can recognize them when they appear in the problem.

- **Base32**: strings encoded in base32 contains capital alphabets A-Z, numerals 2-7 (variations exists) and = for padding.
- **Base64**: strings encoded in base64 contains alphabets A-Za-z, numerals 0-9, 2 more special characters +/ (variations exists) and = for padding.
- **Percent encoding**: common for URI/URL, use hex ASCII values, and UTF-8 for non-ASCII characters. E.g. %20 for space and %2F for slash (/)

E.g. <https://zh.wikipedia.org/wiki/%E6%96%B0%E5%8A%A0%E5%9D%A1> is the Chinese Wikipedia page for Singapore.

Some [binary-to-text encoding](#) are very common in various ctf problems. Be familiar with them so that you can recognize them when they appear in the problem.

- [Base32](#): strings encoded in base32 contains capital alphabets A-Z, numerals 2-7 (variations exists) and = for padding.
- [Base64](#): strings encoded in base64 contains alphabets A-Za-z, numerals 0-9, 2 more special characters +/ (variations exists) and = for padding.
- Percent encoding: common for URI/URL, use hex ASCII values, and UTF-8 for non-ASCII characters. E.g. %20 for space and %2F for slash (/)

Challenge: Some problems use relative uncommon encodings. Try to [solve this](#). (Hackergame2018 of USTC, released under CC BY-NC-SA 4.0 license.)

Hint: come back to this after reading page [23](#) if you cannot solve it now.

Endianness (or *byte ordering*) refers to the order of bytes within a multi-bytes number in memory.

C-type	variable	value	Memory at offset							
			Big Endian				Little Endian			
			+0	+1	+2	+3	+0	+1	+2	+3
int32_t	longVar	0x0a0b0c0d	0a	0b	0c	0d	0d	0c	0b	0a
int16_t	shortVar	0x0c0d	0c	0d			0d	0c		

Endianness

Endianness (or *byte ordering*) refers to the order of bytes within a multi-bytes number in memory.

C-type	variable	value	Memory at offset							
			Big Endian				Little Endian			
			+0	+1	+2	+3	+0	+1	+2	+3
int32_t	longVar	0x0a0b0c0d	0a	0b	0c	0d	0d	0c	0b	0a
int16_t	shortVar	0x0c0d	0c	0d			0d	0c		

Little-Endian: the **Least significant** byte goes into the **Lowest-address**.

Binaries of today's PC are mostly little-endian and network packets are big-endian.

Endianness (cont'd)

```
// simple.c
int main(void) {
    long a = 1234567890;
    return 0;
}
```

```
python3 -c "print(hex(1234567890))"
```

```
=> 0x499602d2
```

compile and run with debugger

```
gcc -g simple.c -o simple && gdb simple
```

```
(gdb) break 3
```

```
Breakpoint 1 at 0x1131: file simple.c, line 3.
```

```
(gdb) run
```

```
Breakpoint 1, main () at simple.c:3
```

```
(gdb) p/d a
```

```
$1 = 1234567890
```

```
(gdb) p &a
```

```
$2 = (long *) 0x7fffffffdba8
```

```
(gdb) x/4xb &a
```

```
0x7fffffffdba8: 0xd2    0x02    0x96    0x49
```

GDB Manual:
[Examine](#)
[Memory](#)

Example: a BMP image file beginning with 42 4D EA 9B 22 00 (first 6 bytes)

- the first 2 bytes (42 4d: BM) are file signature (magic bytes)
- bytes 3-6 (EA 9B 22 00) represents file size.
 - the file size is actually 00229BEA (2268138) bytes.

`hexdump` is a utility for “display file contents in hexadecimal, decimal, octal, or ascii”. When using without any options, it will display bytes as they are 2-bytes (16-bits) words in little-endian.

While `xxd` gives a human-friendly order by default.

```
echo -n "abcd" > tmp
# hexdump tmp
00000000 6261 6463
# hexdump -C tmp
00000000 61 62 63 64 |abcd|
# xxd tmp
00000000: 6162 6364  abcd
```

Reverse Engineering

Two syntax exist for x86 assembly.

[web resource: a detailed comparison](#)

```
movl (%eax), %ecx  
addl $12, %edx  
movl %ecx, -12(%edx)
```

```
mov ecx, DWORD PTR [eax]  
add edx, 12  
mov DWORD PTR [edx-12], ecx
```

AT&T syntax (left), used by GNU Assembler

- instructions suffixed with letters indicating operand sizes
- \$ before intermediate numbers
- % before register names, src operands before dest
- order of operands is *src, dst*

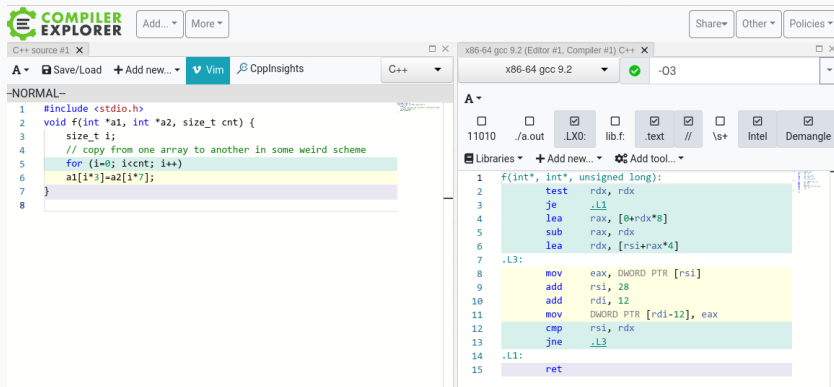
Intel syntax (right)

- order of operands is *dst, src*
- different indirect addressing format

x86 assembly: [Wikibooks:X86 Assembly](#), a one-page quick [tutorial](#).

Tool: Godbolt Compiler Explorer

Whenever you are interested in code generated by compiler x in y architecture with z options, check [Godbolt Compiler Explorer](#).



The screenshot displays the Godbolt Compiler Explorer interface. The left pane shows the C++ source code for a function `f` that copies data from one array to another. The right pane shows the assembly code generated by x86-64 gcc 9.2 with the `-O3` optimization level. The assembly code includes instructions for testing registers, calculating addresses, and moving data.

Compiler Explorer Interface:

- Top Bar:** Compiler Explorer logo, "Add..." and "More" buttons, "Share", "Other", and "Policies" buttons.
- Language/Compiler:** "C++ source #1" and "x86-64 gcc 9.2 (Editor #1, Compiler #1) C++".
- Options:** "Save/Load", "Add new...", "Vim", "Cpplnsights", "C++", "x86-64 gcc 9.2", and "-O3".
- Source Code (Left Pane):**

```
1 #include <stdio.h>
2 void f(int *a1, int *a2, size_t cnt) {
3     size_t i;
4     // copy from one array to another in some weird scheme
5     for (i=0; i<cnt; i++)
6         a1[i*3]=a2[i*7];
7 }
8
```
- Assembly Code (Right Pane):**

```
1 f(int*, int*, unsigned long):
2     test    rdx, rdx
3     je      .L1
4     lea     rax, [0+rdx*8]
5     sub     rax, rdx
6     lea     rdx, [rsi+rax*4]
7 .L3:
8     mov     eax, DWORD PTR [rsi]
9     add     rsi, 28
10    add     rdi, 12
11    mov     DWORD PTR [rdi-12], eax
12    cmp     rsi, rdx
13    jne     .L3
14 .L1:
15    ret
```


A list of command line tools for basic inspection of binary files.

- `file` shows filetype and some basic information about the file
- `strings` prints the sequences of printable characters in files, can help solve easy warmup challenges quickly.
- `xxd`, `hexdump` can create hexdump of files.
- `bvi` is a hex editor, there is also a good non-free GUI hex editor [010 Editor](#).
- `objdump` displays information from object files.
- `readelf` displays information about ELF files.
- [checksec](#) checks the properties of executables.

[Ghidra](#) is a powerful reverse-engineering tool, released under Apache-2.0 license.

Note: On a recent version of macOS, you may need to deal with gatekeeper. (refer to [github/ghidra/issue/1559](https://github.com/NationalSecurityAgency/ghidra/issues/1559))

Resources:

- In-application help appears automatically the first time you start ghidra, or access it through [menu]->[help]->[contents].
- [Presentations](#) by NSA introduce features.
- [Cheat Sheet](#) of shortcuts.

Some alternatives are:

- [Radare2](#) is an open source toolset.
- [IDA Pro](#), the long time state-of-the-art product, but proprietary and expensive. An old feature-limited version is provided [for free](#).
- [Binary Ninja](#) is also not free but comes with a demo version.

Can you write C programs without using headers such as stdio, stdlib?
Try to get flag from [this file](#).

(Hackergame2019 of USTC, released under CC BY-NC-SA 4.0 license.)

Can you write C programs without using headers such as `stdio`, `stdlib`?
Try to get flag from [this file](#).

(Hackergame2019 of USTC, released under CC BY-NC-SA 4.0 license.)

First we can try to run this file.
It just prompts a string.

```
% ./tinyELF
please input flag:
```

`objdump` also gives nothing
useful.

```
% objdump -D tinyELF
tinyELF:      file format elf64-x86-64
```

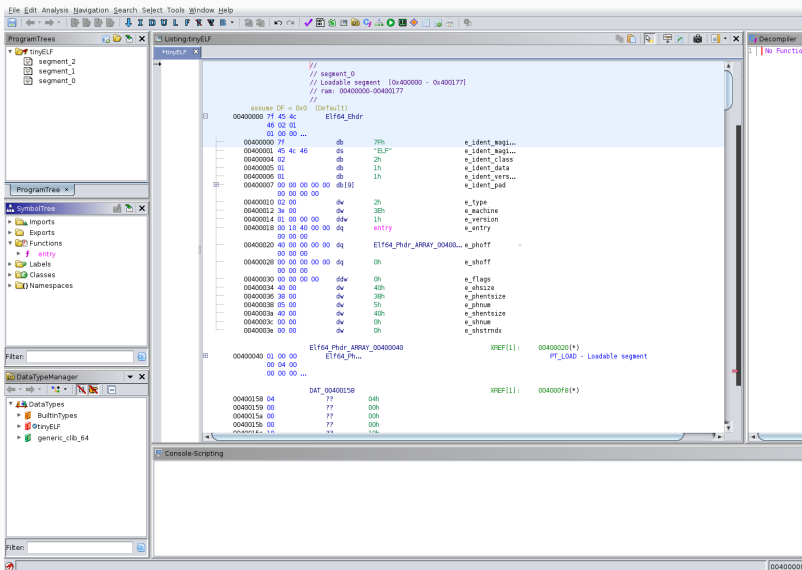
`strings` gives out some hint,
possibly we need to reach the
“correct”.

```
% strings tinyELF
please input flag:
correct
```


Then start Ghidra:

- [new project]->[non-shared project], select a path and name for your project.
- [File]->[Import file], add tinyELF you just downloaded.
- In the active project window, double-click file you want to analyze, tinyELF here.
- A new code browser will appear and ask if you want to analyze this file, click Yes.
- Then you will see disassembler in the middle and decompiler on the right.
- To quickly find solutions, make use of the strings we found using strings.
- [Search]->[For Strings], click search, and a new window pops up, click “correct” and go back to disassembler window.

Usage of Ghidra (cont'd)



Usage of Ghidra (cont'd)



After clicking "correct", ghidra will locate it in the disassembler.

```
s_correct_00402014      xREF[1]:  entry:004011ad(*)
63 6f 72              ds      "correct"
72 65 63
74 00
```

Double clicking entry: 004011ad will bring you to the address in the entry function.

Once you are in a function in the disassembler, ghidra will automatically show decompiled code on the right side window.

Go through the decompiled entry. Focus on the following part:

```
syscall(); syscall(); local_c = 0;
while(local_c < 0x2e) {
    bVar1 = (byte)local_c;
    local_48[local_c] = bVar1 * '\x02' + local_48[local_c];
    local_48[local_c] = bVar1 ^ local_48[local_c];
    local_48[local_c] = local_48[local_c] - bVar1;
    local_c = local_c + 1;
}
local_10 = 0;
while(local_10 < 0x2e) {
    if (local_48[local_10] != local_78[local_10]) {syscall();}
    local_10 = local_10 + 1;
}
syscall(); syscall();
```

In the code, `local_48` is an un-initialized array, while `local_78` is filled with specific value from `local_78[0]` to `local_78[45]`.

Usage of Ghidra (cont'd)

There are several `syscall()`s, we can check what they actually did in the assembly. Each `syscall` has a number, and is called by set specific register with the number. Since this is a x86_64 ELF, we have:

Syscall #	Param 1	Param 2	Param 3	Param 4	Param 5	Param 6
rax	rdi	rsi	rdx	r10	r8	r9

For example, look at assembly around the first `syscall()`:

```
LEA RDX,[s_please_in_put_flag]
MOV RDI,0x1 ; param 1
MOV RSI,RDX ; param 2
MOV RDX,0x13 ; param 3
MOV RAX,0x1
SYSCALL
```

```
// /usr/include/asm/unistd_64.h
#define __NR_read 0
#define __NR_write 1
#define __NR_exit 60
```

Recall the write system call or `man 2 write`:

```
write(int fd, const void *buf, size_t count);
```

So this syscall is actually `write(1, "please in put flag", 19)`.

fd 1 is stdout, so that is how the "please in put flag" appeared.

After working out all the `syscall()`s, we can have clearer code.

```
write(1, "please input flag", 19)
read(0, user_input, 46)
for (i = 0; i < 46; ++i) {
    solve[i] += i*2;
    solve[i] ^= i;
    solve[i] -= i;
}
for (i = 0; i < 46; ++i) {
    if (solve[i] != goal[i])
        _exit(0);
}
write(1, "correct", 7)
_exit(0);
```

Values of `goal[]` are known, and can be found in assembly code.

We just need to find the 46 characters for `solve[0]` to `solve[45]`, so that they equal to `goal[]` after the transformations.

Try to write program to do the inverse calculations.

And the correct flag is `flag{Linux_Syst3m_C4ll_is_4_f4scin4ting_t00ls}`.

(This challenge can also be solved by using [angr](#).)

Try to get flag from [this file](#).

[from picoCTF 2021]

Solving Puzzles with GDB

Try to get flag from [this file](#).

[from picoCTF 2021]

```
.LC1: .string      "Correct! You entered the flag."
.LC2: .string      "No, that's not right."
.L2:
...
leaq   -272(%rbp), %rcx ; load rbp-272 addr in %rcx
leaq   -208(%rbp), %rax ; load rbp-208 addr in %rax
movl   $49, %edx      ; set %edx to number 49
movq   %rcx, %rsi      ; set %rsi to the content in %rcx
movq   %rax, %rdi      ; set %rdi to the content in %rax
call   memcmp@PLT      ; memcmp() compare %rsi and %rdi
testl  %eax, %eax      ; check return val of memcmp()
jne     .L4 ; if memcmp() does not ret 0, i.e., %rsi != %rdi (49 bytes)
leaq   .LC1(%rip), %rdi
call   puts@PLT ; print out "Correct"
movl   $0, %eax

.L4:
leaq   .LC2(%rip), %rdi
call   puts@PLT ; print out "No"
movl   $1, %eax
```

By locating the two strings indicating a "correct" or "incorrect" branch, we discover the call to `memcmp()` which is condition determining which branch to take.

Go back to check the content on stack, user inputs are read using `fgets()` and stored starting from `rbp-208`.

```
movq stdin(%rip), %rdx ; %rdx: FILE *st
leaq -208(%rbp), %rax ; load rbp-208 address in %rax movl $50, %esi ; %esi: int size
movq %rax, %rdi ; %rdi: char* s
call fgets@PLT ; fgets(char*s, int size, FILE *st)
```

So what contents do `rbp-272` hold?

We can analyze the rest assembly to work out the transformations on `rbp-272`, but a dynamic solution is easier for this puzzle.

Solving Puzzles with GDB (cont'd)

Go back to check the content on stack, user inputs are read using `fgets()` and stored starting from `rbp-208`.

```
movq stdin(%rip), %rdx ; %rdx: FILE *st
leaq -208(%rbp), %rax ; load rbp-208 address in %rax movl $50, %esi ; %esi: int size
movq %rax, %rdi ; %rdi: char* s
call fgets@PLT ; fgets(char*s, int size, FILE *st)
```

So what contents do `rbp-272` hold?

We can analyze the rest assembly to work out the transformations on `rbp-272`, but a dynamic solution is easier for this puzzle.

Note a difference between this puzzle and the previous ghidra one is:

In this puzzle, the transformations are done on the constants while in the ghidra one, the transformations are done on the user input. Therefore, this gdb method cannot be applied directly on that puzzle we solved using ghidra.

Solving Puzzles with GDB (cont'd)

```
% gdb chall
(gdb) run # program wait for user input, press Control-C here
^C Program received signal SIGINT, Interrupt.
(gdb) next
asdfasdf # random input
(gdb) finish # use finish until we are back in the main()
(gdb) disassemble # check the assembly code to get desired breakpoint
# You can find sth. like this:
# 0000555555552f0 <+379>: call 0x55555555060 <memcmp@plt>
(gdb) break *0x55555555060
(gdb) c # continue to the breakpoint
(gdb) x/49xb $rsi # refer to previous slides on why the length is 49
0x7fffffff6b0: 0x70 0x69 0x63 0x6f 0x43 0x54 0x46 0x7b
0x7fffffff6b8: 0x64 0x79 0x6e 0x34 0x6d 0x31 0x63 0x5f
0x7fffffff6c0: 0x34 0x6e 0x34 0x6c 0x79 0x31 0x73 0x5f
0x7fffffff6c8: 0x31 0x73 0x5f 0x35 0x75 0x70 0x33 0x72
0x7fffffff6d0: 0x5f 0x75 0x73 0x33 0x66 0x75 0x6c 0x5f
0x7fffffff6d8: 0x36 0x30 0x34 0x34 0x65 0x36 0x36 0x30
0x7fffffff6e0: 0x7d
```

Convert hex to ASCII characters, you can get the flag:

picoCTF{dyn4m1c_4n4ly1s_1s_5up3r_us3ful_6044e660}.

Forensics

Steganography usually involves finding information hidden in various types of files. The most common steganography challenges are

- including the flags in an archive file (try to [solve this](#) and [this from Google CTF 2018 QualsBeginners Quest](#))
- modifying the length field in the metadata (similar to the challenge using an uncommon encoding on page [5](#)).
- Hiding flags in image files is quite popular. Information can be inserted in not only metadata of images but also [lower bits of “actual image data”](#).(Click “how it works”, if you do not understand the explanation, read about [RGB](#))

Try this: You are given an [image](#). Try to extract the flag. (Hint: use the lower bits trick) [Solution \[PDF\]](#) [↗](#)

Another Simple Challenge

Try to extract flag from this [filesystem data](#).

(Google CTF 2018 Quals Beginners Quest.)

Another Simple Challenge

Try to extract flag from this [filesystem data](#).

(Google CTF 2018 Quals Beginners Quest.)

- file tells that it is a gzip, `gzip -d challenge.ext4.gz`
- mount the ext4 file, `sudo mount -o loop challenge.ext4 /tmp/mount`
(Always be careful when you use sudo and only when you know what you are doing.)
- See the `.mediapc_backdoor_password.gz` file, uncompress it using `gzip -d` and you get the flag.

Appendix

- [RE4B](#) is a free book about reverse engineering.
- Refer to [list of file signatures](#) when checking file magic bytes.
- [Awesome-ctf on github](#) is a curated list of CTF-related resources
- [A YouTube channel](#) related to CMU picoCTF competition.
- [TLDR](#) is an utility for showing common usage of command line tools.